



Recommendation Technologies for Multimedia Content

Xiangnan He

Hanwang Zhang

Tat-Seng Chua

National University
of Singapore

Nanyang Technological
University

National University of
Singapore

Slides are available: <http://comp.nus.edu.sg/~xiangnan/icmr18-recsys.pdf>

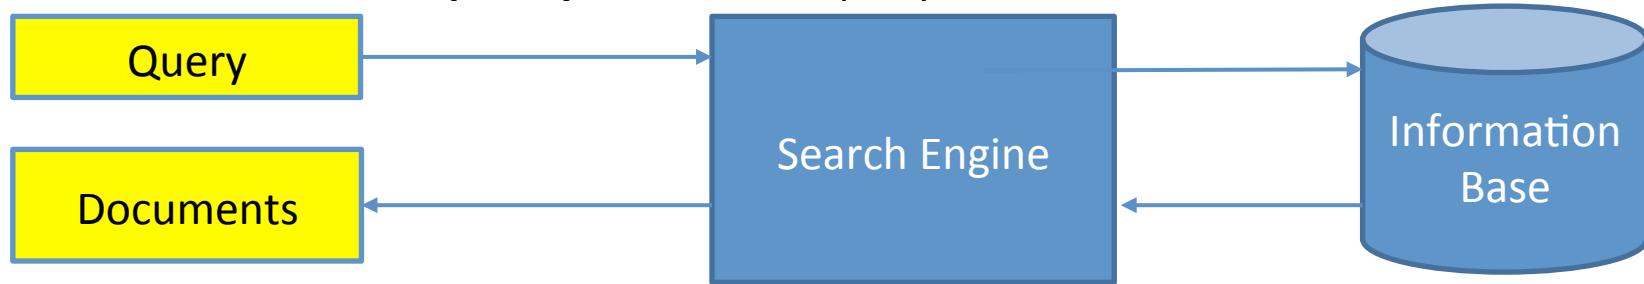
Outline of Tutorial

- Background (Xiangnan, 10 mins)
- Basics & Advances in Recommendation (Xiangnan, 50 mins)
- Visually-aware Product Recommendation (Xiangnan, 30 mins)
- Break (15 mins)
- Visual Representation (Hanwang, 45 mins)
- Image/Video Recommendation (Hanwang, 25 mins)
- Conclusion (Hanwang, 5 mins)

Retrieval vs. Recommendation

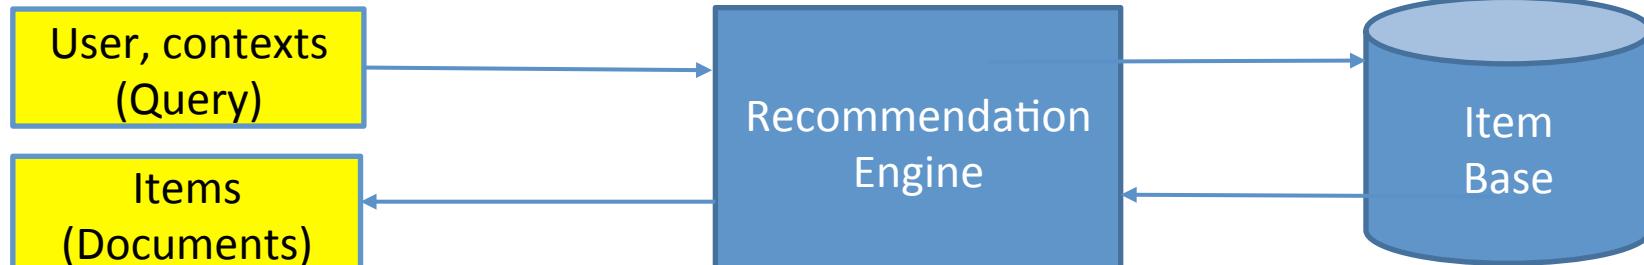
- Retrieval is **information pull**:
 - User **pulls** desired information by making a specific request

User intent is **explicitly** reflected in query



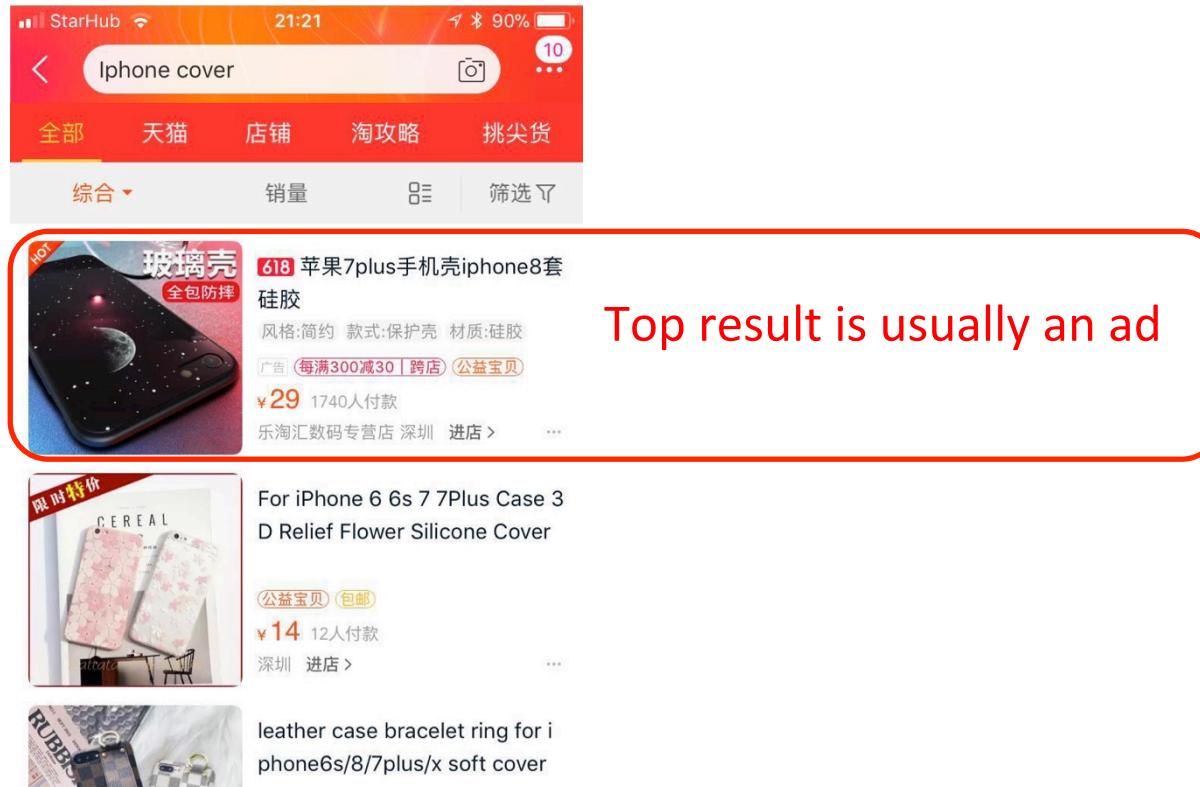
- Recommendation is **information push**:
 - System **pushes** desired information to a user by guessing her interest

User intent is **implicitly** reflected in interaction history, profile, contexts etc.



Importance of Recommendation

- Retrieval mostly exists in search engines
- But, recommendation exists everywhere...
 - When you search for a product <= Ad recommendation



Search results of Taobao

Slides: <http://comp.nus.edu.sg/~xiangnan/icmr18-recsys.pdf>

Importance of Recommendation

- Retrieval mostly exists in search engines
- But, recommendation exists everywhere...
 - When you search for a product => Ad recommendation
 - When you open a product page => Product recommendation



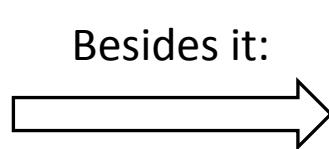
Screenshot of Amazon

Importance of Recommendation

- Retrieval mostly exists in search engines
- But, recommendation exists everywhere...
 - When you search for a product => Ad recommendation
 - When you open a product page => Product recommendation
 - When you watch a video => Video Recommendation



President Trump reads letter from Kim Jong Un
(YouTube)



Up next

AUTOPLAY

'Kim Jong Un': Trump, I'm here in Singapore!
The Star Online
108K views
New
2:28

How Donald Trump Answers A Question
Nerdwriter1
7.3M views
6:54

North Korea's Kim Jong Un Calls President Donald Trump...
TODAY
3M views
3:03

Importance of Recommendation

- Retrieval mostly exists in search engines
- **But, recommendation exists everywhere...**
 - When you search for a product => Ad recommendation
 - When you open a product page => Product recommendation
 - When you watch a video => Video Recommendation
 - When you read a news => News recommendation
 - When you book a flight => Hotel Recommendation
 - When you use social network => Friend Recommendation
 - When you are hungry => Restaurant Recommendation
 - When you open any webpage/app, there maybe a recommendation list.

... ...

Value of Recommender System (RecSys)

- RecSys has become a major **monetization** tool for customer-oriented online services
 - E.g., E-commerce, News Portal, Social Networks, etc.
- Ad systems are technically supported by recommendation solutions.
 - The key is Click-Through Rate (CTR) prediction
- Some statistics:
 - YouTube homepage: 60%+ clicks [Davidson et al. 2010]
 - Netflix: 80%+ movie watches, 1billion+ value/year [Gomze-Uribe et al 2016]
 - Amazon: 30%+ page views [Smith and Linden, 2017]

Why RecSys + MultiMedia?

- Multimedia contents are prevalent in Web, generated and consumed in a fast speed.
 - YouTube, Pinterest, Snap etc.

=> Need dedicated RecSys for multimedia contents (e.g., images, videos)
- A picture is worth a thousand words
 - Visual signal is crucial to attract users perform an action.
 - Many non-multimedia items are affiliated with image/video for better explanation, e.g., news, products, ads

=> Need enhanced RecSys to incorporate visual signal

Some Useful Resources

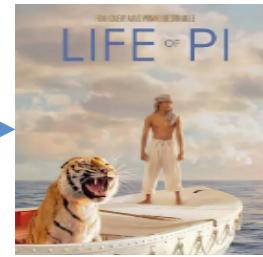
- Recent challenges:
 - MediaEval 2018 on content-based movie recommendation:
<http://www.multimediaeval.org/mediaeval2018/content4recsys/>
 - ACM MM Challenge 2018 on social media headline prediction:
<https://social-media-prediction.github.io/PredictionChallenge/>
- Datasets:
 - Pinterest images:
<https://sites.google.com/site/xueatalphabeta/academic-projects>
 - Amazon products (with images):
<http://jmcauley.ucsd.edu/data/amazon/>

Outline of Tutorial

- Background (Xiangnan, 10 mins)
- **Basics & Advances in Recommendation (Xiangnan, 50 mins)**
 - Traditional Shallow Learning Methods
 - Recent Deep Learning Methods
- Visually-aware Product Recommendation (Xiangnan, 30 mins)
- Break (15 mins)
- Visual Representation (Hanwang, 45 mins)
- Visual Recommendation (Hanwang, 25 mins)
- Conclusion (Hanwang, 5 mins)

Problem Formulation

- Recommendation solves a **matching** problem.



User Profile (query):

- User ID
- Rating history
- Age, Gender
- Income level
- Time of the day

.....

Item Profile (document):

- Item ID
- Description
- Category
- Price
- Image

.....

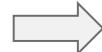
Challenge: **no overlap** between user features and item features
Matching can't be done on the superficial feature level!

Collaborative Filtering

- Collaborative Filtering (CF) is the most well-known technique for recommendation.
*“CF makes predictions (**filtering**) about a user’s interest by collecting preferences information from many users (**collaborating**)” ---Wikipedia*
- Math formulation: matrix completion problem

User	Movie	Rating
Alice	Titanic	5
Alice	Notting Hill	3
Alice	Star Wars	1
Bob	Star Wars	4
Bob	Star Trek	5
Charlie	Titanic	1
Charlie	Star Wars	5
...

Input Tabular data



		Movie				
		TI	NH	SW	ST	...
User	A	5	3	1	?	...
	B	?	?	4	5	...
	C	1	?	5	?	...

Rating Matrix
(Interaction Matrix)

Solving Matrix Completion

- Singular Value Decomposition (SVD) is the most well-known technique for matrix completion

	Movie					
	TI	NH	SW	ST	...	
A	5	3	1	?	...	
B	?	?	4	5	...	
C	1	?	5	?	...	
...	

Rating Matrix

$$\mathbf{M}_{m \times n} = \mathbf{U}_{m \times m} \mathbf{\Sigma}_{m \times n} \mathbf{V}^*_{n \times n}$$

Orthonormal Orthonormal
Diagonal

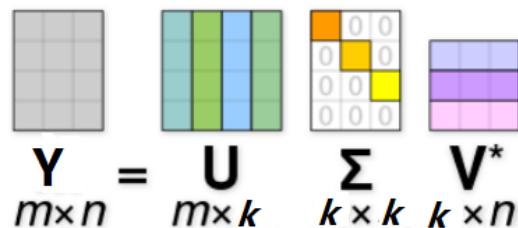
Steps to use SVD for CF:

1. Impute missing data to 0 in \mathbf{Y}
2. Solving the SVD problem
3. Using only K dimensions in \mathbf{U} and \mathbf{V} to obtain a low rank model to estimate \mathbf{Y}

Figure adopted from: https://en.wikipedia.org/wiki/Singular-value_decomposition

14

SVD is Suboptimal for CF

$$\mathbf{Y}_{m \times n} = \mathbf{U}_{m \times k} \mathbf{\Sigma}_{k \times k} \mathbf{V}^*_{k \times n}$$


- In essence, SVD is solving the problem:

$$\begin{aligned} & \arg \min_{\mathbf{U}, \mathbf{\Sigma}, \mathbf{V}} (\mathbf{Y} - \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T)^2 \\ &= \arg \min_{\mathbf{U}, \mathbf{\Sigma}, \mathbf{V}} \sum_{i=1}^m \sum_{j=1}^n (\text{Label } [y_{ij}] - \text{Model Prediction } (\mathbf{U}\mathbf{\Sigma}\mathbf{V}^T)_{ij})^2 \end{aligned}$$

Training instance

- Several Implications (weaknesses):

1. Missing data has the same weight as observed data (>99% sparsity)
2. No regularization is enforced – easy to overfit

Adjust SVD for CF

- The “SVD” model in the context of recommendation:

$$\hat{y}_{ui} = \underbrace{\mathbf{v}_u^T \mathbf{v}_i}_{\substack{\text{User latent vector} \\ \text{Item latent vector}}}$$

- Regularized Loss function:

$$L = \underbrace{\sum_u \sum_i w_{ui} (y_{ui} - \hat{y}_{ui})^2}_{\substack{\text{Prediction error}}} + \lambda \left(\underbrace{\sum_u \|\mathbf{v}_u\|^2 + \sum_i \|\mathbf{v}_i\|^2}_{\substack{\text{L2 regularizer}}} \right)$$

- This method is also called *Matrix Factorization* (MF) in RecSys:
 - It represents a user and an item as a latent vector (**ID embedding**).
 - The interaction between user and item is modelled using **inner product** (measure how much user latent “preferences” match with item “properties”)
 - Besides L2 regularized loss, other loss can also be used, e.g., cross-entropy, margin-based pairwise loss, etc.

Why MF Can Capture CF

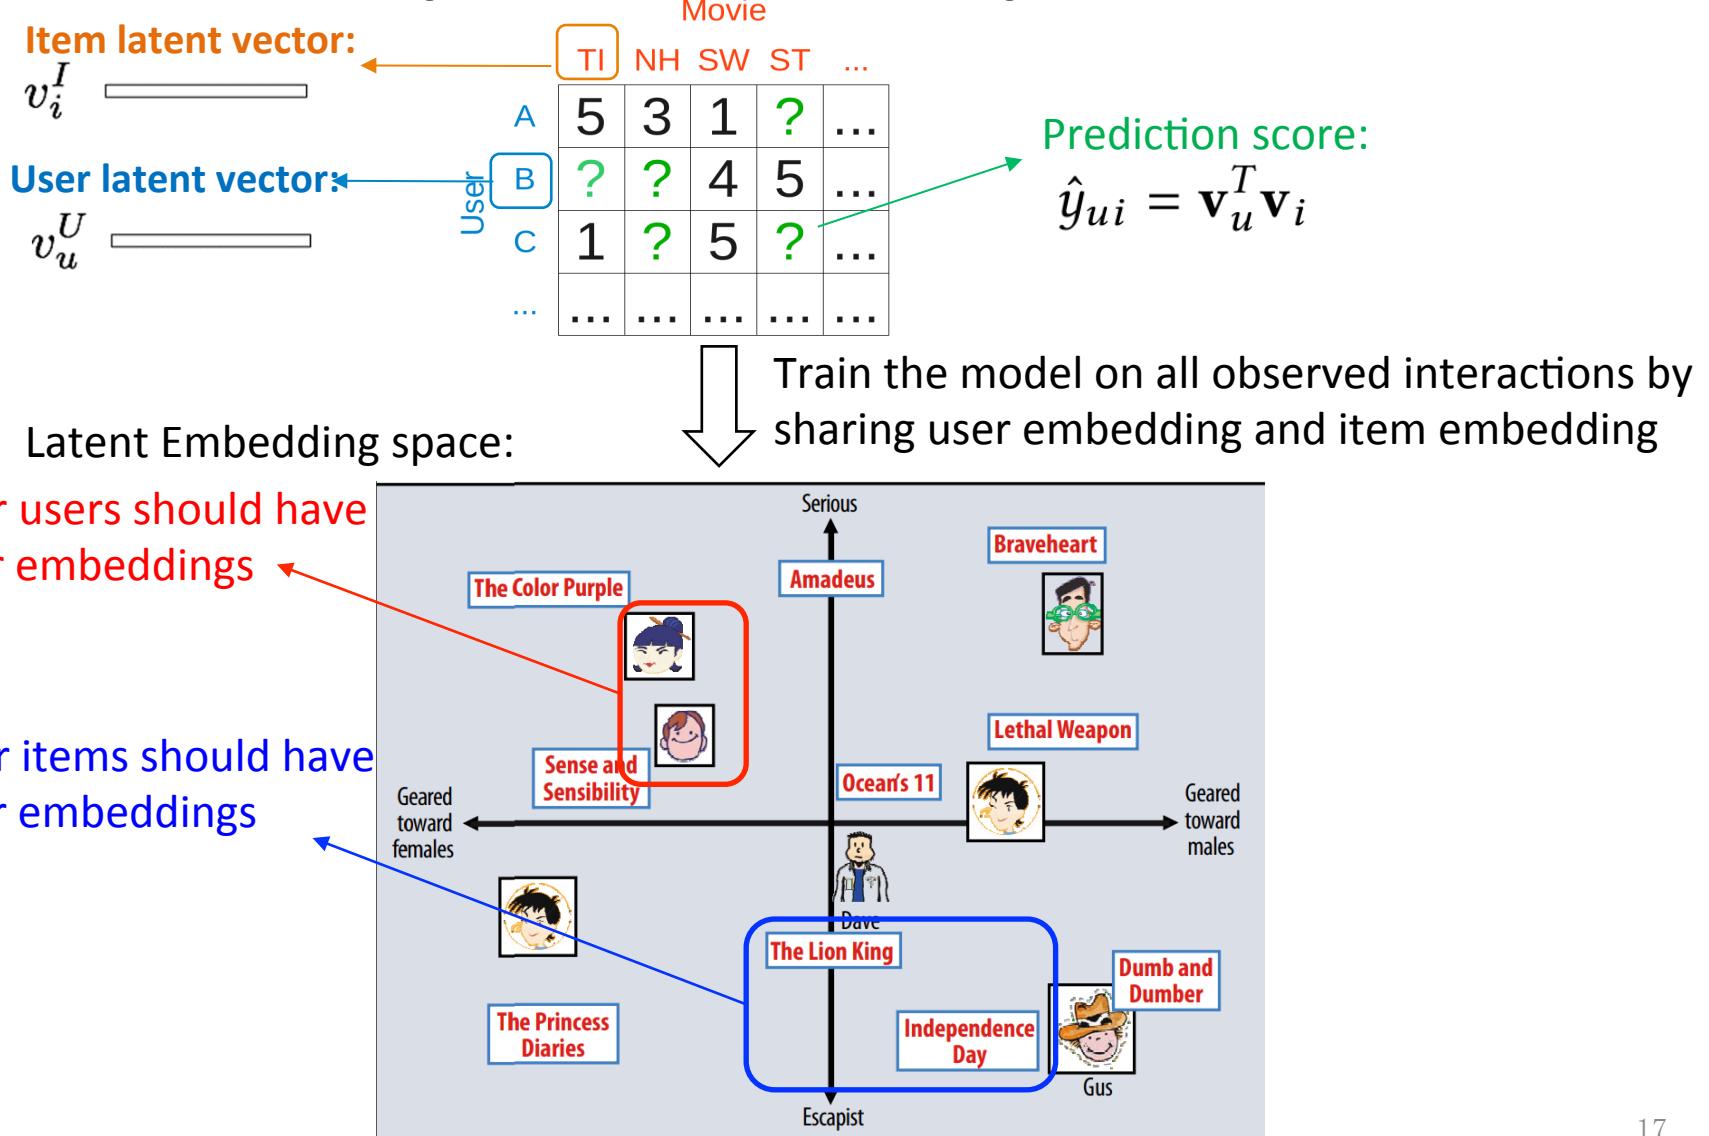
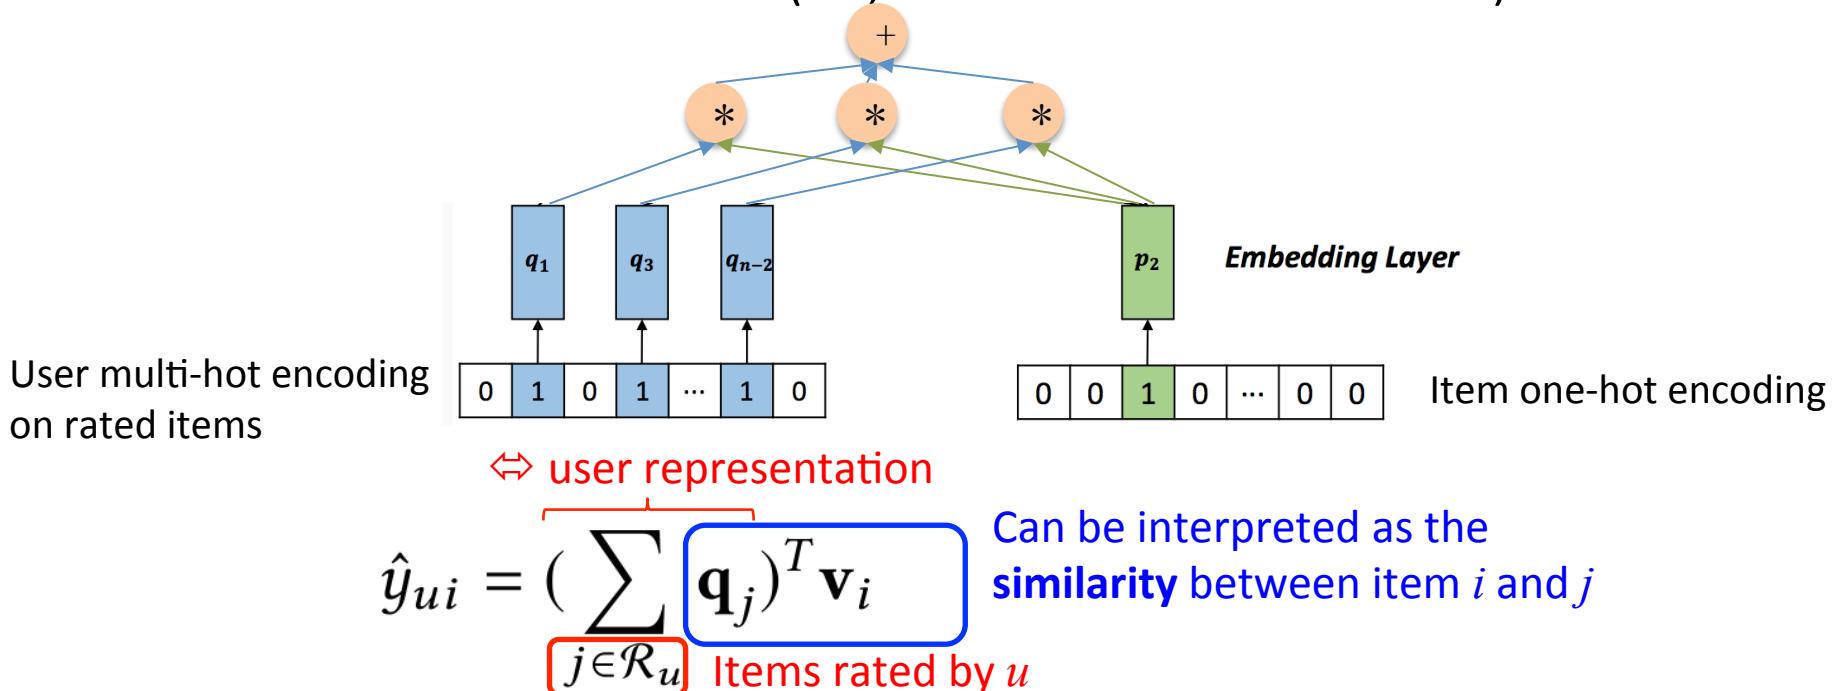


Figure adopted from: <https://datajobs.com/data-science-repo/Recommender-Systems-%5BNetflix%5D.pdf>

Factored Item Similarity Model

(Kabbur et al., KDD'14)

- MF encodes a user with an ID, and projects it to embedding.
 - Also called as user-based CF (i.e., find similar users for recom)
- Another more **meaningful** encoding is to use **rated items** of the user.
 - Also called as item-based CF (i.e., find similar items for recom)



SVD++: Fusing User-based and Item-based CF (Koren, KDD'08)

- MF (user-based CF) represents a user as her ID.
 - Directly projecting the ID into latent space
- FISM (item-based CF) represents a user as her interacted items.
 - Projecting interacted items into latent space
- SVD++ fuses the two types of models in the latent space:

$$\hat{y}_{ui} = (\mathbf{v}_u + \underbrace{\sum_{j \in \mathcal{R}_u} \mathbf{q}_j}_{\text{User representation in latent space}})^T \mathbf{v}_i$$

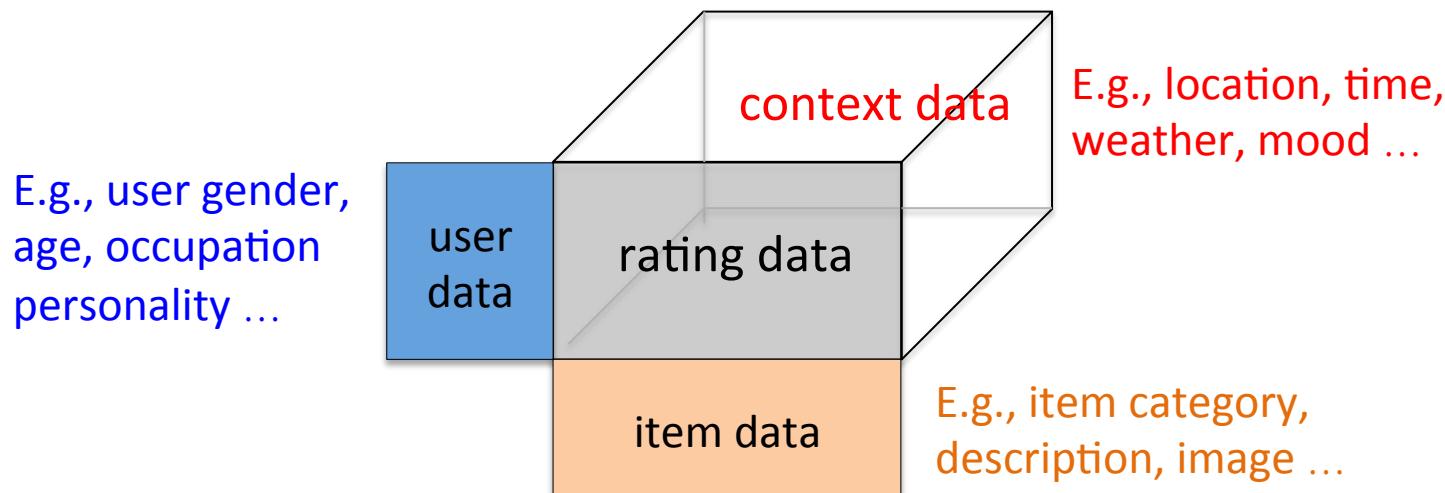
User representation in latent space

- This is the best single model for rating prediction in the Netflix challenge.

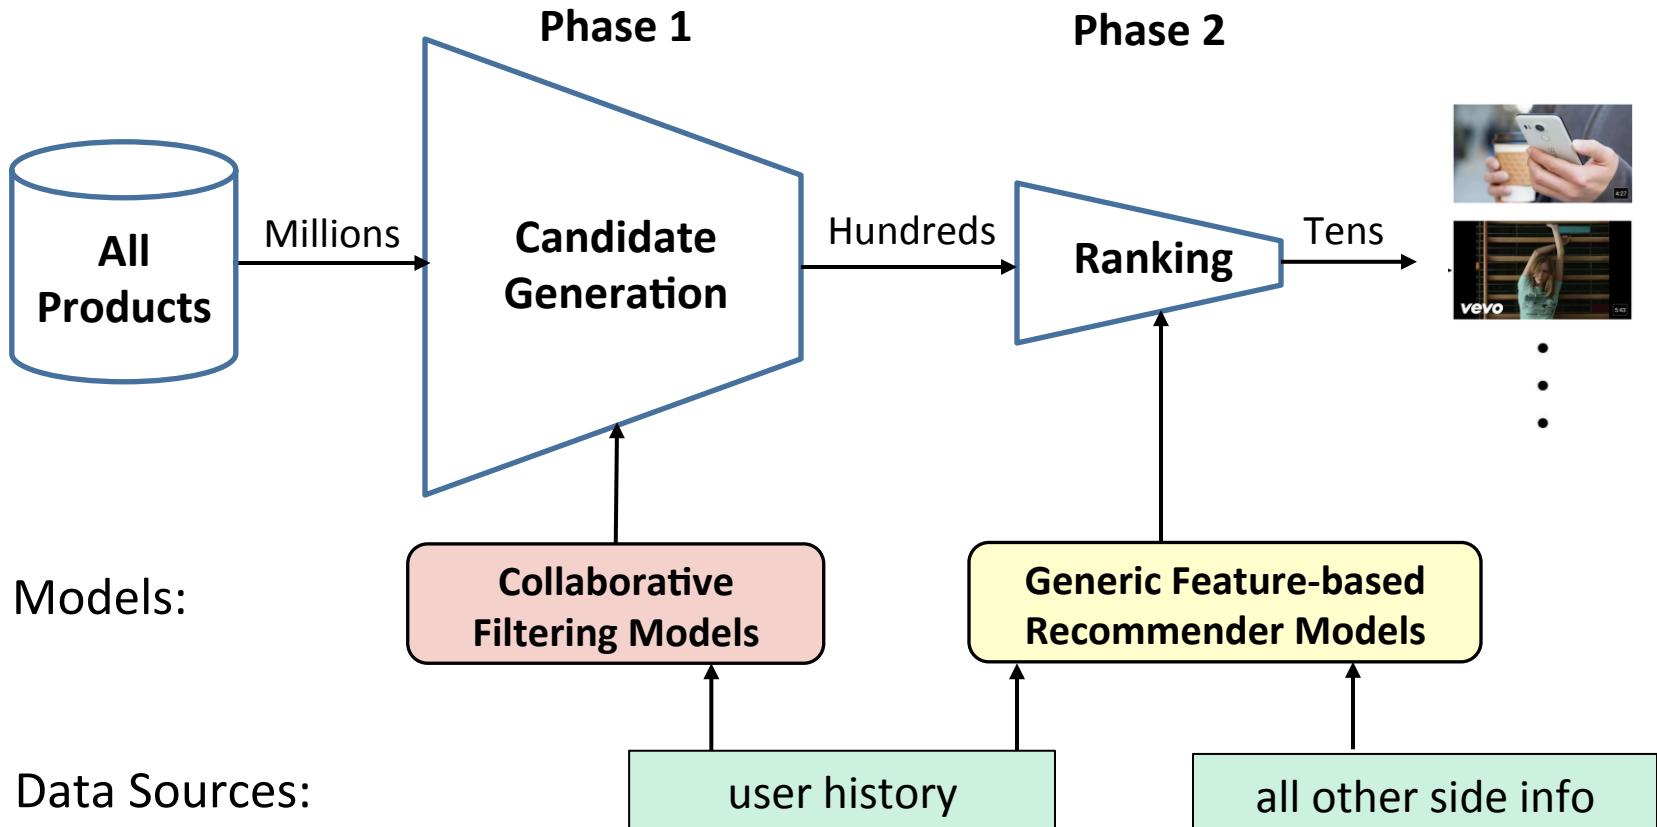
Note: the normalization terms are discarded for clarity.

Generic Feature-based Recommendation

- CF utilizes only the interaction matrix only to build the predictive model.
- How about other information like user/item attributes and contexts?
- Example data used for building a RecSys:

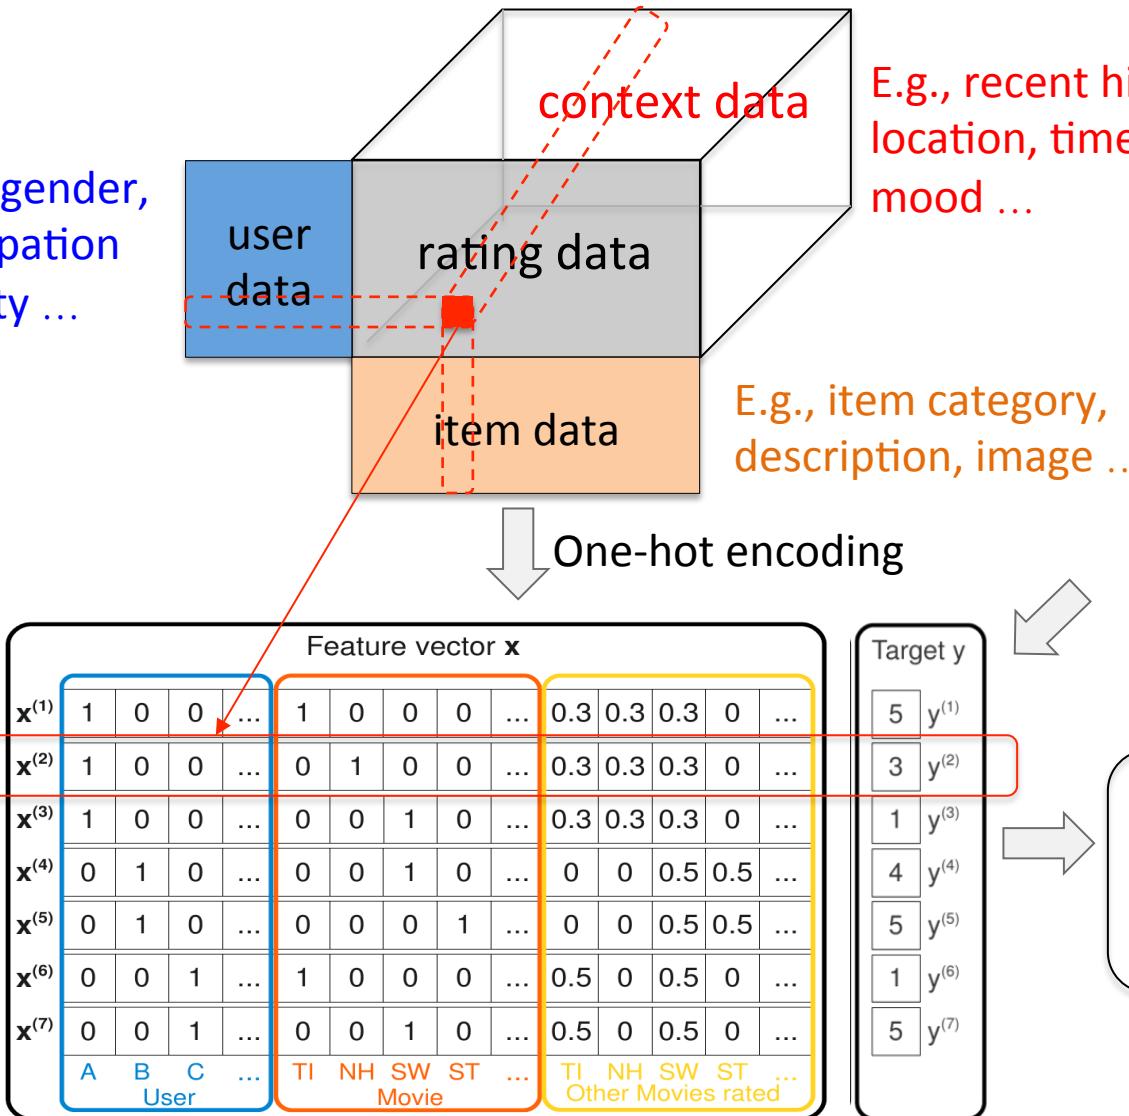


Morden RecSys Architecture



Generic Feature-based Recommendation

E.g., user gender,
age, occupation
personality ...



E.g., recent history,
location, time, weather,
mood ...

E.g., item category,
description, image ...

- Input Features:
1. Categorical features:
user/item ID, bag-of-words, historical features...
 2. Numerical features:
textual/visual embeddings, converted features (e.g. TFIDF, GBDT)...

Sparse Predictive Model

FM: Factorization Machine (Rendle, ICDM'10)

- FM is inspired from previous factorization models
- It represents each feature an embedding vector, and models the second-order feature interactions:

$$\hat{y}(\mathbf{x}) = w_0 + \sum_{i=1}^p w_i x_i + \underbrace{\sum_{i=1}^p \sum_{j>i}^p \langle \mathbf{v}_i, \mathbf{v}_j \rangle}_{\text{Second-order: pair-wise interactions between features}} x_i x_j$$

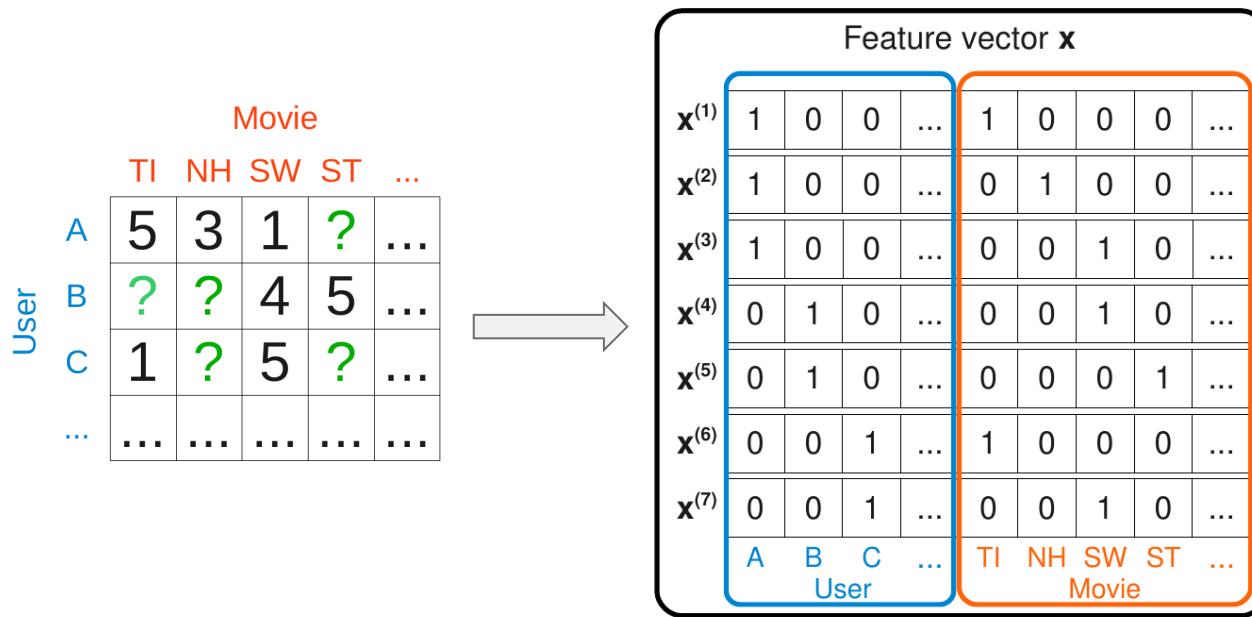
Only nonzero features are considered

First-order: Linear Regression

- Note: self-interaction is not included: $\cancel{\langle \mathbf{v}_i, \mathbf{v}_i \rangle}$.
- FM allows easy feature engineering for recommendation, and can mimic many existing models (that are designed for a specific task) by inputting different features.
 - E.g., MF, SVD++, timeSVD (Koren, KDD'09), PITF (Rendle, WSDM'10) etc.

Matrix Factorization with FM

- Input: 2 variables <user (ID), item (ID)>.

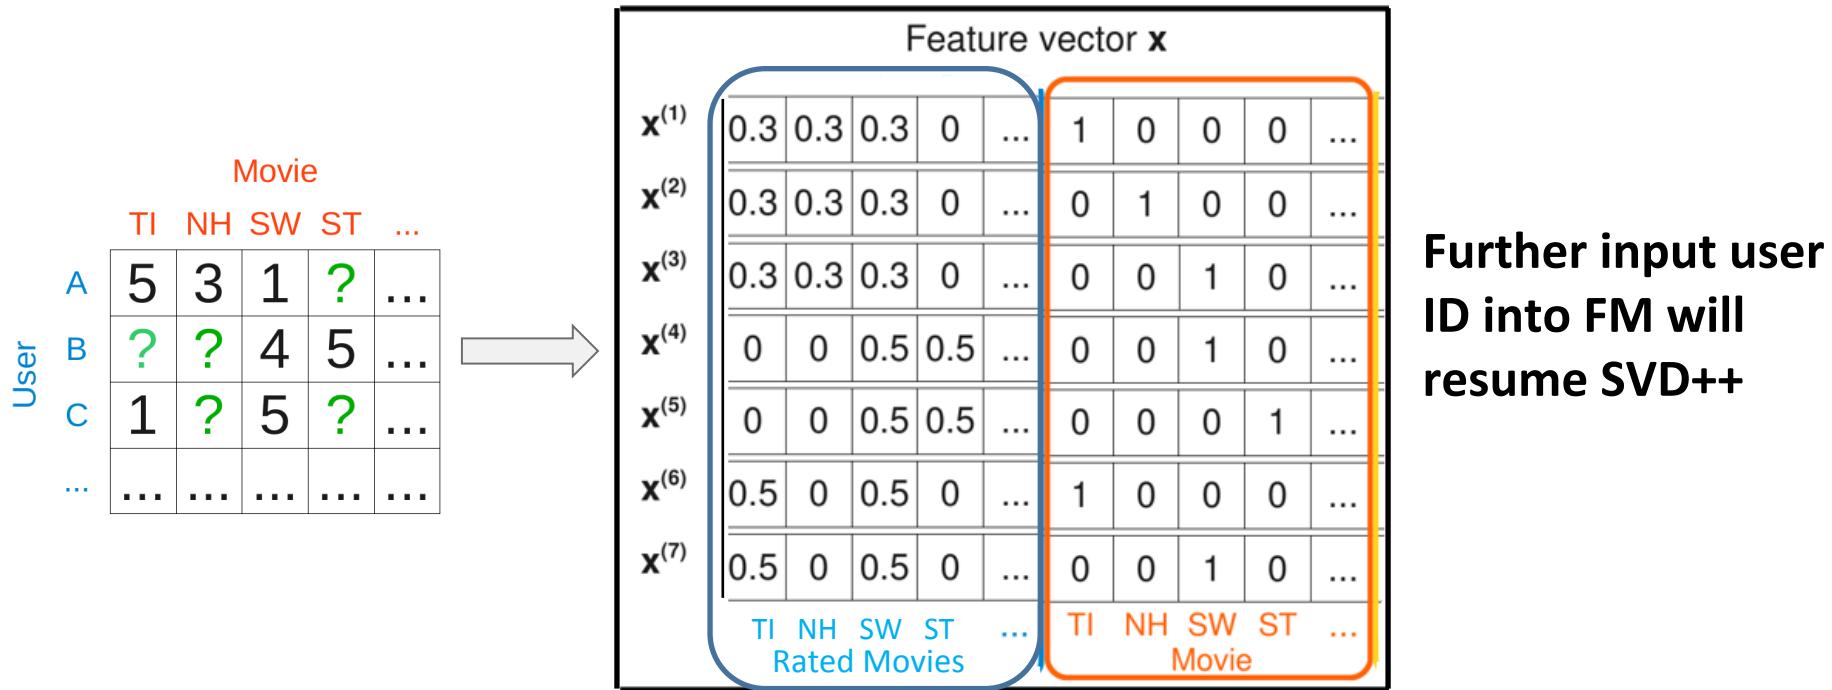


With this input, FM is identical to MF with bias:

$$\hat{y}(\mathbf{x}) = w_0 + w_u + w_i + \underbrace{\langle \mathbf{v}_u, \mathbf{v}_i \rangle}_{\text{MF}}$$

Factored Item Similarity Model with FM

- Input: 2 variables <user (historical items ID), item (ID)>.



With this input, FM subsumes FISM with additional terms:

$$\hat{y}(\mathbf{x}) = bias + \underbrace{\sum_{j \in \mathcal{R}_u} \langle \mathbf{v}_j, \mathbf{v}_i \rangle}_{\text{FISM}} + \sum_{j \in \mathcal{R}_u, j' > j} \langle \mathbf{v}_j, \mathbf{v}_{j'} \rangle$$

Explicit Feedback vs. Implicit Feedback

Explicit Feedback

		Movie				
		TI	NH	SW	ST	...
User	A	5	3	1	?	...
B	?	?	4	5	...	
C	1	?	5	?	...	
...

Ratings

Implicit Feedback

		Movie				
		TI	NH	SW	ST	...
User	A	1	1	1	?	...
B	?	?	1	1	...	
C	1	?	1	?	...	
...

Watches, Clicks, Purchases ...

Explicit Feedback conveys user preference **explicitly**:

- Higher scores carry positive signal
- Lower scores carry negative signal

Implicit Feedback conveys user preference **implicitly**:

- Observed interactions do not mean positive signal
- Unobserved interactions do not mean negative signal

Rating Prediction vs. Ranking

- Old work on recommendation optimize L2 loss:

$$L = \sum_u \sum_i w_{ui} (y_{ui} - \hat{y}_{ui})^2 + \lambda (\sum_u \|\mathbf{v}_u\|^2 + \sum_i \|\mathbf{v}_i\|^2)$$

- But many empirical evidence show that:
 - A lower error rate does not lead to a good ranking performance...
- Possible Reasons:
 - 1) Discrepancy between error measure (e.g., RMSE) and ranking measure.
 - 2) Observation bias – users tend to rate on the items they like.

- Modern work on recommendation optimize pairwise ranking loss:

$$L_{BPR} = \arg \max_{\Theta} \sum_{(u, i, j) \in \mathcal{R}_B} \ln \sigma(\hat{y}_{ui} - \hat{y}_{uj}) - \lambda \|\Theta\|^2$$

sigmoid Positive prediction Negative prediction

Pairwise training examples: u prefers i over j

- Known as the *Bayesian Personalized Ranking* loss (Rendle, UAI'09).
- It optimizes relative ranking between two items, rather than absolute scores

References

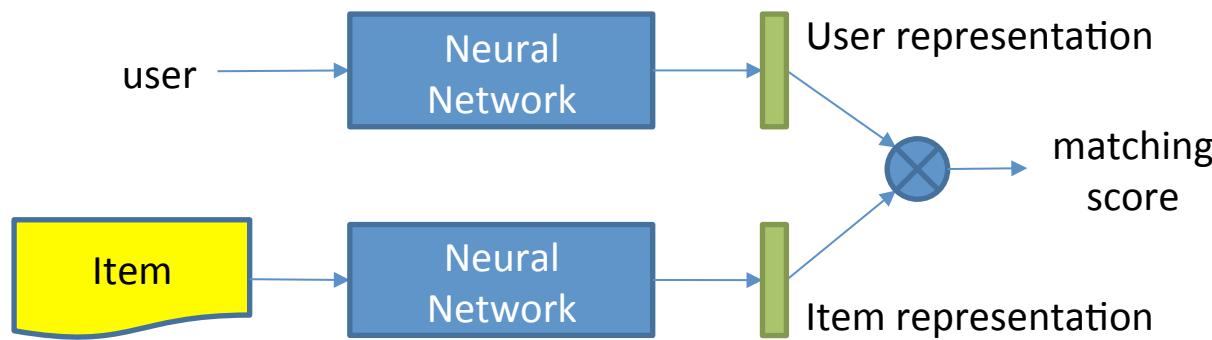
- https://en.wikipedia.org/wiki/Collaborative_filtering
- Xiangnan He, Hanwang Zhang, Min-Yen Kan, and Tat-Seng Chua. Fast matrix factorization for online recommendation with implicit feedback. In SIGIR 2016.
- Yehuda Koren, and Robert Bell. Advances in collaborative filtering. Recommender systems handbook. Springer, Boston, MA, 2015. 77-118.
- Santosh Kabbur, Xia Ning, and George Karypis. Fism: factored item similarity models for top-n recommender systems. In KDD 2013.
- Yehuda Koren. Factorization meets the neighborhood: a multifaceted collaborative filtering model. In KDD 2018.
- Steffen Rendle. Factorization machines. In ICDM 2010.
- Yehuda Koren. Collaborative filtering with temporal dynamics. Communications of the ACM 53, no. 4 (2010): 89-97.
- Steffen Rendle, and Lars Schmidt-Thieme. Pairwise interaction tensor factorization for personalized tag recommendation. In WSDM 2010.
- Immanuel Bayer, Xiangnan He, Bhargav Kanagal, and Steffen Rendle. A generic coordinate descent framework for learning from implicit feedback. In WWW 2017.
- Xiangnan He, Lizi Liao, Hanwang Zhang, Liqiang Nie, Xia Hu, and Tat-Seng Chua. Neural collaborative filtering. In WWW 2017.
- Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and Lars Schmidt-Thieme. BPR: Bayesian personalized ranking from implicit feedback. In UAI 2009.

Outline of Tutorial

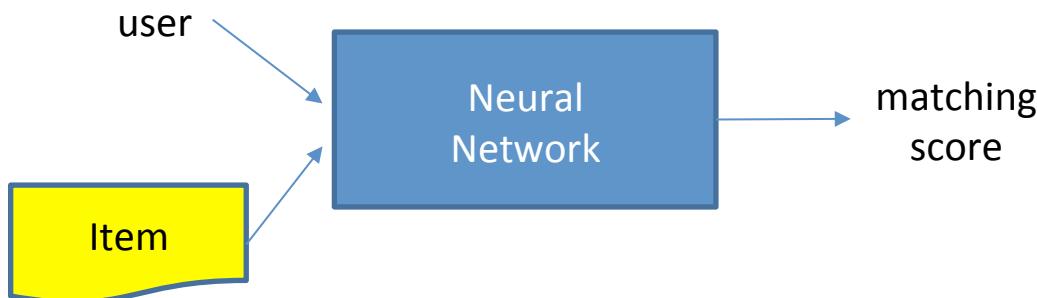
- Background (Xiangnan, 10 mins)
- **Basics & Advances in Recommendation (Xiangnan, 50 mins)**
 - Traditional Shallow Learning Methods
 - Recent Deep Learning Methods
 - 1) Deep Collaborative Filtering methods**
 - 2) Deep Feature-based Recommender Models**
- Visually-aware Product Recommendation (Xiangnan, 30 mins)
- Break (15 mins)
- Visual Representation (Hanwang, 45 mins)
- Image/Video Recommendation (Hanwang, 25 mins)
- Summary (Hanwang, 5 mins)

Deep Learning Models for CF

- Methods of **representation learning**

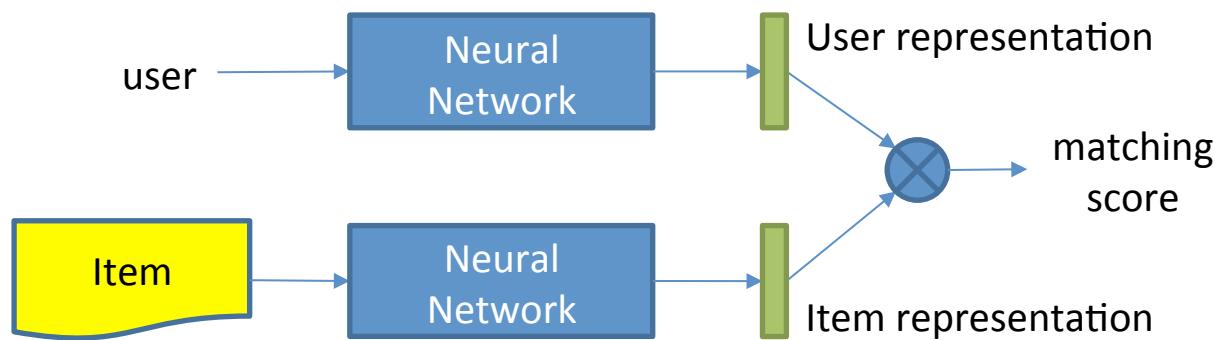


- Methods of **matching function learning**



Next ...

- Methods of **representation learning**

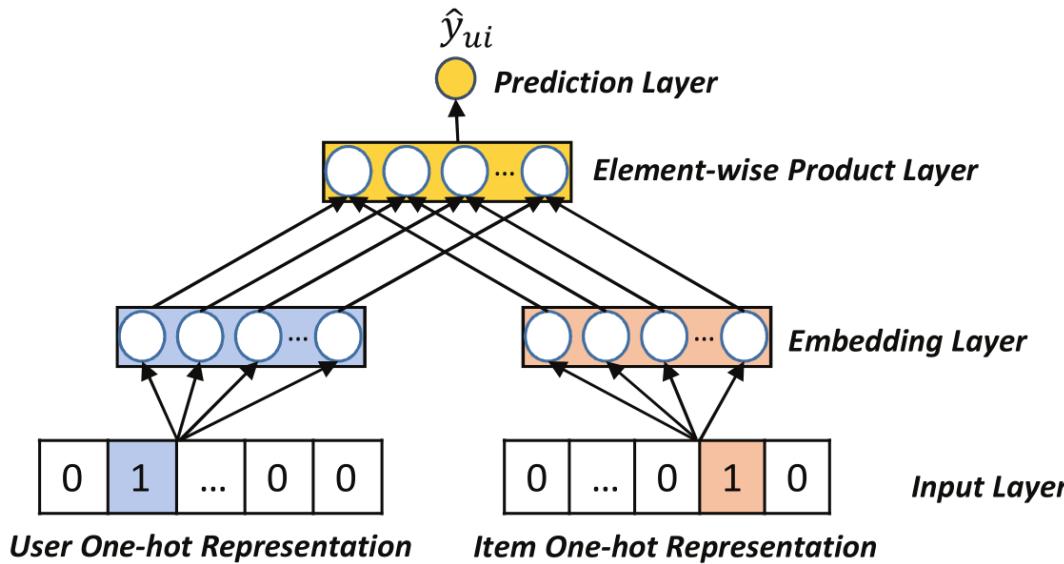


- **DeepMF**: Deep Matrix Factorization (Xue et al, IJCAI'17)
- **AutoRec**: Autoencoders Meeting CF (Sedhain et al, WWW'15)
- **CDAE**: Collaborative Denoising Autoencoder (Wu et al, WSDM'16)

Matrix Factorization as a Neural Network (Wang et al, SIGIR'17)

- **Input:** user \rightarrow ID (one-hot), item \rightarrow ID (one-hot).
- **Representation Function:** linear embedding layer.
- **Matching Function:** inner product.

$$f_{MF}(u, i | \mathbf{p}_u, \mathbf{q}_i) = \mathbf{p}_u^\top \mathbf{q}_i = \sum_{k=1}^K p_{uk} q_{ik},$$



Deep Matrix Factorization (Xue et al, IJCAI'17)

- **Input:**

user -> **items that she has rated** (multi-hot), i.e., row vector of \mathbf{Y}

indicates the user's global preference

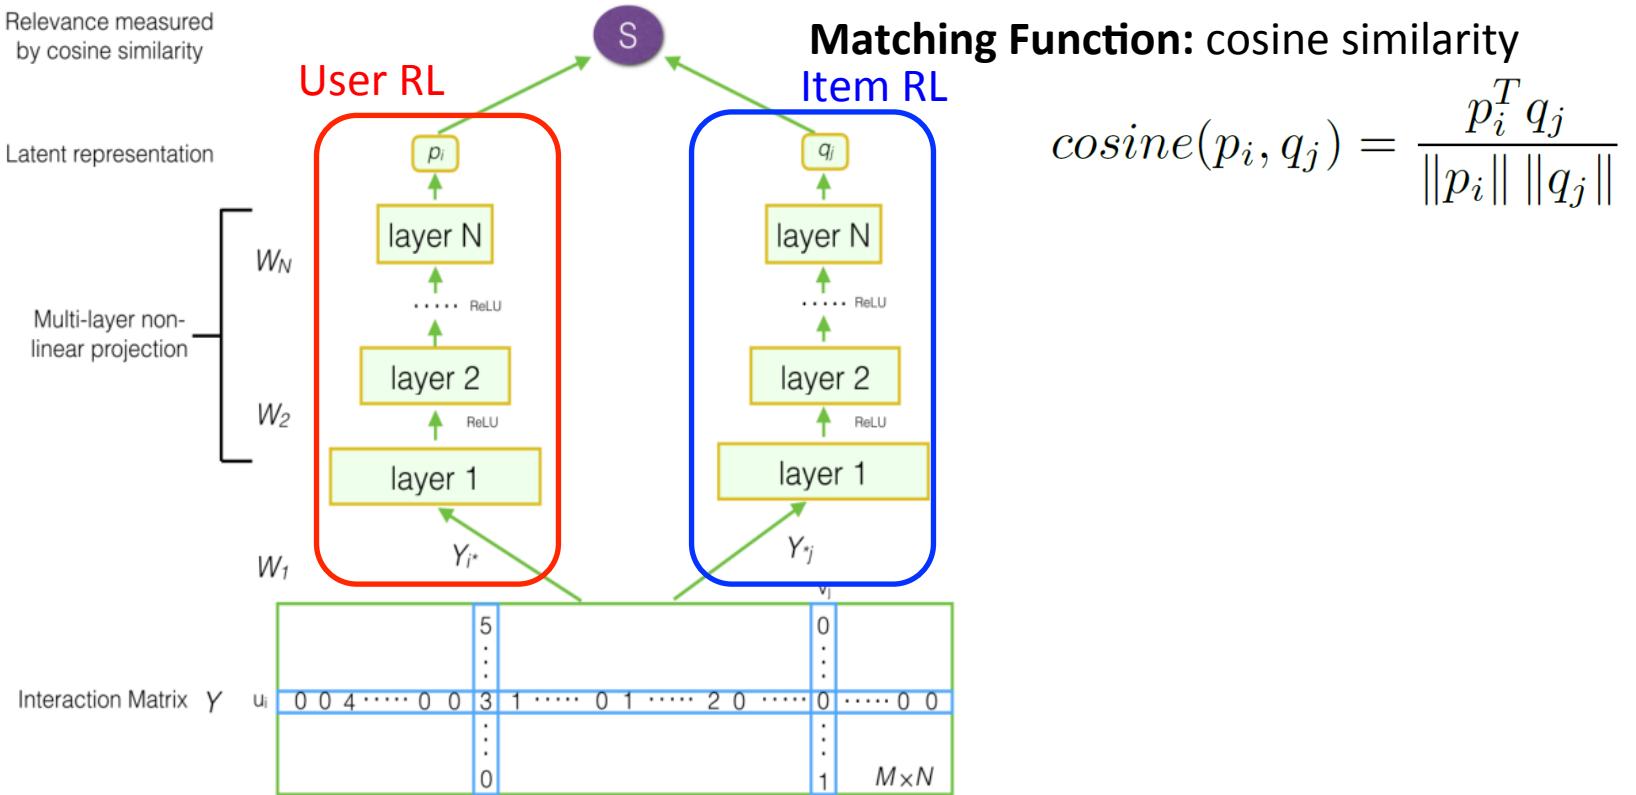
item -> **users who have rated it** (multi-hot), i.e., column vector of \mathbf{Y}

indicates the item's rating profile.

Interaction Matrix \mathbf{Y}		u_i	$0 \ 0 \ 4 \ \dots \ 0 \ 0$	$3 \ 1 \ \dots \ 0 \ 1 \ \dots \ 2 \ 0 \ \dots \ 0 \ \dots \ 0 \ 0$	0	$M \times N$
5	\vdots				\vdots	

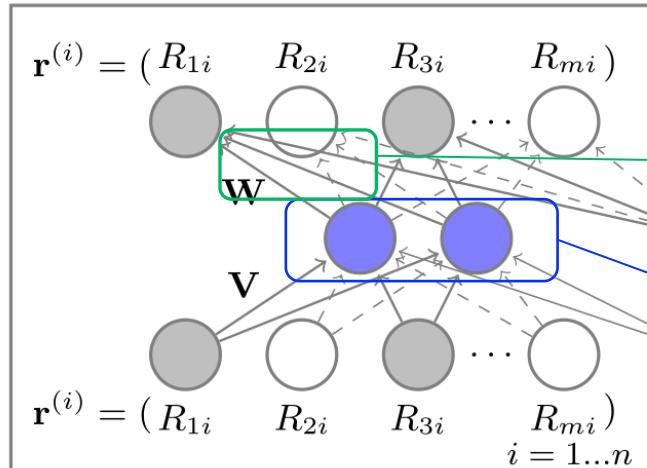
Deep Matrix Factorization (Xue et al, IJCAI'17)

- **Representation Function:**
 - Multi-Layer Perceptron



AutoRec (Sedhain et al, WWW'15)

- **Input:** (similar to DeepMF)
 - user -> historically rated items (user-based autoencoder).
 - item-> ID
- **Representation Function:** Multi-Layer Perceptron
- **Matching Function:** inner product



Input reconstruction: $h(\mathbf{r}; \theta) = f(\mathbf{W} \cdot g(\mathbf{V}\mathbf{r} + \boldsymbol{\mu}) + \mathbf{b})$

$$\min_{\theta} \sum_{i=1}^n \|\mathbf{r}^{(i)} - h(\mathbf{r}^{(i)}; \theta)\|_{\mathcal{O}}^2 + \frac{\lambda}{2} \cdot (\|\mathbf{W}\|_F^2 + \|\mathbf{V}\|_F^2),$$

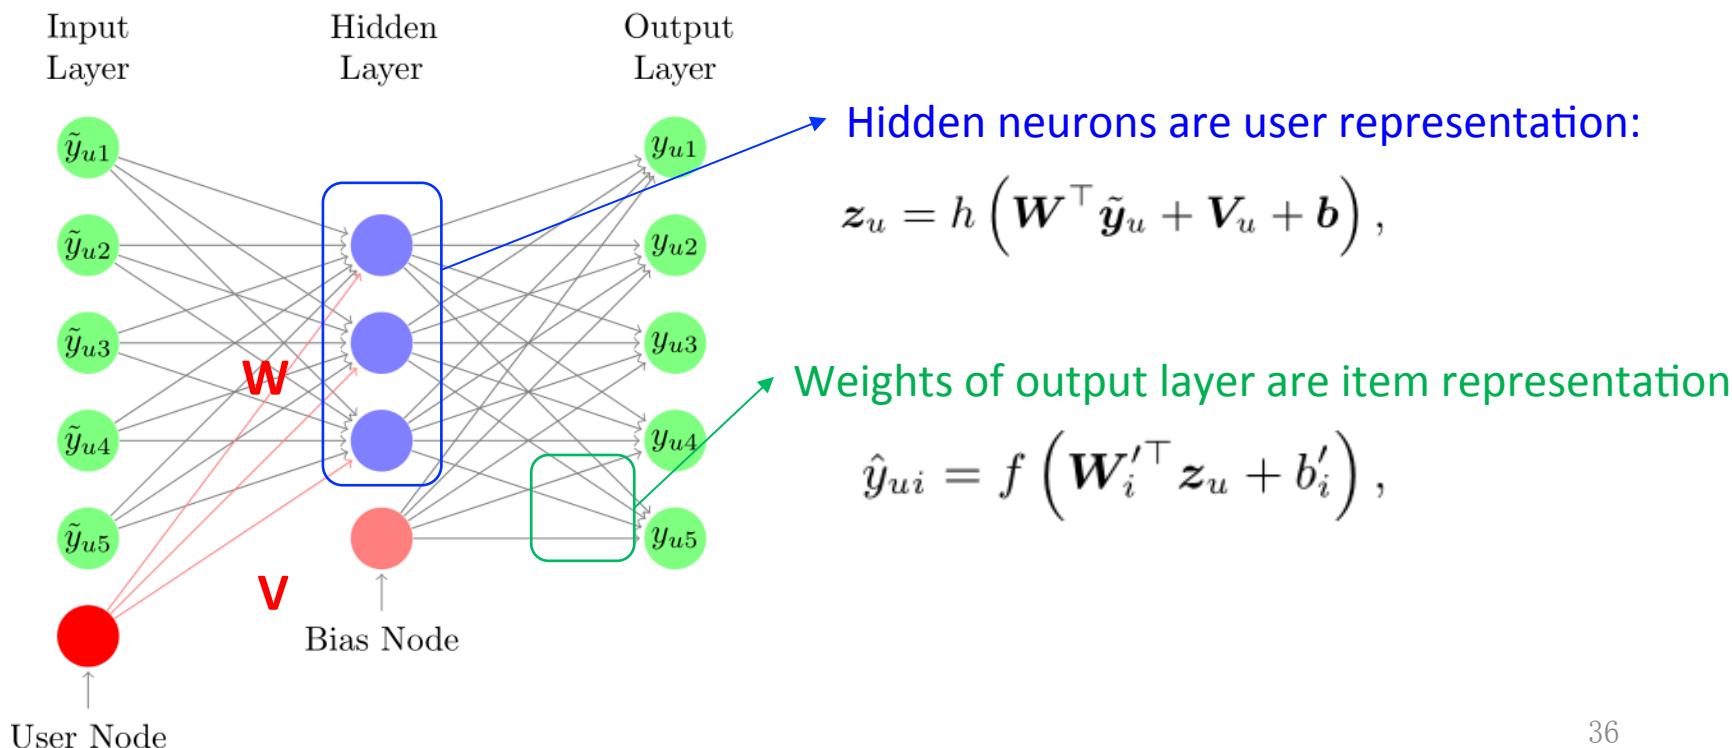
Output weights denote item representation

Hidden neurons denote user representation

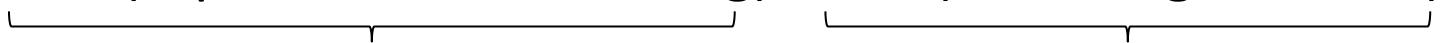
user-based autoencoder

Collaborative Denoising Auto-Encoder (Wu et al, WSDM'16)

- **Input:**
 - user \rightarrow ID & historically rated items (similar to SVD++)
 - item \rightarrow ID
- **Representation Function:** Multi-Layer Perceptron



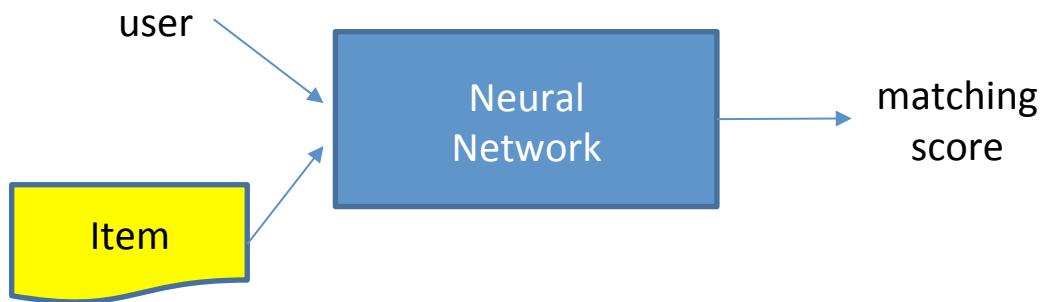
Short Summary

- Either ID or history is used as the profile of user/item
- Models with history as input are more expressive, but are also more expensive to train.
- The Auto-Encoder architecture is essentially identical to
MLP (representation learning) + MF (matching function).


The diagram illustrates the Auto-Encoder architecture. It consists of two main components: 'Nonlinear' (MLP) and 'Linear' (MF). The 'Nonlinear' component is represented by a bracket under the text 'MLP (representation learning)', and the 'Linear' component is represented by a bracket under the text 'MF (matching function)'.

Next...

- Methods of **matching function learning**:



- Based on **Neural Collaborative Filtering (NCF)** framework:
 - **NeuMF**: Neural Matrix Factorization (He et al, WWW'17)
 - **NNCF**: Neighbor-based NCF (Bai et al, CIKM'17)
 - **ConvNCF**: Convolutional NCF (He et al, IJCAI'18)

Why Using Neural Networks to Learn the Matching Function?

- The simple choice of **inner product** can limit the **expressiveness** of an embedding-based matching model.

$$\hat{y}_{ui} = \mathbf{U}_i^T \mathbf{V}_j \simeq \cos(\mathbf{U}_i, \mathbf{V}_j) \quad (\text{E.g., assuming a unit length})$$

- Example:

	i_1	i_2	i_3	i_4	i_5
u_1	1	1	1	0	1
u_2	0	1	1	0	0
u_3	0	1	1	1	0
u_4	1	0	1	1	1

$$\text{sim}(u_1, u_2) = 0.5$$

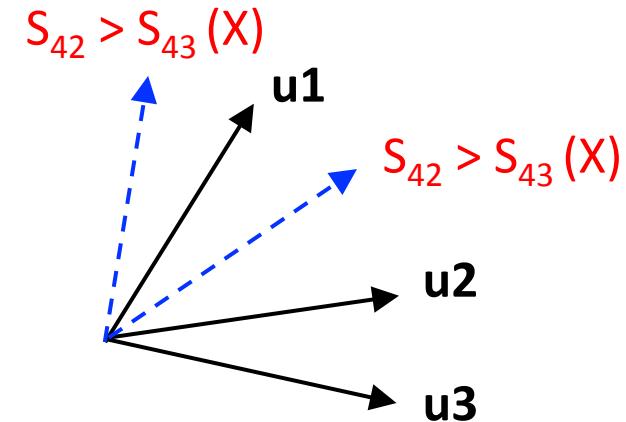
$$\text{sim}(u_3, u_1) = 0.4$$

$$\text{sim}(u_3, u_2) = 0.66$$

$$\text{sim}(u_4, u_1) = 0.6 \quad *****$$

$$\text{sim}(u_4, u_2) = 0.2 \quad *$$

$$\text{sim}(u_4, u_3) = 0.4 \quad ***$$



(He et al, WWW'17)

$$\text{Jaccard Similarity: } s_{ij} = \frac{|\mathcal{R}_i \cap \mathcal{R}_j|}{|\mathcal{R}_i \cup \mathcal{R}_j|}$$

Neural Collaborative Filtering Framework (He et al, WWW'17)

- NCF is a general framework that replaces the inner product with a neural network to learn the matching function.

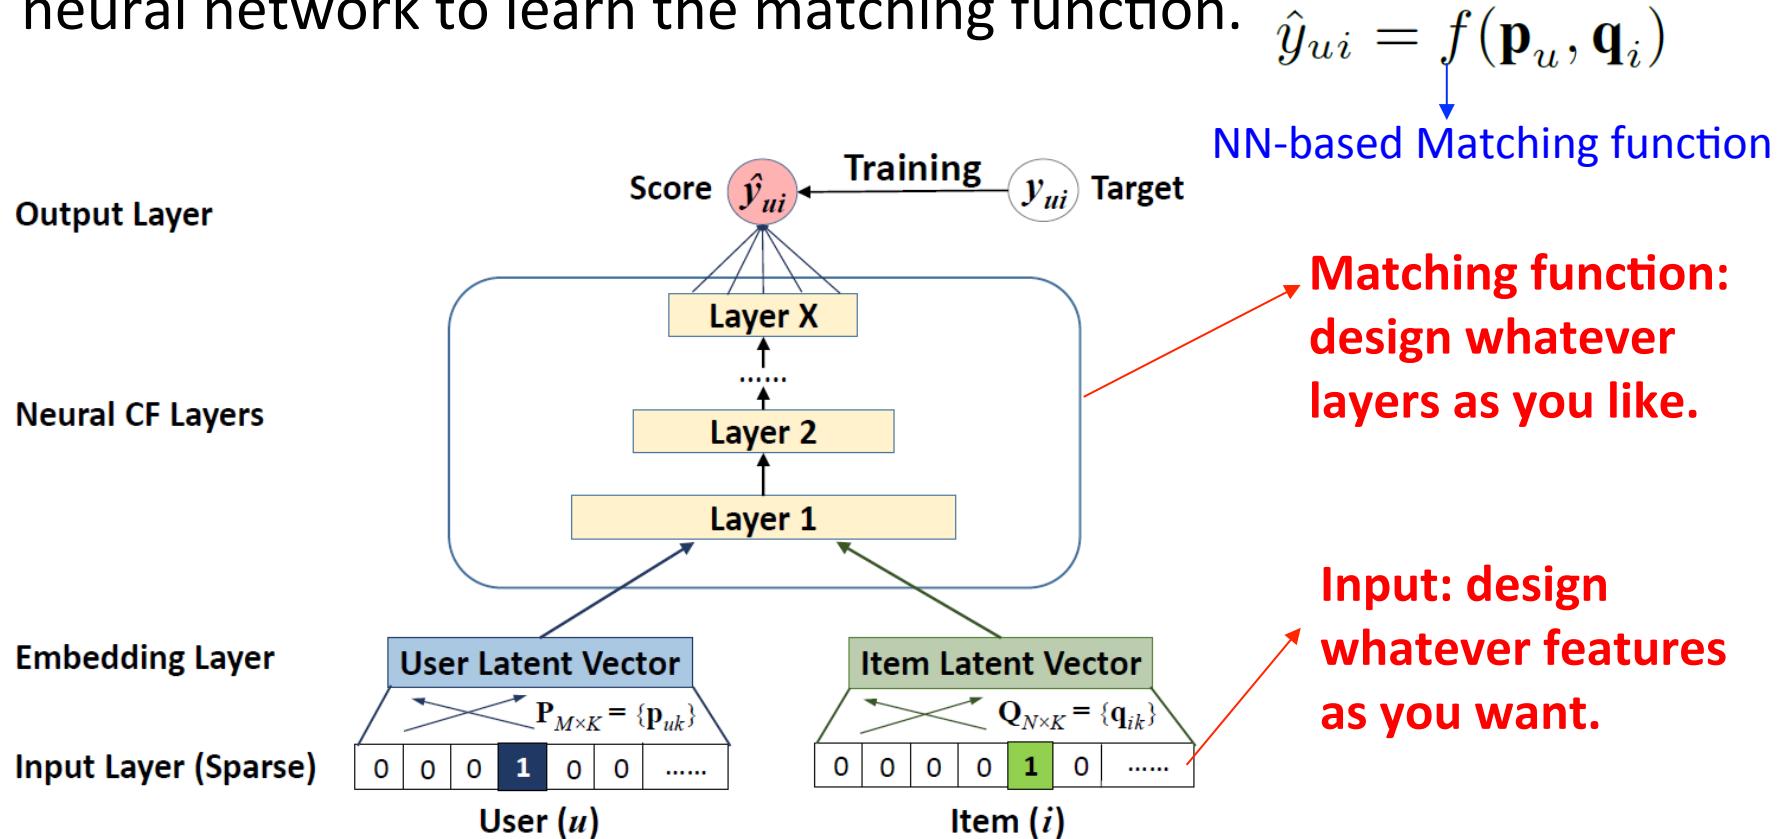
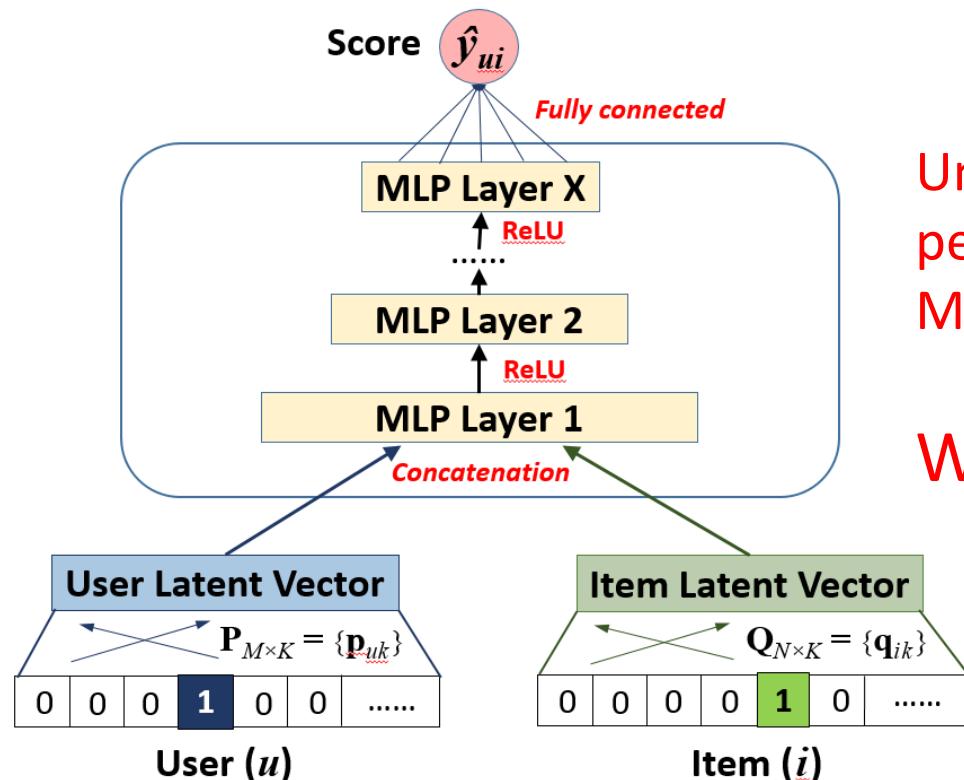


Figure 2: Neural collaborative filtering framework

Multi-Layer Perceptron for CF

- The most intuitive idea is to use a Multi-Layer Perceptron as the matching function.



Unfortunately, MLP doesn't perform well and underperforms MF.

Why?

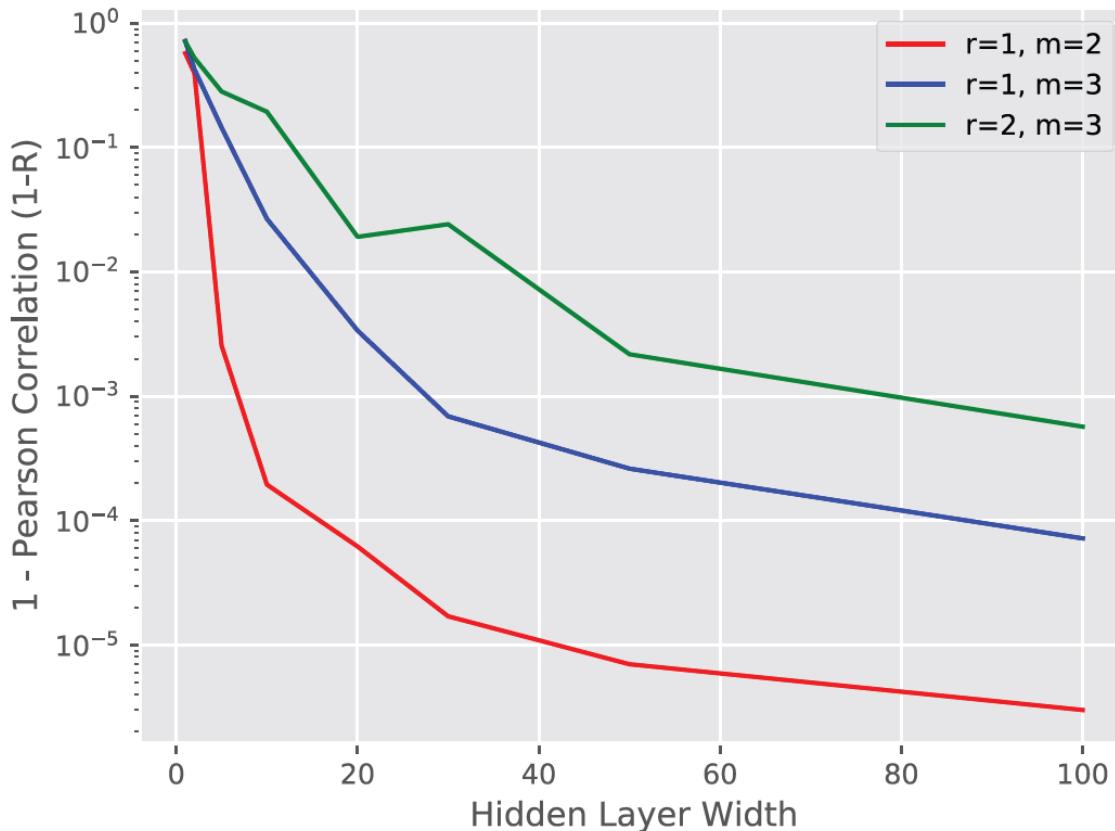
(He et al, WWW'17)

Slides: <http://comp.nus.edu.sg/~xiangnan/icmr18-recsys.pdf>

MLP is Weak in Capturing Low-Rank Relation (Beutel et al, WSDM'18)

Setting: Generating low-rank data, and using one-layer MLP to fit it.

r : rank size; m : data dimension (2 -> matrix; 3 -> 3D tensor).

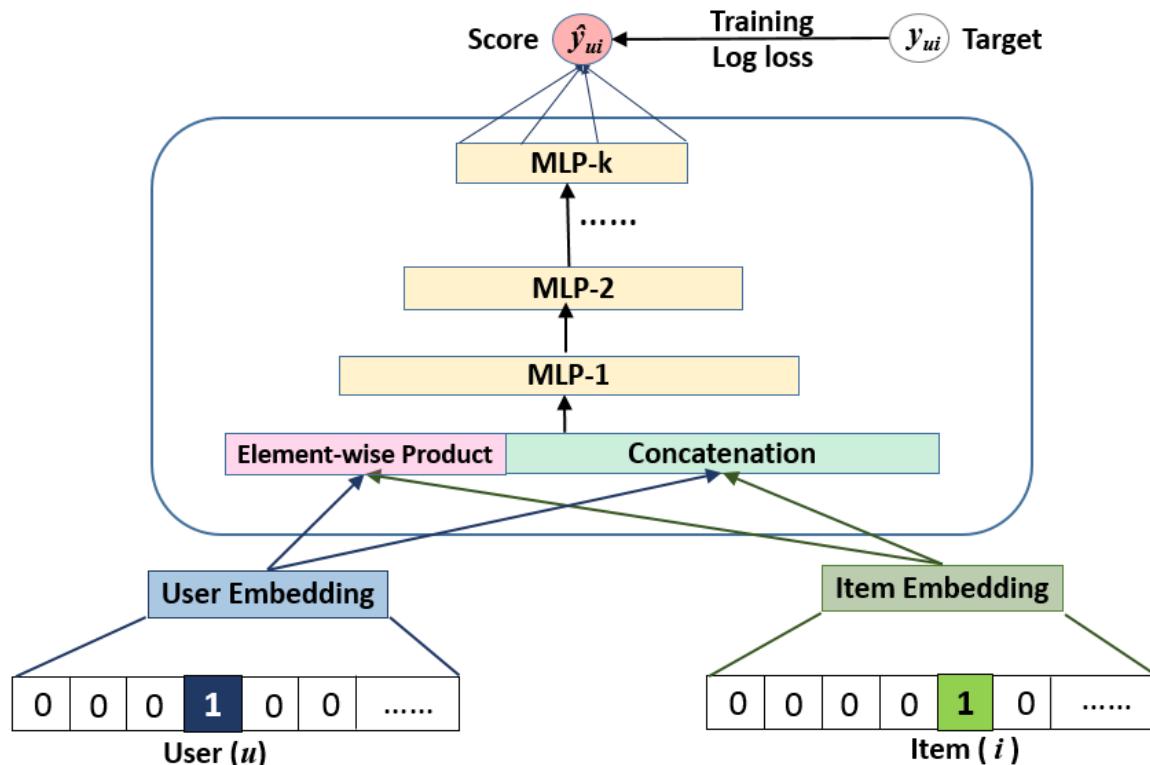


We have to design
more effective
models to make
DNN work for CF!

MLP can learn to approximate the low-rank relation, but is inefficient in doing so.

NeuMF: Neural Matrix Factorization (He et al, WWW'17)

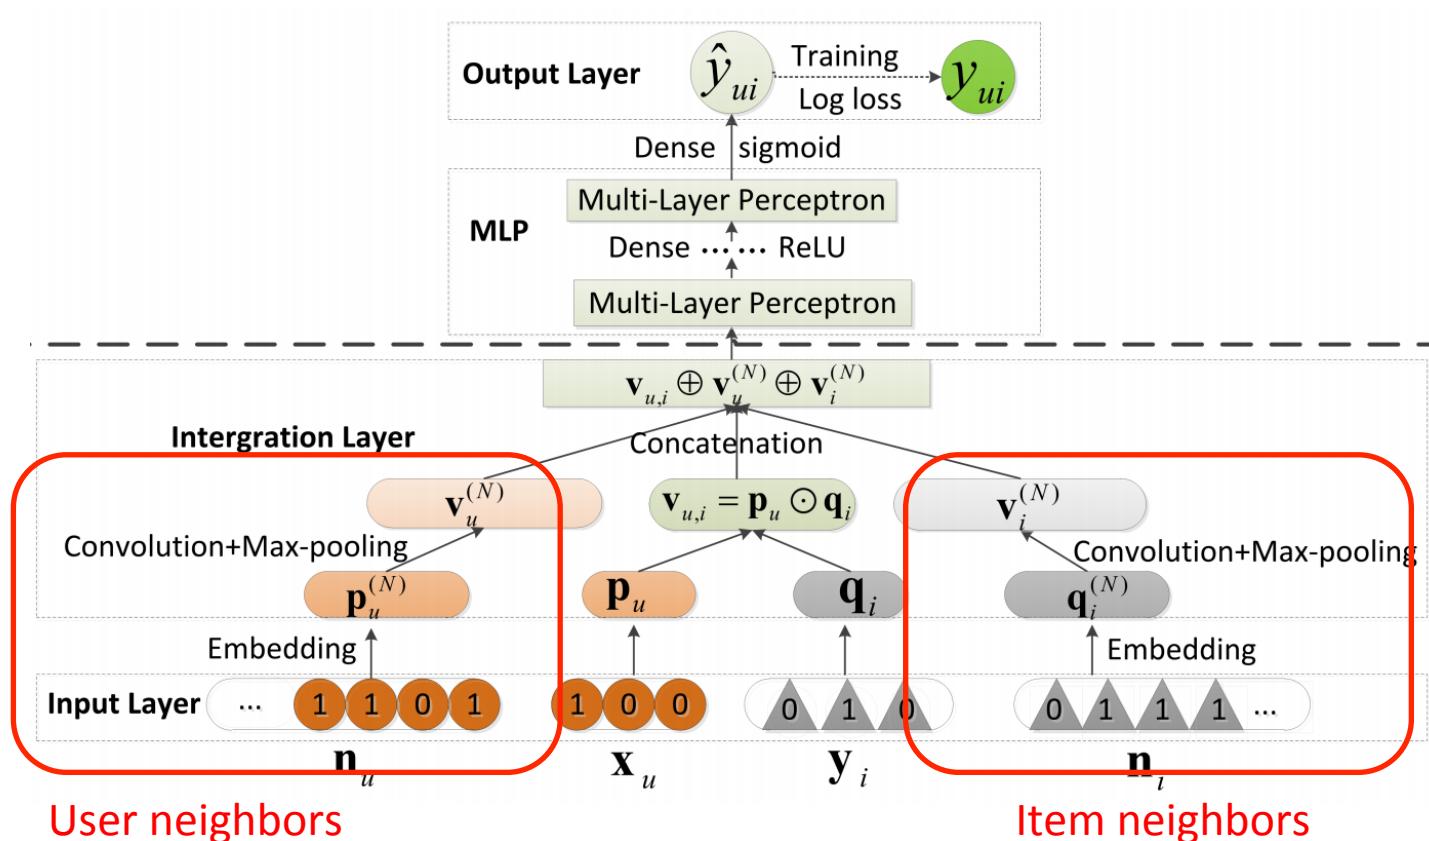
- NeuMF unifies the strengths of MF and MLP in learning the matching function:
 - MF uses inner product to capture the **low-rank** relation
 - MLP is more **flexible in using DNN** to learn the matching function.



NNCF: Neighbor-based NCF

(Bai et al, CIKM'17)

- Feeding user and item **neighbors** into the NCF framework
 - Direct neighbors or community neighbors are considered.



Experiment Evidence

Datasets	#Interaction	# Users	#Items	Sparsity
Delicious	437,593	1,867	69,223	99.66%
MovieLens	1,000,209	3,706	6,040	95.53%

Performance Comparison on Item Recommendation (%)

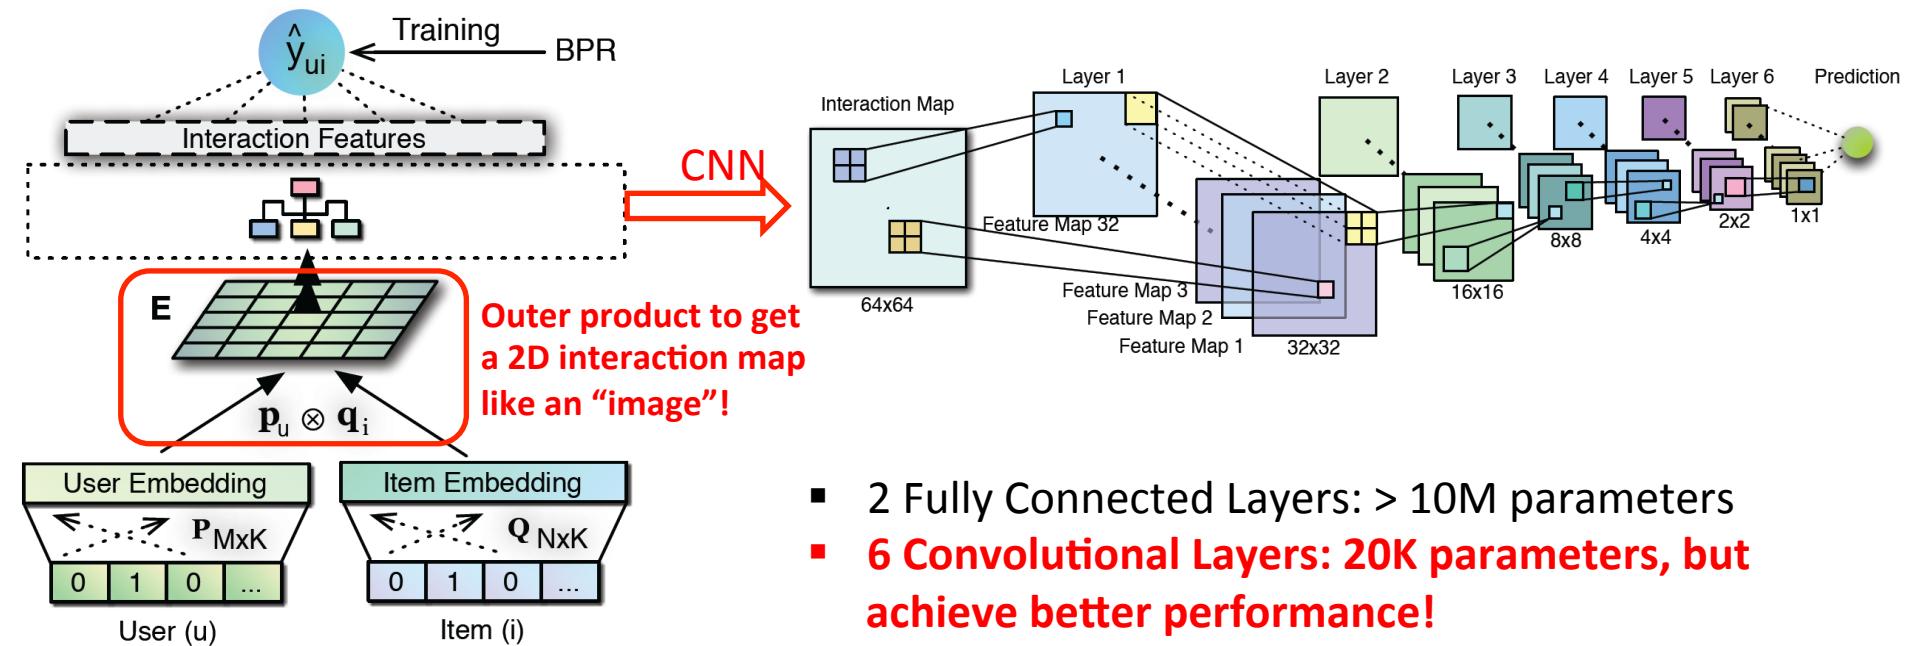
Datasets	Delicious		MovieLens		
	Models	HR@5	NDCG@5	HR@5	NDCG@5
ItemPop	5.41	3.22	31.49	20.18	CF method is better than non-personalized method
ItemKNN	59.69	55.90	45.01	30.14	Model-based CF is better than memory-based CF
MF-BPR	73.77	74.11	51.03	36.21	
NeuMF	85.53	80.68	56.55	38.30	
NNCF	87.31	84.58	62.00	42.21	Deep NCF models are better than shallow MF models by a large margin.

(Bai et al, CIKM'17)

Slides: <http://comp.nus.edu.sg/~xiangnan/icmr18-recsys.pdf>

Convolutional NCF (He et al, IJCAI'18)

- Although fully connected layers are popular in learning the matching function, they have too many parameters.
- Recently, we propose to use the **locally connected CNN** to build deeper NCF models.



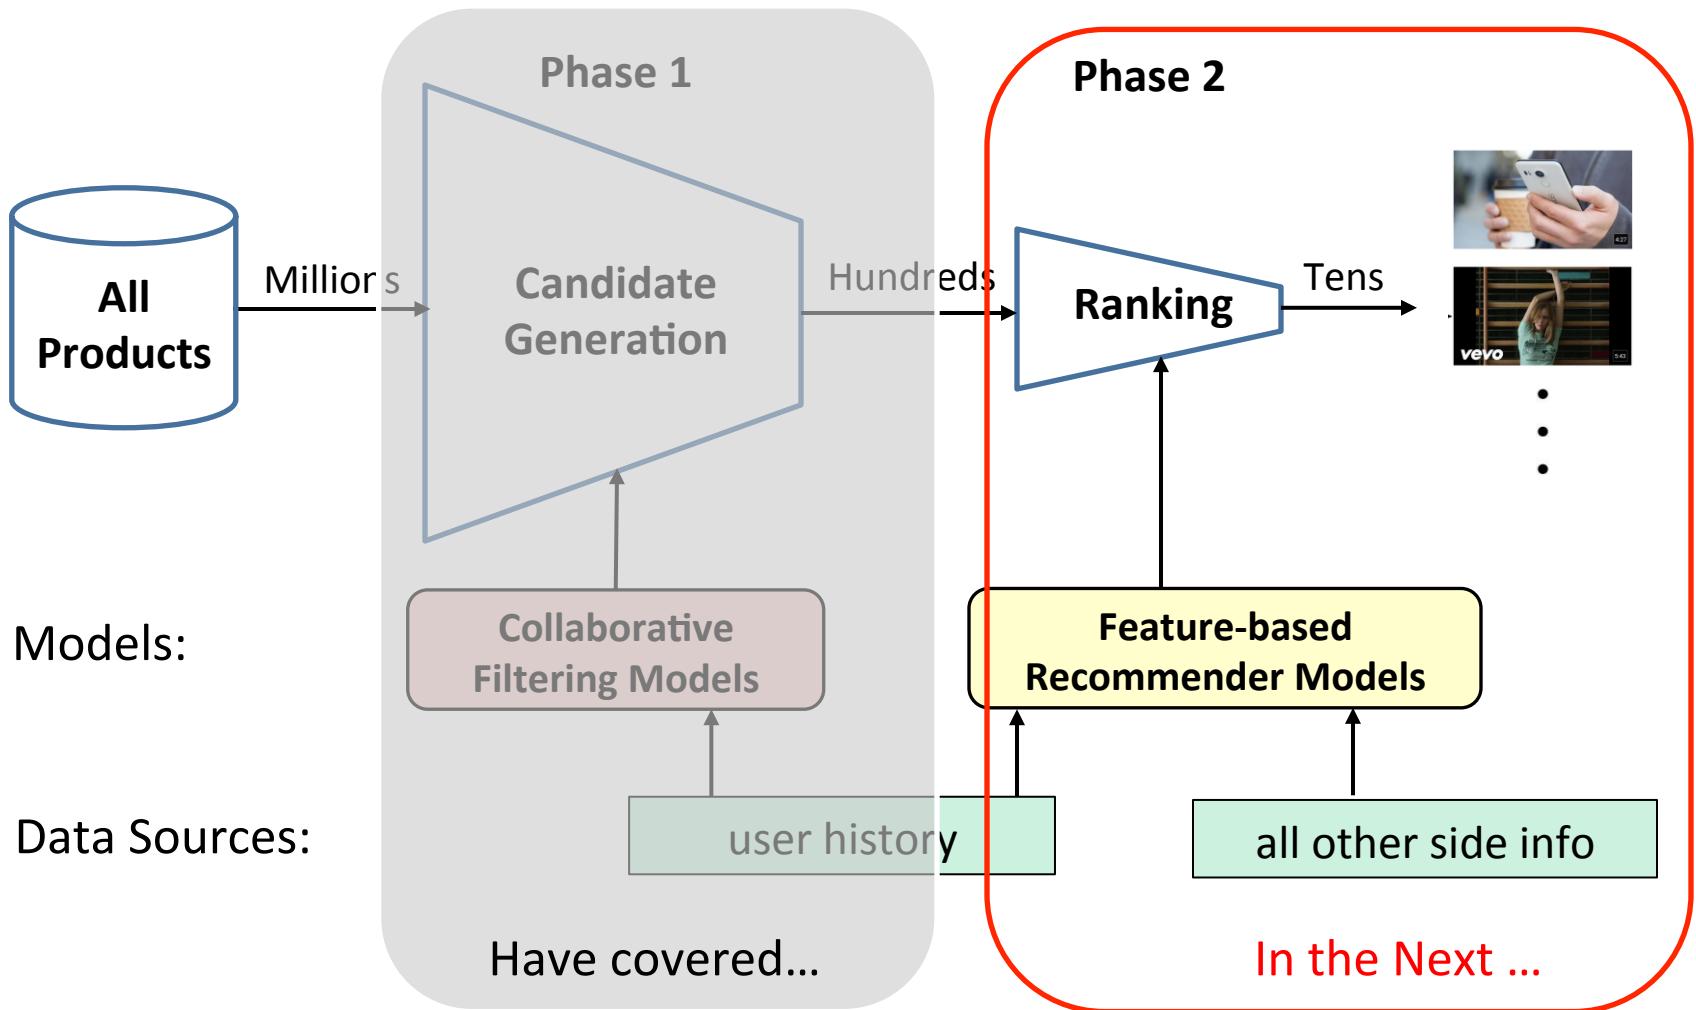
Experiment Evidence

Datasets	#Interactions	#Users	#Items	Sparsity
Yelp	730,791	25,815	25,677	99.89%
Gowalla	1,249,703	54,156	52,400	99.95%

Datasets	Gowalla		Yelp		
	Models	HR@5	NDCG@5	HR@5	NDCG@5
ItemPop	20.03	10.99	7.10	3.65	
MF-BPR	62.84	48.25	17.52	11.04	
MLP	63.59	48.02	17.66	11.03	
IRGAN	63.89	49.58	18.61	11.98	
NeuMF	67.44	53.19	18.81	11.89	
ConvNCF	69.14	54.94	19.06	12.09	

ConvNCF are better than NeuMF and MLP with much fewer parameters.

Morden RecSys Architecture



Recall: Input to Feature-based Models

Feature vector \mathbf{x}												Target y		
$\mathbf{x}^{(1)}$	1	0	0	...	1	0	0	0	...	0.3	0.3	0.3	0	...
$\mathbf{x}^{(2)}$	1	0	0	...	0	1	0	0	...	0.3	0.3	0.3	0	...
$\mathbf{x}^{(3)}$	1	0	0	...	0	0	1	0	...	0.3	0.3	0.3	0	...
$\mathbf{x}^{(4)}$	0	1	0	...	0	0	1	0	...	0	0	0.5	0.5	...
$\mathbf{x}^{(5)}$	0	1	0	...	0	0	0	1	...	0	0	0.5	0.5	...
$\mathbf{x}^{(6)}$	0	0	1	...	1	0	0	0	...	0.5	0	0.5	0	...
$\mathbf{x}^{(7)}$	0	0	1	...	0	0	1	0	...	0.5	0	0.5	0	...
	A	B	C	...	TI	NH	SW	ST	...	TI	NH	SW	ST	Other Movies rated
	User				Movie				Other Movies rated					

Raw features:

1. Categorical features
One-hot encoding on ID features
2. Continuous features
E.g., time, frequency.
Need feature normalization

Transformed features:

1. Categorical features
Cross features are important
(e.g., AND (A=true, B=true))
2. Continuous features
E.g., outputs of other models like
visual embeddings.

Importance of Cross Feature

An example that cross feature (**feature interaction**) is important for prediction:

Task: predict customer *income*.

Two input variables:

- 1) occupation = {banker, engineer,...}
- 2) Level = {junior, senior}

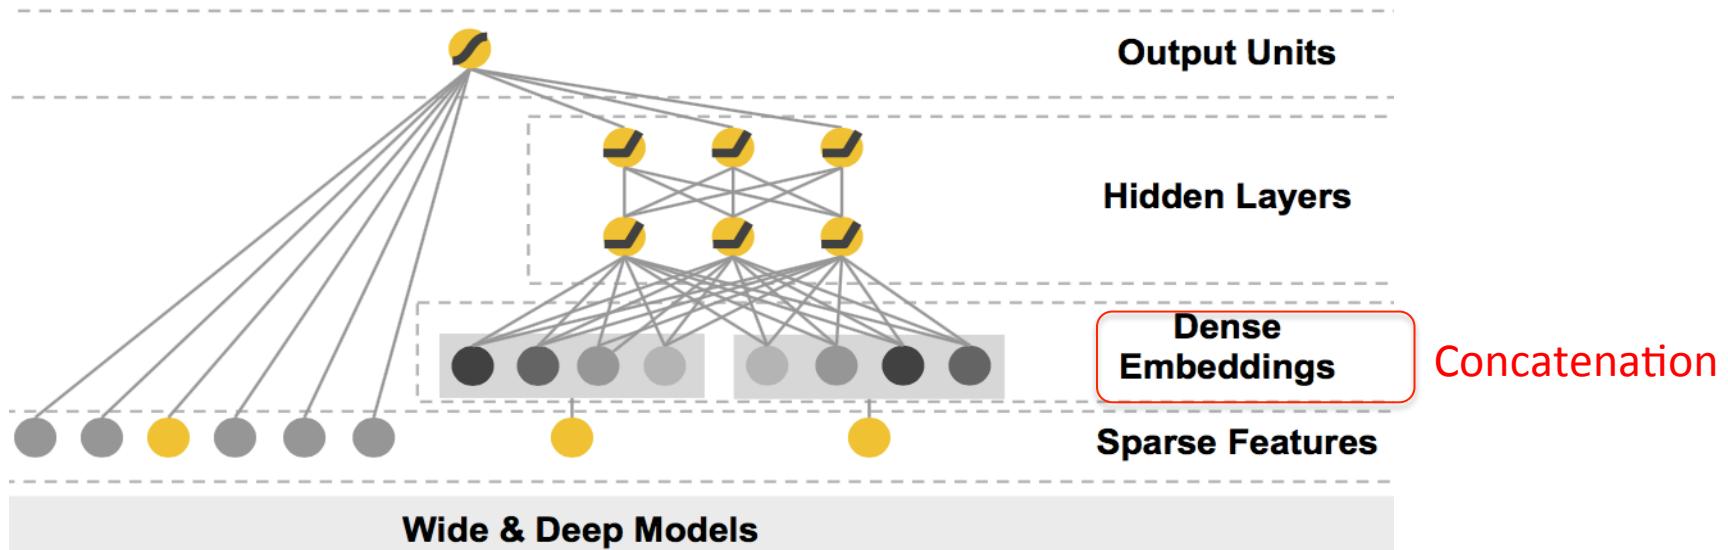
Facts:

income(junior, banker) < income(junior, engineer)

but,

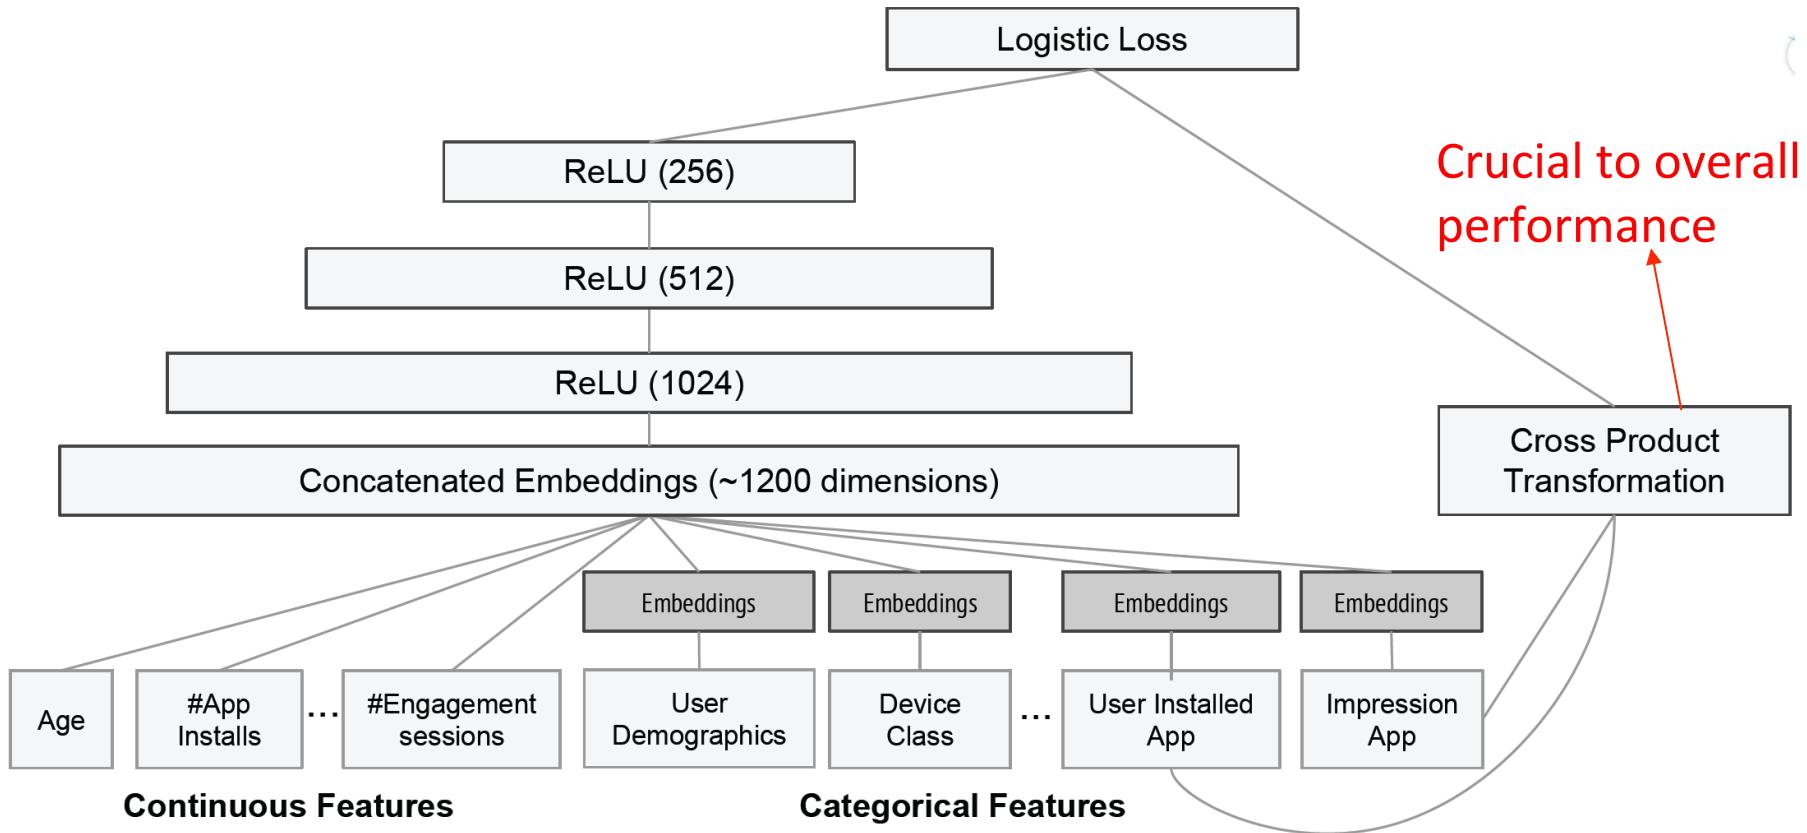
Income(senior, banker) > income(senior, engineer)

Wide&Deep (Cheng et al, Recsys'16)



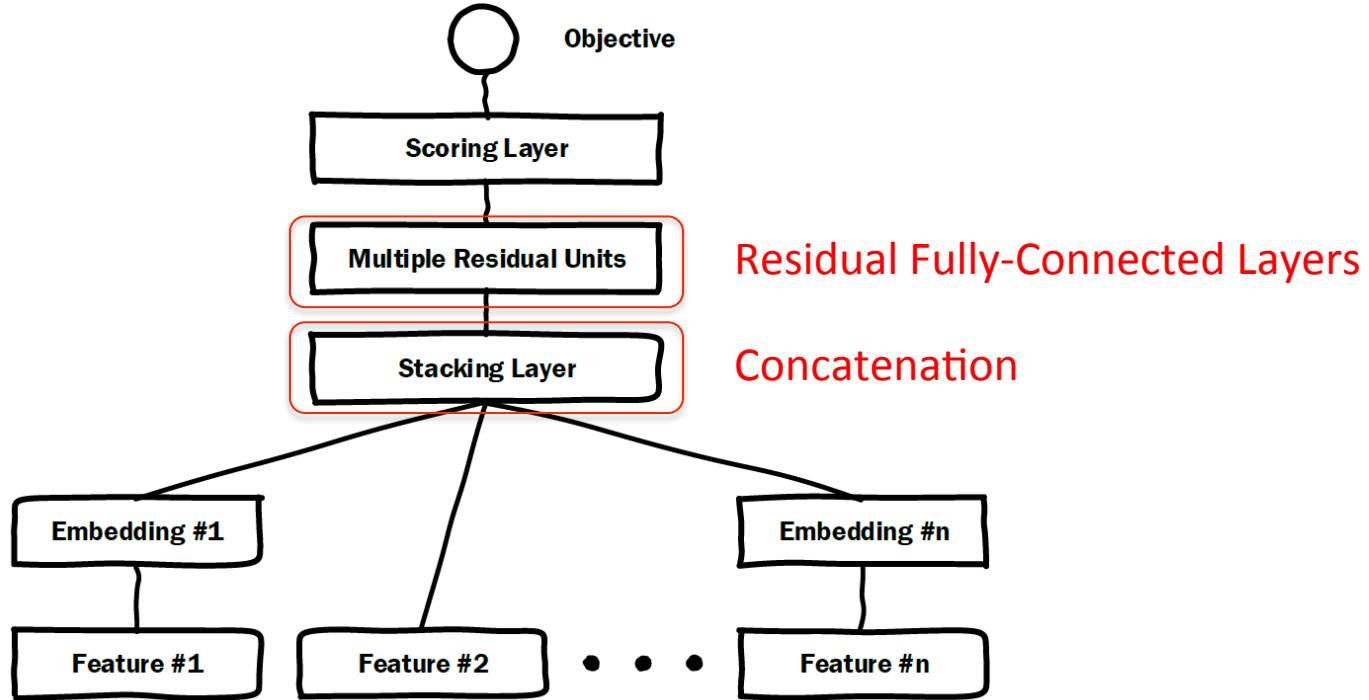
- The wide part is **linear regression** for **memorizing seen feature interactions**, which requires careful engineering on cross features.
E.g., $AND(gender=female, language=en)$ is 1 iff both single features are 1
- The deep part is **DNN** for **generalizing to unseen feature interactions**. Cross feature effects are captured in an implicit way.

Wide&Deep for Google App Recommendation (Cheng et al, Recsys'16)



Deep Crossing (Shan et al, KDD'16)

Microsoft's Sponsor Search Solution in 2016:



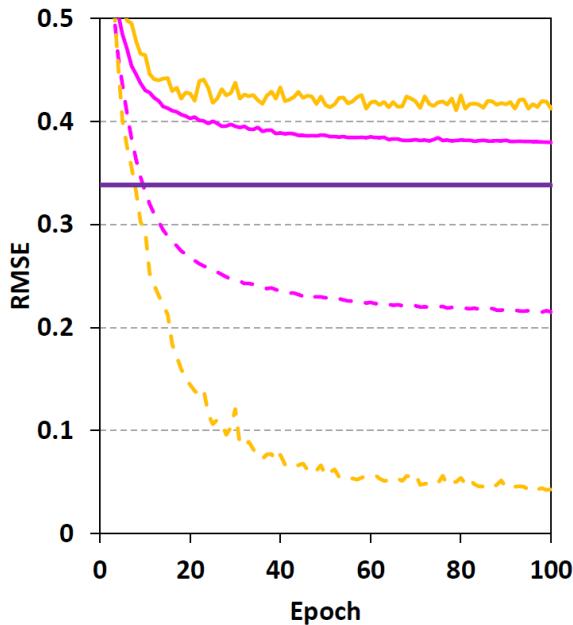
The use of residual layers makes the network be deep.

Empirical Evidence

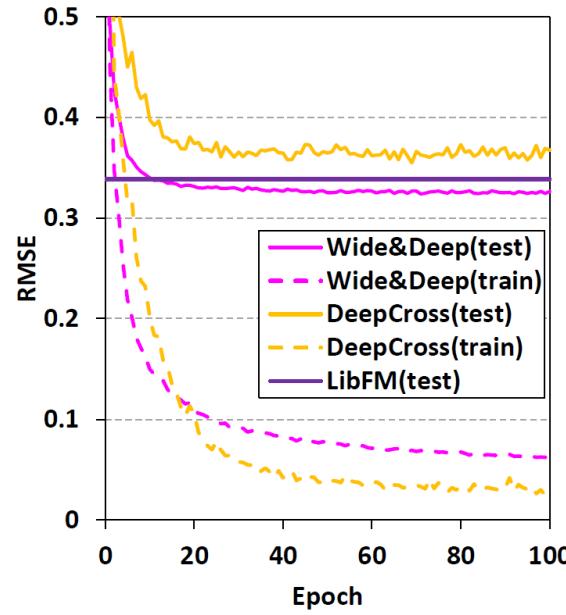
- However, when only **raw features** are used, both DL models don't perform well in learning unseen feature interactions.

Solid line: testing loss;

Dashed line: training loss



(a) Random initialization

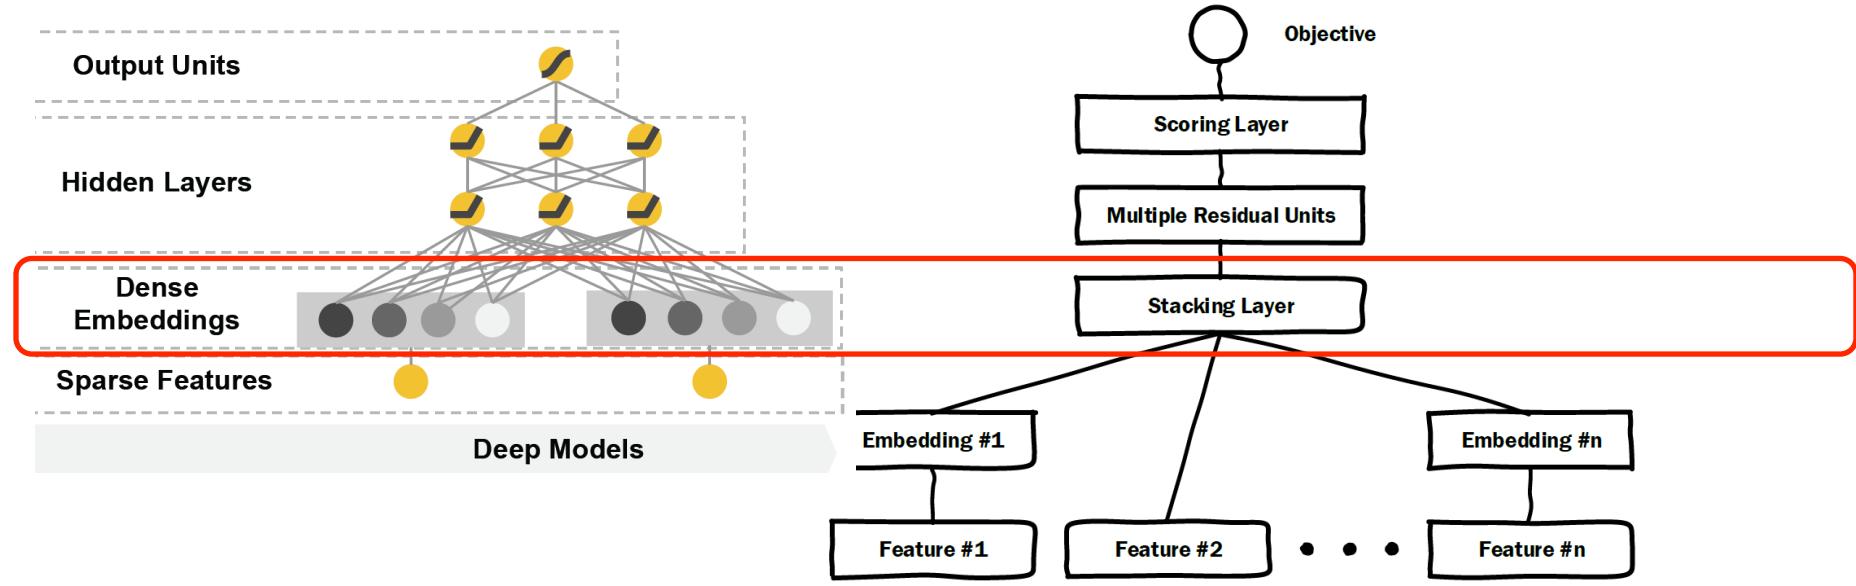


(b) FM as pre-training

(He and Chua, SIGIR'17)

Why MLP is Ineffective?

Besides optimization difficulties, one reason might be model design:



1. Embedding concatenation carries little information about feature interactions in the low level!
2. The structure of Concat+MLP is ineffective to learn the multiplicative relation (Beutel et al, WSDM'18).

NFM: Neural Factorization Machine (He and Chua, SIGIR'17)

- Inspired by FM, NFM models **pairwise interactions** between feature embeddings with **multiplication**.

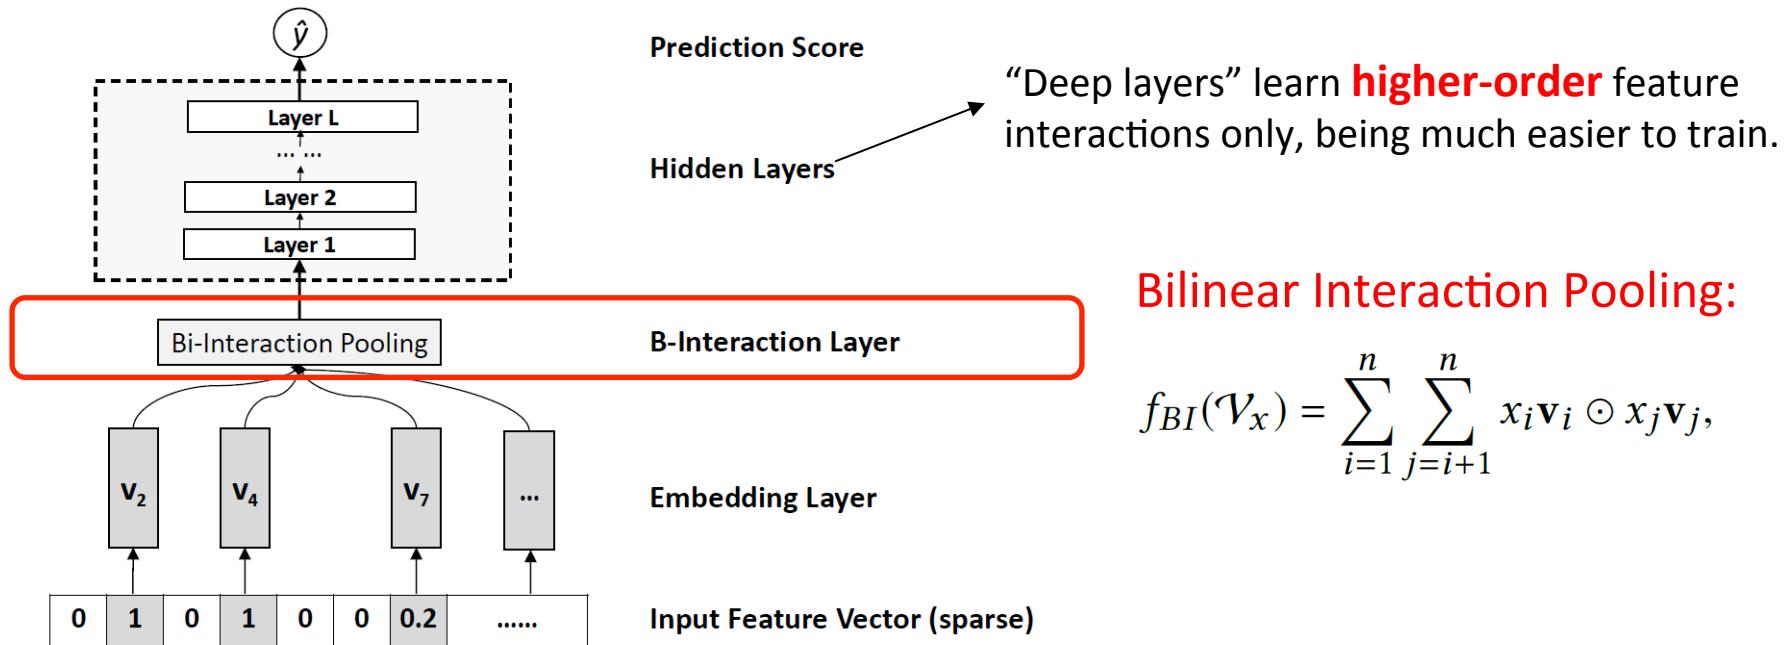


Figure 2: Neural Factorization Machines model (the first-order linear regression part is not shown for clarity).

Experiment Evidence

Task #1: Context-aware App Usage Prediction

- Frappe data: instance #: 288,609, feature #: 5,382

Task #2: Personalized Tag Recommendation

- MovieLens data: Inst #: 2,006,859, Feat #: 90,445

Table: Parameter # and testing RMSE at embedding size 128

	Frappe		MovieLens	
Method	Param#	RMSE	Param#	RMSE
Logistic Regression	5.38K	0.5835	0.09M	0.5991
FM	1.38M	0.3385	23.24M	0.4735
High-order FM	2.76M	0.3331	46.40M	0.4636
Wide&Deep (3 layers)	4.66M	0.3246	24.69M	0.4512
DeepCross (10 layers)	8.93M	0.3548	25.42M	0.5130
Neural FM (1 layer)	1.45M	0.3095	23.31M	0.4443

Codes: github.com/hexiangnan/neural_factorization_machine

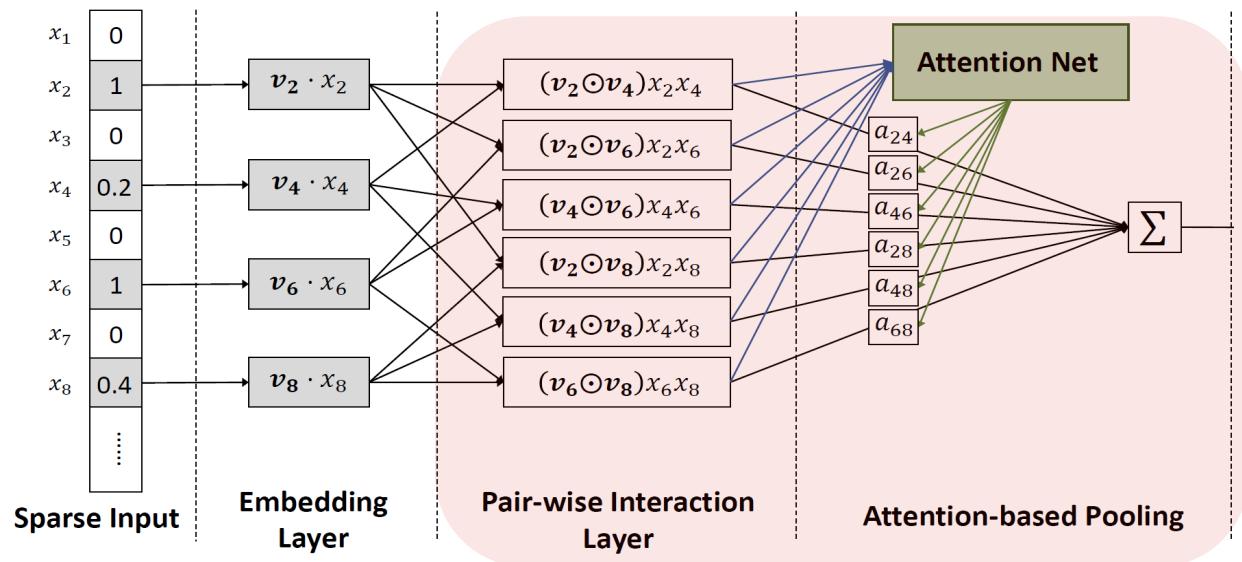
1. Shallow embedding methods learn interactions, better than simple linear models

2. Deep embedding methods:
Wide&Deep = Concat+3 layers
DeepCross = Concat+10 layers

3. Our methods:
Neural FM = BI pooling + 1 layer
Shallower but outperforming existing deeper methods with less parameters.

AFM: Attentional Factorization Machine (Xiao et al, IJCAI'17)

- Neural FM treats all second-order feature interactions as **contributing equally**.
- Attentional FM uses an **attention network** to learn the **weight** of a feature interaction.



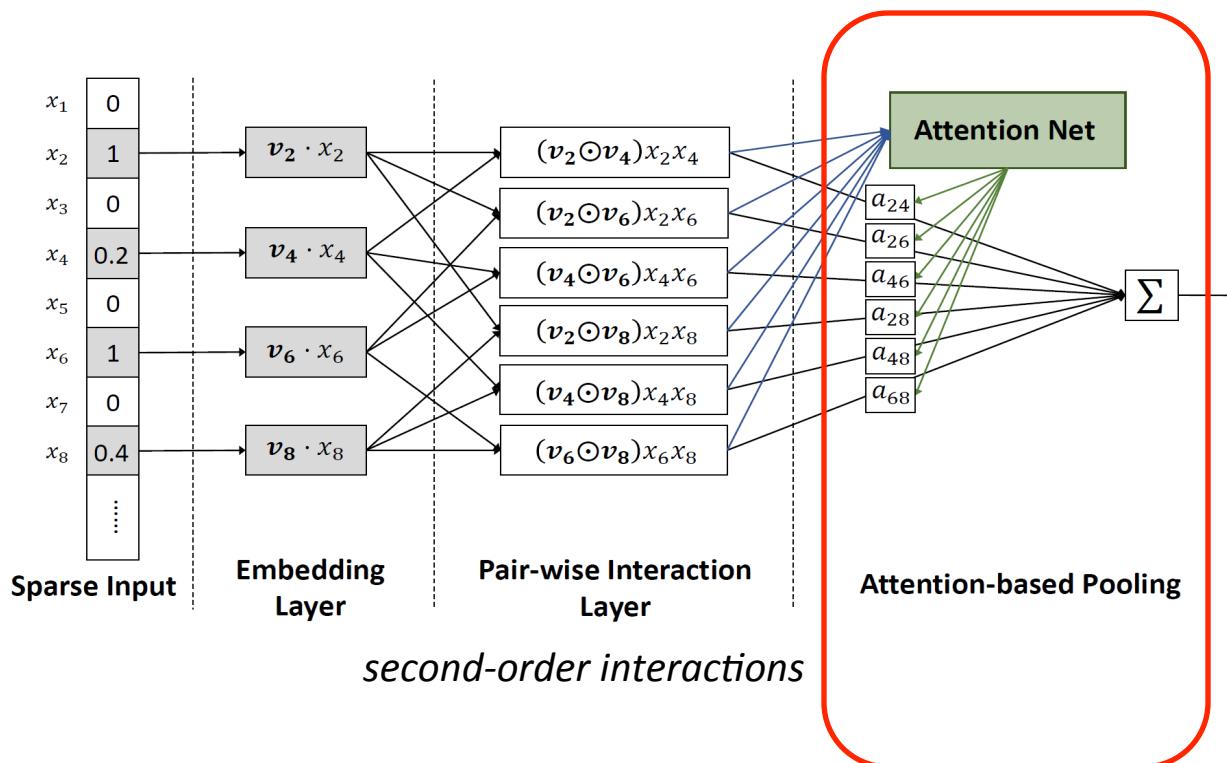
$$f_{ABI}(\mathcal{V}_x) = \sum_{i=1}^n \sum_{j=i+1}^n (x_i \mathbf{v}_i \odot x_j \mathbf{v}_j) a_{ij}$$

$$a'_{ij} = \mathbf{h}^T \text{ReLU}(\mathbf{W}(\mathbf{v}_i \odot \mathbf{v}_j)x_i x_j + \mathbf{b}),$$

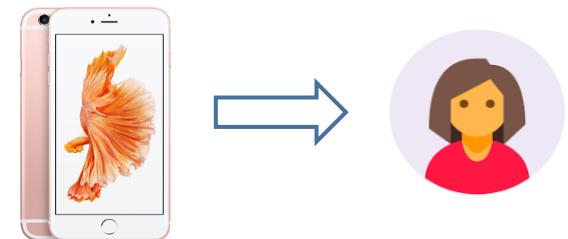
$$a_{ij} = \frac{\exp(a'_{ij})}{\sum_{(i,j) \in \mathcal{R}_x} \exp(a'_{ij})},$$

Explaining Recommendation with AFM

The attention scores can be used to select the most predictive second-order feature interactions as explanations.



Example: **explainable recommendation** with second-order cross features:



<Female, Age 20>
<Age 20, iPhone>
<Female, Color Pink>
.....

Experiment Evidence

Task #1: Context-aware App Usage Prediction

- Frappe data: instance #: 288,609, feature #: 5,382

Task #2: Personalized Tag Recommendation

- MovieLens data: Inst #: 2,006,859, Feat #: 90,445

Table: Parameter # and testing RMSE at embedding size 128

Method	Frappe		MovieLens	
	Param#	RMSE	Param#	RMSE
Logistic Regression	5.38K	0.5835	0.09M	0.5991
FM	1.38M	0.3385	23.24M	0.4735
High-order FM	2.76M	0.3331	46.40M	0.4636
Wide&Deep (3 layers)	4.66M	0.3246	24.69M	0.4512
DeepCross (10 layers)	8.93M	0.3548	25.42M	0.5130
Neural FM (1 layer)	1.45M	0.3095	23.31M	0.4443
Attentional FM (0 layer)	1.45M	0.3102	23.26M	0.4325

AFM without hidden
layers can be better than
NFM with 1 hidden
layer.

Codes: github.com/hexiangnan/attentional_factorization_machine

Short Summary

- ✓ Deep methods for collaborative filtering
 - User/Item representation learning
 - Matching function learning
- ✓ Deep methods for feature-based recommendation
 - Cross features are important to encode
 - DNN can effectively learn feature interactions, but require careful design
- It remains challenging to do **explainable** learning on **high-order** feature interaction.

Reference

- Xiangnan He, Lizi Liao, Hanwang Zhang, Liqiang Nie, Xia Hu, and Tat-Seng Chua. Neural collaborative filtering. In WWW 2017.
- Xiangnan He, Xiaoyu Du, Xiang Wang, Feng Tian, Jinhui Tang, and Tat-Seng Chua. Out Product-based Neural Collaborative Filtering. In IJCAI 2018.
- Xiangnan He, and Tat-Seng Chua. Neural factorization machines for sparse predictive analytics. In SIGIR 2017.
- Jun Xiao, Hao Ye, Xiangnan He, Hanwang Zhang, Fei Wu, and Tat-Seng Chua. Attentional factorization machines: Learning the weight of feature interactions via attention networks. In IJCAI 2017.
- Ting Bai, Ji-Rong Wen, Jun Zhang, and Wayne Xin Zhao. A Neural Collaborative Filtering Model with Interaction-based Neighborhood. In CIKM 2017.
- Paul Covington, Jay Adams, and Emre Sargin. Deep neural networks for youtube recommendations. In Recsys 2016.
- Xiang Wang, Xiangnan He, Liqiang Nie, and Tat-Seng Chua. Item silk road: Recommending items from information domains to social users. In SIGIR 2017.
- Hong-Jian Xue, Xin-Yu Dai, Jianbing Zhang, Shujian Huang, and Jiajun Chen. Deep matrix factorization models for recommender systems. IJCAI 2017.
- Suvash Sedhain, Aditya Krishna Menon, Scott Sanner, and Lexing Xie. Autorec: Autoencoders meet collaborative filtering. In WWW 2015.
- Yao Wu, Christopher DuBois, Alice X. Zheng, and Martin Ester. Collaborative denoising auto-encoders for top-n recommender systems. In WSDM 2016.
- Heng-Tze Cheng, Levent Koc, Jeremiah Harmsen, Tal Shaked, Tushar Chandra, Hrishi Aradhye, Glen Anderson et al. Wide & deep learning for recommender systems. In DLRS 2016.
- Ying Shan, T. Ryan Hoens, Jian Jiao, Haijing Wang, Dong Yu, and J. C. Mao. Deep crossing: Web-scale modeling without manually crafted combinatorial features. In KDD 2016.
- Alex Beutel, Paul Covington, Sagar Jain, Can Xu, Jia Li, Vince Gatto, and Ed H. Chi. 2018. Latent Cross: Making Use of Context in Recurrent Recommender Systems. In WSDM 2018.

Outline of Tutorial

- Background (Xiangnan, 10 mins)
- Basics & Advances in Recommendation (Xiangnan, 50 mins)
- Visually-aware Product Recommendation (Xiangnan, 30 mins)
- Break (15 mins)
- Visual Representation (Hanwang, 45 mins)
- Image/Video Recommendation (Hanwang, 25 mins)
- Conclusion (Hanwang, 5 mins)

Problem Formulation

I have purchased:

U.S. Polo Assn. Men's V-Neck T-Shirt
★★★★★ 781
\$9.25 - \$40.95

Nautica Men's Short Sleeve Crew Neck Cotton T-Shirt
★★★★★ 12
\$13.30 - \$39.10

Nautica Men's Solid Crew Neck Short Sleeve Pocket T-Shirt
★★★★★ 34
\$18.17 - \$65.57

Will I like this one?



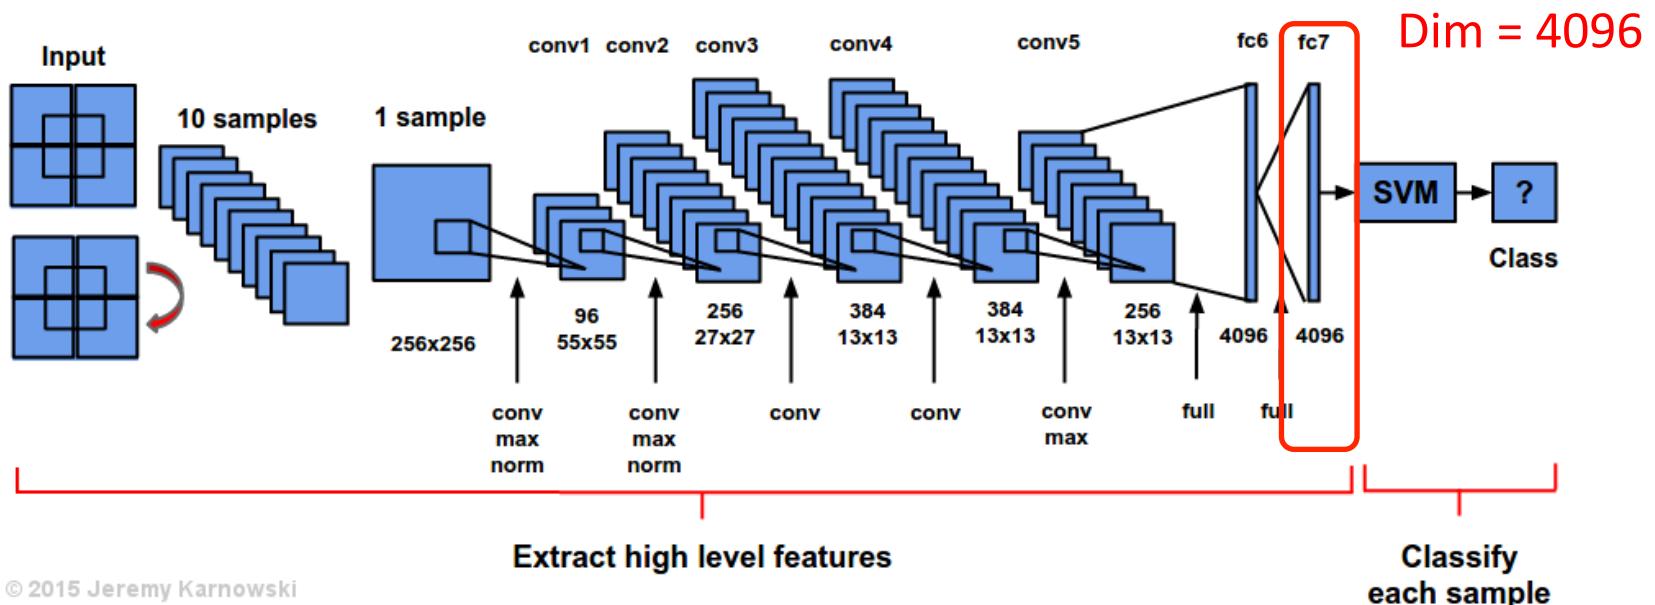
Each item = ID + image

Two key questions:

1. How to understand image?
2. How to integrate image feature into CF model?

Image Understanding

- Traditional (low-level) image features:
 - Pixels, Color histograms
 - SIFT descriptors
 - Speeded Up Robust Features (SURF)
- Gap between low-level features and real semantics.
- Recent work uses deep CNN as feature extractor.



Collaborative Filtering with CNN Features

- Let \mathbf{f}_i be CNN features for image i :
 - Usually of thousands dimension: AlexNet: 4096, ResNet: 2048

- MF predicts user rating on image i :

$$\hat{y}_{ui} = \langle \mathbf{p}_u, \mathbf{f}_i \rangle = \mathbf{p}_u^T \mathbf{f}_i$$

User preference on image CNN features

- Problem:
 - \mathbf{p}_u has to be of the **same dimension** as \mathbf{f}_i
 - Too big for CF latent space: **too many parameters => overfitting**
E.g., 100 million users * 4096 * 8 B = 3.28 TB
 - Typically, the dimension of CF latent space is hundreds (128, 256) at most.

Deep CNN Features => CF Latent Space

- An intuitive solution is to do dimension reduction on CNN features, e.g., PCA
- However, it will lose signal in CNN features.
 - The objective of dimension reduction is not recommendation.
- Solution: learning a **transformation matrix** to do the projection based on user-item interactions:

$$\hat{y}_{ui} = \mathbf{p}_u^T (\mathbf{E} \mathbf{f}_i)$$

Transformation matrix that projects
CNN features to CF latent space

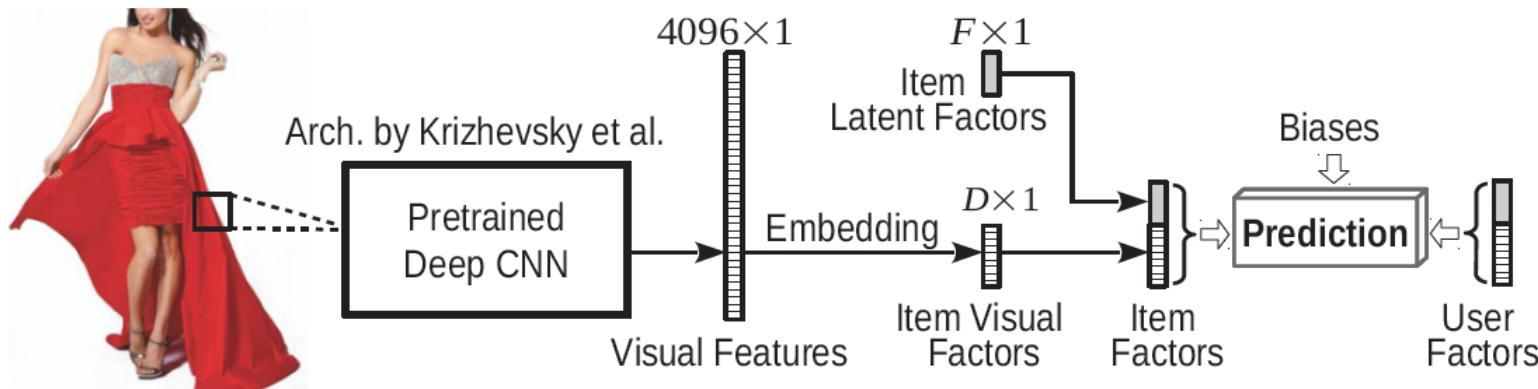
- \mathbf{E} is optimized for the recommendation task.

VBPR: Visual Bayesian Personalized Ranking (He et al, AAA'16)

- The first visually-aware recommendation method based on Deep CNN features:

$$\hat{y}_{ui} = \boxed{b_u + b_i} + \boxed{\mathbf{v}_u^T \mathbf{v}_i} + \boxed{\mathbf{p}_u^T (\mathbf{E} \mathbf{f}_i)}$$

Bias terms CF prediction Image feature-based Prediction



VBPR: Visual Bayesian Personalized Ranking (He et al, AAA'16)

$$\hat{y}_{ui} = b_u + b_i + \mathbf{v}_u^T \mathbf{v}_i + \mathbf{p}_u^T (\mathbf{E} \mathbf{f}_i)$$

- To learn model parameters, VBPR optimizes BPR pairwise loss:

$$L_{BPR} = \arg \max_{\Theta} \sum_{(u, i, j) \in \mathcal{R}_B} \ln \sigma(\hat{y}_{ui} - \hat{y}_{uj}) - \lambda \|\Theta\|^2$$

Positive prediction Negative prediction

Pairwise training examples: u prefers i over j

sigmoid

- Pairwise loss on the model:

$$\begin{aligned} \hat{y}_{ui} - \hat{y}_{uj} &= b_u + b_i + \mathbf{v}_u^T \mathbf{v}_i + \mathbf{p}_u^T (\mathbf{E} \mathbf{f}_i) \\ &\quad - [b_u + b_j + \mathbf{v}_u^T \mathbf{v}_j + \mathbf{p}_u^T (\mathbf{E} \mathbf{f}_j)] \end{aligned}$$

$$= (b_i - b_j) + \mathbf{v}_u^T (\mathbf{v}_i - \mathbf{v}_j) + \mathbf{p}_u^T \mathbf{E} (\mathbf{f}_i - \mathbf{f}_j)$$

Margin on
item bias

Margin on item
embeddings

Margin on projected
image features

Experimental Results

AUC score on personalized ranking

Dataset	Setting	(a) RAND	(b) MP	(e) BPR-MF	(f) VBPR	improvement f vs. best	improvement f vs. e
<i>Amazon Women</i>	All Items	0.4997	0.5772	0.7020	0.7834	9.4%	11.6%
	<i>Cold Start</i>	0.5031	0.3159	0.5281	0.6813	2.1%	29.0%
<i>Amazon Men</i>	All Items	0.4992	0.5726	0.7100	0.7841	9.1%	10.4%
	<i>Cold Start</i>	0.4986	0.3214	0.5512	0.6898	1.6%	25.1%
<i>Amazon Phones</i>	All Items	0.5063	0.7163	0.7918	0.8052	1.2%	1.7%
	<i>Cold Start</i>	0.5014	0.3393	0.5346	0.6056	-4.2%	13.3%
<i>Tradesy.com</i>	All Items	0.5003	0.5085	0.6198	0.7829	26.3%	26.3%
	<i>Cold Start</i>	0.4972	0.3721	0.5241	0.7594	44.9%	44.9%

1. RAND < MP (popularity) < BPR-MF < VBPR
2. VBPR has more improvements for **cold-start** items
3. Visual signal are **more important for Clothing** than Phones.

Visualization on Transformed CNN Features

$$\hat{y}_{ui} = \mathbf{p}_u^T (\mathbf{E}\mathbf{f}_i)$$

2D visualization (with t-SNE) on Amazon women dataset.



Items of the **same subcategory** are mapped to **nearby locations**.

DeepStyle (Liu et al, SIGIR'17)

- Drawback of VBPR: the transformed features mainly encode the **category** info, rather than the **style** info.
 - Example categories: ups, dresses, shoes, watches ...
 - Example styles: casual, athletic, formal ...

Each row is a cluster based on learned features.



Each cluster mixes items of different styles

DeepStyle (Liu et al, SIGIR'17)

- Item representation should be two-facet:
 $item = \text{category} + \text{style}$
- User rating on a fashion product is mainly determined by its **style**, rather than **category**.
- But how do we get style representation for an image?
 - A dedicated style-CNN needs labeled data.
- Solution: learn it from user-item interaction data!

DeepStyle (Liu et al, SIGIR'17)

Representation: $item = \text{style} + \text{category}$

Embedding:

$$\mathbf{Ef}_i = \mathbf{s}_i + \mathbf{c}_i$$

Transformed Style Category
CNN features vector vector



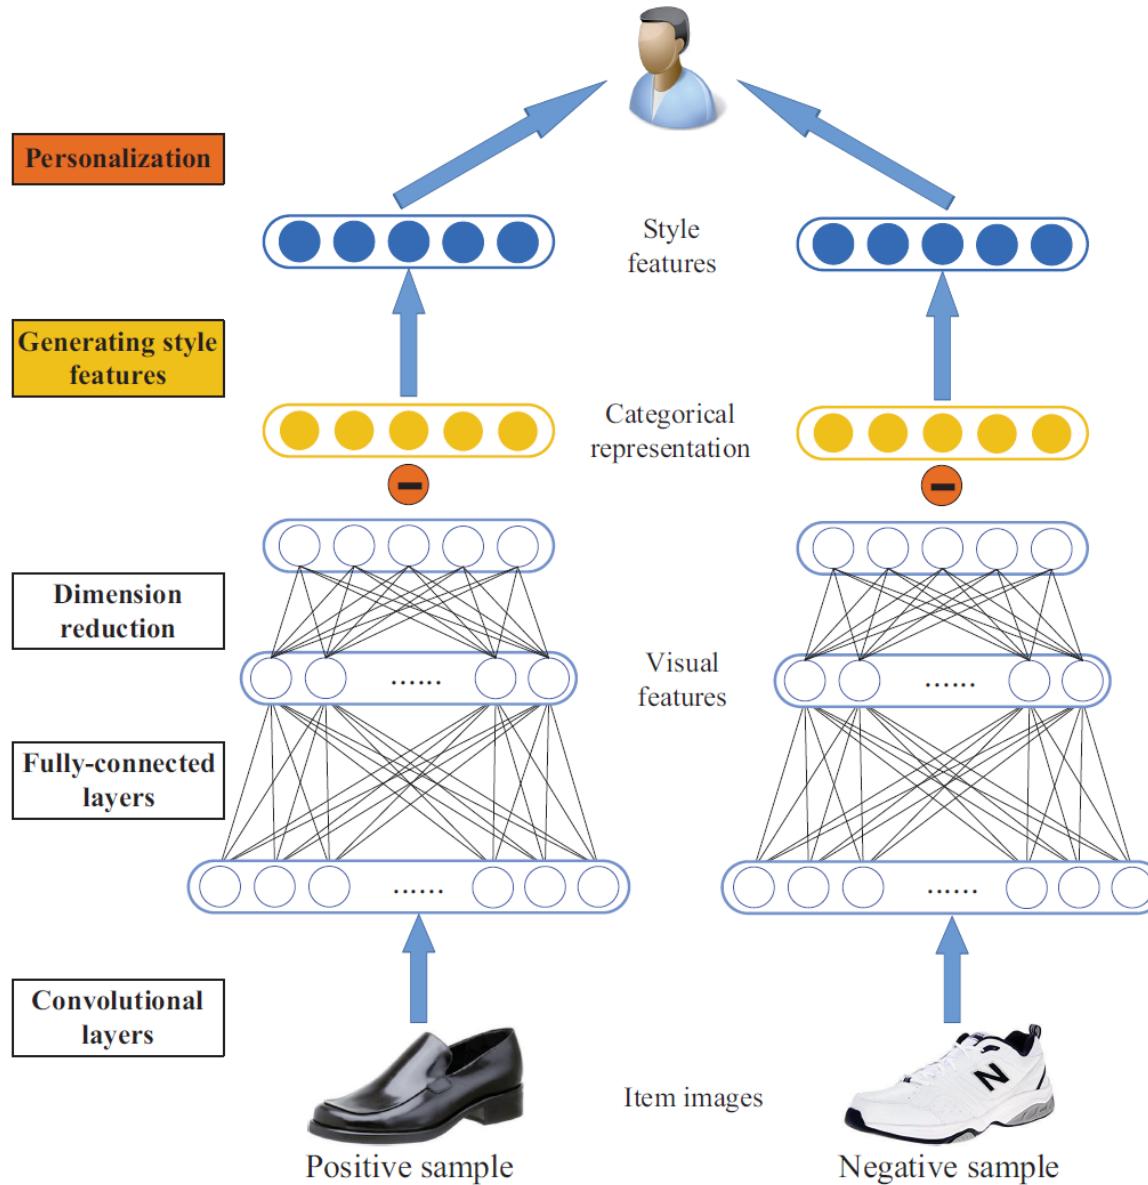
$$\mathbf{s}_i = \mathbf{Ef}_i - \mathbf{c}_i$$

Prediction model:

$$\hat{y}_{ui} = b_u + b_i + \mathbf{v}_u^T \mathbf{v}_i + \mathbf{p}_u^T (\mathbf{Ef}_i - \mathbf{c}_i)$$

Remove the effect of
category info in prediction

DeepStyle (Liu et al, SIGIR'17)



Experimental Results

AUC score on personalized ranking

dataset	setting	BPR	VBPR	DeepStyle
Clothing	warm-start	0.6243	0.7441	0.7961
	cold-start	0.5037	0.6915	0.7317
Home	warm-start	0.5848	0.6845	0.7155
	cold-start	0.5053	0.6140	0.6396

Room of improvement by learning better image representation.

Visualization on clusters of item style vector: $\mathbf{s}_i = \mathbf{E}\mathbf{f}_i - \mathbf{c}_i$



Each cluster mixes item of different categories, but they follow the same style!

Visually Explainable CF (Chen et al, 2018)

- Users care about different visual features even on the same item.



A

★★★★★ Great material, loose fit around the waist

By Maureen Button on November 2, 2017

Size: Medium/US 8-10 | Color: Black | Verified Purchase

Great material, loose fit around the waist. **Nice wide neck opening, very stylish looking.**

B

★★★★★ I absolutely love this tunic

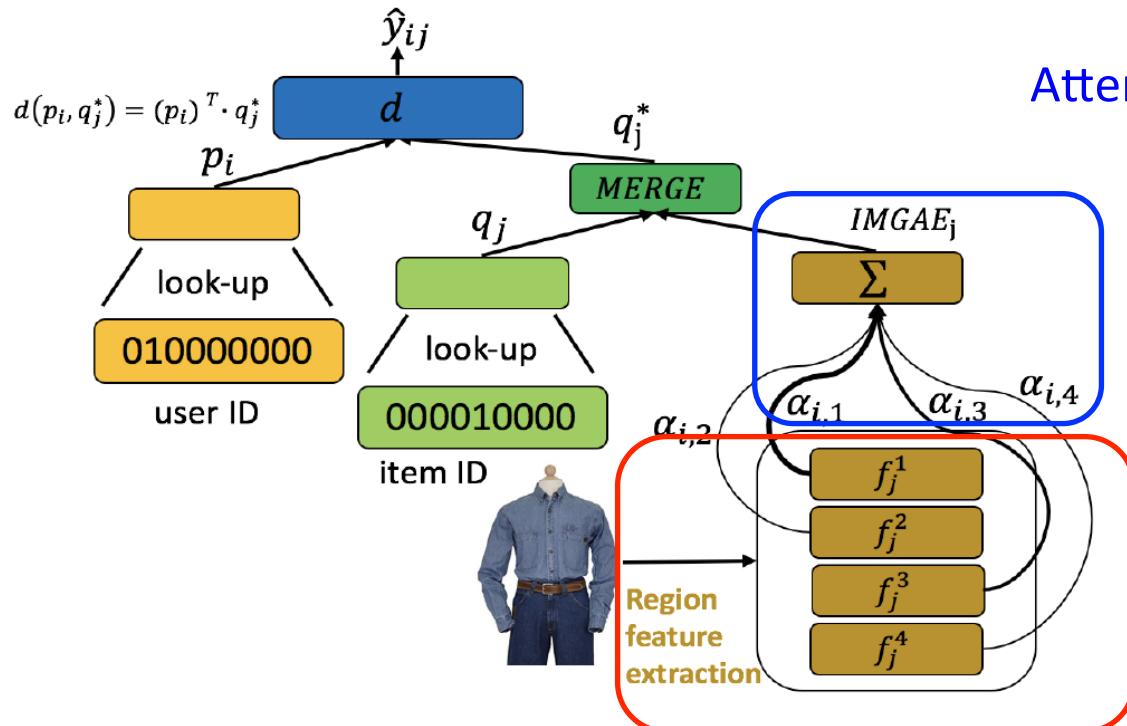
By Amazon Customer on November 30, 2017

Size: Small/US 4-6 | Color: Wine | Verified Purchase

The M fits more like a tunic where I'm fine wearing tights/leggings underneath. Nice quality, incredibly soft (especially the blue one) and **really nice pocket size**. Received numerous compliments on this.

Visually Explainable CF (Chen et al, 2018)

- Uncover **region-level** user preference with attention
 - Different **regions** of an image affect a user differently.
 - Originated from Attentive CF (Chen et al, SIGIR'17)



(a) Visually explainable collaborative filtering (VECF)

Attentive image representation:

$$IMAGE_i = \sum_{r=1}^R \alpha_{u,i,r} \cdot \mathbf{f}_r^i$$

$$\alpha_{u,i,r} = \text{softmax}(\mathbf{w}^T \mathbf{p}_u + \mathbf{h}^T \mathbf{q}_i + b)$$

Divide image by 14×14 , each region is fed into a pre-trained VGG to get a 512-dim vector.

Visually Explainable CF + User Reviews (Chen et al, 2018)

- Attention Net can learn user preference on regions, but may not be accurate in doing this.
- User review may help!

levaca Womens Long Sleeve Button Cowl Neck Casual Slim Tunic Tops With Pockets

Price: \$19.99 & Free Return on some sizes and colors

By Maureen Button on November 2, 2017

Size: Medium/US 8-10 | Color: Black | Verified Purchase

Great material, loose fit around the waist. **Nice wide neck opening, very stylish looking.**

By Amazon Customer on November 30, 2017

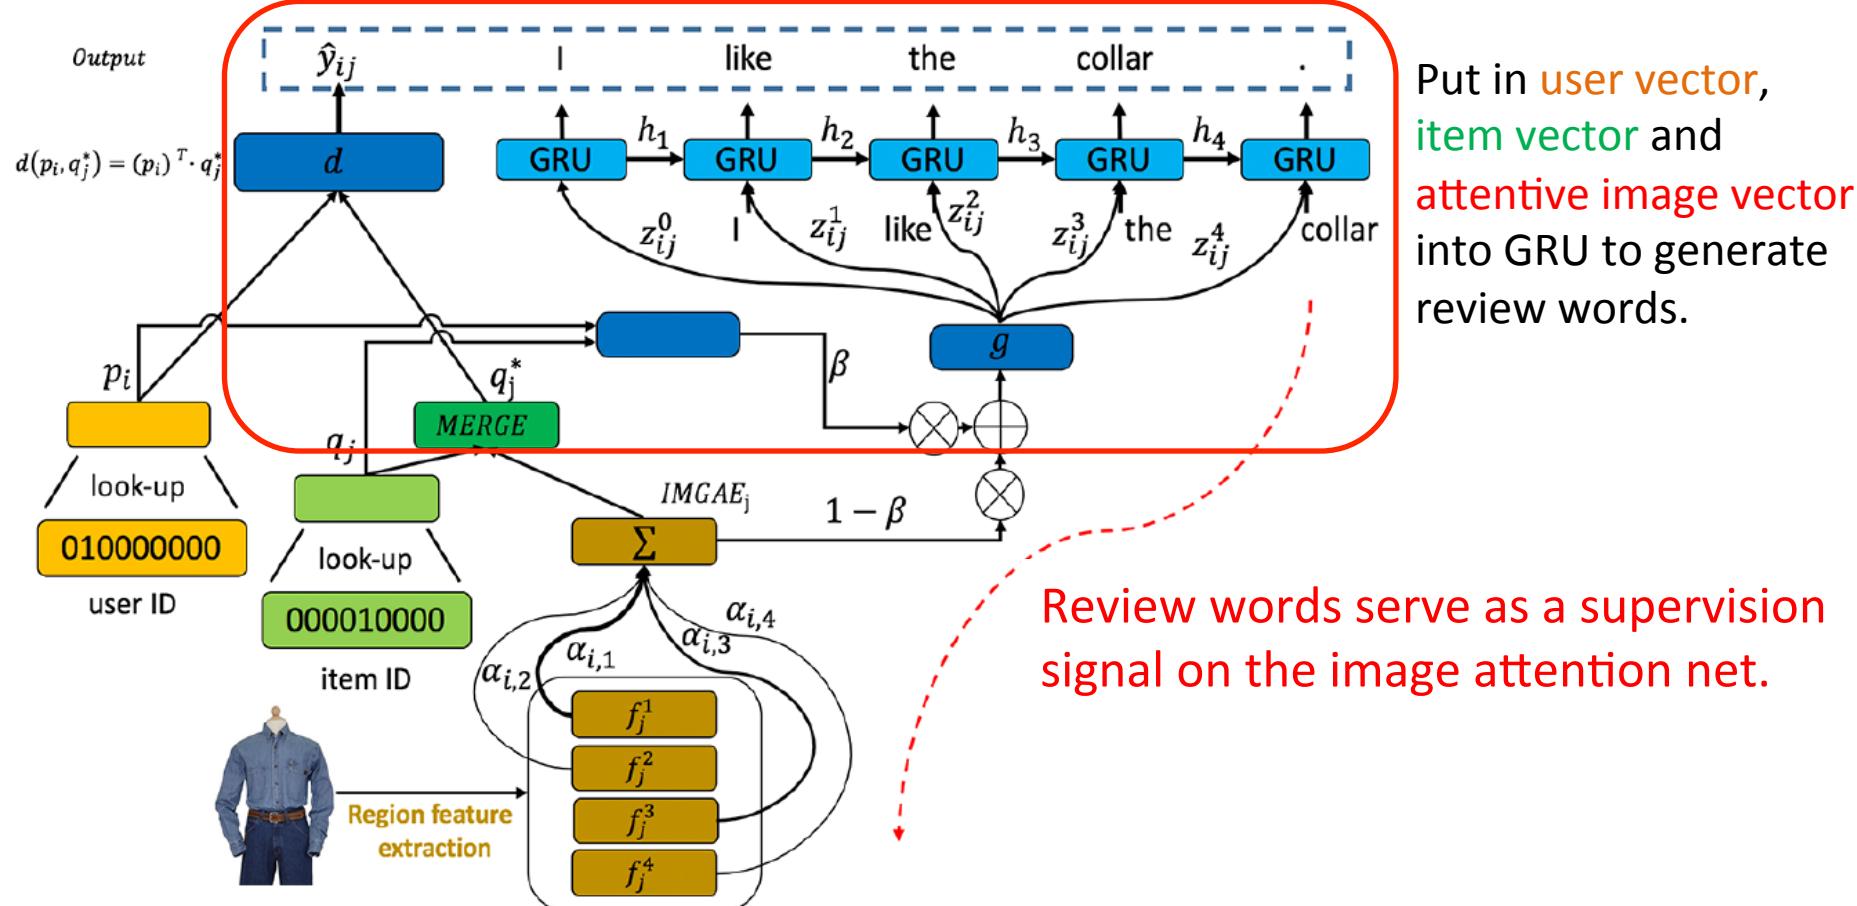
Size: Small/US 4-6 | Color: Wine | Verified Purchase

The M fits more like a tunic where I'm fine wearing tights/leggings underneath. Nice quality, incredibly soft (especially the blue one) and **really nice pocket size**. Received numerous compliments on this.

Some review words can be aligned with image regions

Visually Explainable CF + User Reviews (Chen et al, 2018)

- Generative model to incorporate review text:



(b) Review-enhanced visually explainable collaborative filtering (Re-VECF)

Visually Explainable CF + User Reviews (Chen et al, 2018)

- Learning model in a multi-task way:

$$l_2 = \delta \sum_{(i,j)} \sum_{t=1}^{l_{ij}} \log p(w_{ij}^t | w_{ij}^{1:t-1}, z_{ij}^{t-1})$$

Task #1: review generation
(error of generating each word)

$$+ (1 - \delta) \left(\sum_{i \in \mathcal{U}} \sum_{j \in \mathcal{V}_+^i} \log \hat{y}_{ij} + \sum_{i \in \mathcal{U}} \sum_{j \in \mathcal{V} / \mathcal{V}_+^i} \log(1 - \hat{y}_{ij}) \right) - \lambda \|\Theta\|_F^2$$

Task #2: user preference prediction
(point-wise cross entropy loss)

Experimental Results

Dataset	Men			Women		
Measure@5(%)	F_1	HR	NDCG	F_1	HR	NDCG
BPR	1.209	3.901	0.740	0.897	3.342	0.611
HFT	1.242	4.243	0.757	0.915	3.371	0.631
VBPR	1.361	4.261	0.773	0.929	3.402	0.648
VECF	1.378	4.373	0.791	0.948	3.523	0.669
Re-VECF	1.442	4.803	0.846	0.985	3.587	0.712

HFT(review) vs. VBPR (image)

VECF (Visually Explainable CF) vs.
Re-VECF (Review-enhanced VECF)

Image Attention with Review

#	Target Item	Historical Records	Textual Review	Visual Explanation	
				VECF	Re-VECF
2		 	<i>this is a really comfortable v-neck i found that the size and location of the v are just right for me. i'm 5'8 & #34, but 200 lbs (and dropping :)</i>		
3		 	<i>Great leggings, perfect for fly fishing or hunting or running. just perfect anytime you are cold!</i>		
4		  	<i>The socks on the shoes are a perfect fit for me. first time with a shoe with the speed laces and i like them a lot</i>		
5		 	<i>Really like these socks! they are really thick woolen socks and are good for cold days. they cover a good portion of your feet as they go a little (halfway) above the calf muscle area.</i>		
6		 	<i>I like the front pocket~. Very cool!</i>		

Short Summary

- The quality of image features is crucial for visually-aware recommendation.
 - Attention over regions is helpful
- After projecting image features to embedding space, it becomes a standard recommendation problem.
 - Any advanced preference learning techniques introduced in Part I can be used.

Reference

- He, Ruining, and Julian McAuley. "VBPR: Visual Bayesian Personalized Ranking from Implicit Feedback." In AAAI 2016.
- Chen, Jingyuan, Hanwang Zhang, Xiangnan He, Liqiang Nie, Wei Liu, and Tat-Seng Chua. "Attentive collaborative filtering: Multimedia recommendation with item- and component-level attention." In SIGIR 2017.
- Liu, Qiang, Shu Wu, and Liang Wang. "DeepStyle: Learning User Preferences for Visual Recommendation." In SIGIR 2017.
- Chen, Xu, Yongfeng Zhang, Hongteng Xu, Yixin Cao, Zheng Qin, and Hongyuan Zha. "Visually Explainable Recommendation." *arXiv preprint arXiv:1801.10288* (2018).
- Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." In NIPS 2012.
- Simonyan, Karen, and Andrew Zisserman. "Very deep convolutional networks for large-scale image recognition." In ICLR 2015.

Outline of Tutorial

- Background (Xiangnan, 10 mins)
- Basics & Advances in Recommendation (Xiangnan, 50 mins)
- Visually-aware Product Recommendation (Xiangnan, 30 mins)
- Break (15 mins)
- Visual Representation (Hanwang, 45 mins)
- Image/Video Recommendation (Hanwang, 25 mins)
- Conclusion (Hanwang, 5 mins)

Social Multimedia Data

- *Interaction with MM Content on Social Networks*



151,974

423 86



Hope Valley Light....

Hope Valley bathed in late afternoon hazy sunlight, taken from the top of Over Owler Tor. Sony A7II/24-70mm.

Press L for a larger and nicer view :-)

8,632
views

514
faves

74
comments



喜欢 收藏 下载 分享

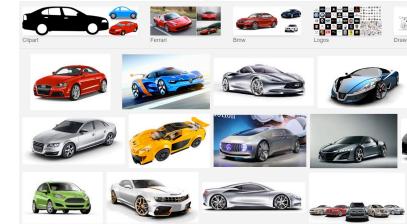
| 听友评论 (已有797条评论)

淘宝网
Taobao.com

买家 [登录后可见]	淘宝价	数量	付款时间	款式和型号
就**0 ❤️❤️❤️ 🎉	¥58	1	2015-12-01 15:52:35	颜色分类:红蓝格纹
t**9 ❤️❤️	专属优惠	1	2015-12-01 15:51:04	颜色分类:红蓝格纹
相**8 🎉	专属优惠	1	2015-12-01 15:49:45	颜色分类:红蓝格纹



Google
images

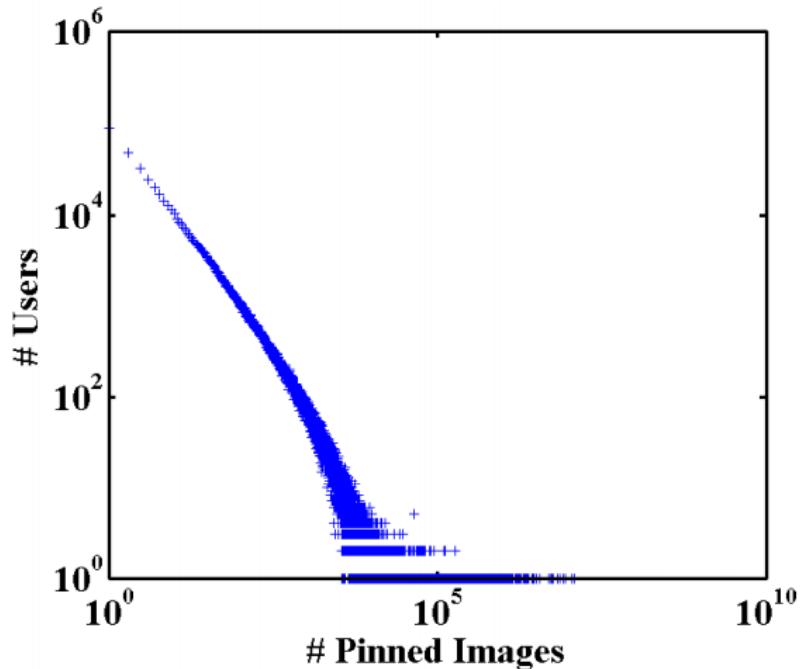


Challenges

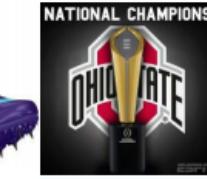
- Large, sparse Image-User matrix

1M Flickr Image and 30K Tags: **99.7%** sparsity

1M Pinterest Image and 10K Users: **99.1%** sparsity



Sports



Travel



Animals



Sparse

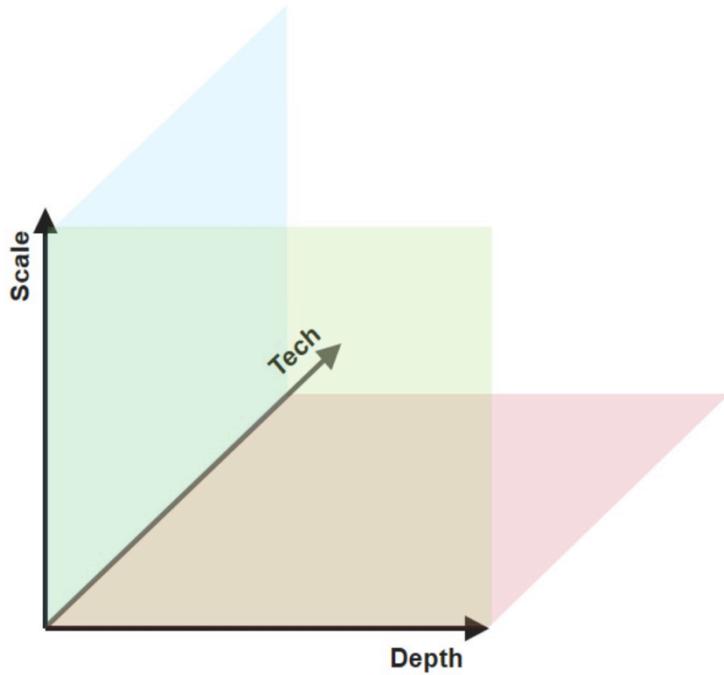
$$X \approx U'V$$

Diverse

Outline

- State-of-the-art visual representations (45min)
 - Image, object feature
 - Video feature
 - Dynamic representation (visual attention)
- Deep embedding visual representation into user-item matrix (25min)
 - Weakly, semi-supervised learning

The Feature Evolution Space

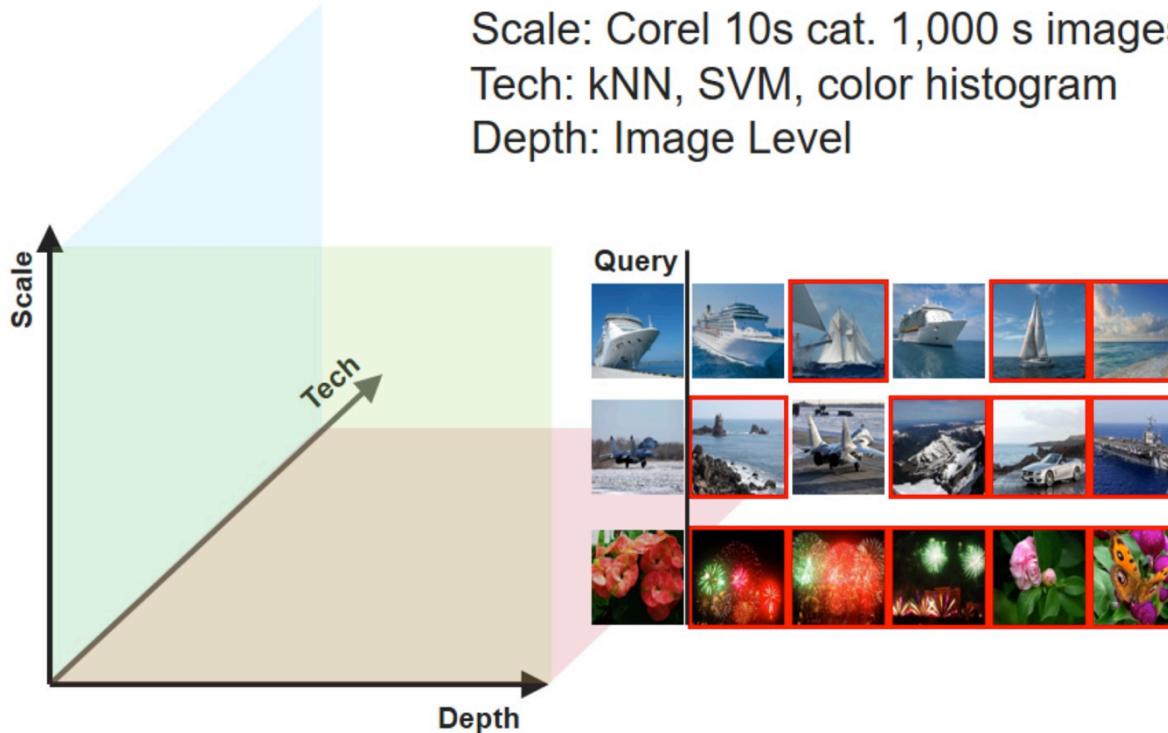


- **Scale**
 - Data, concept
- **Depth**
 - Content understanding
- **Tech**
 - Model

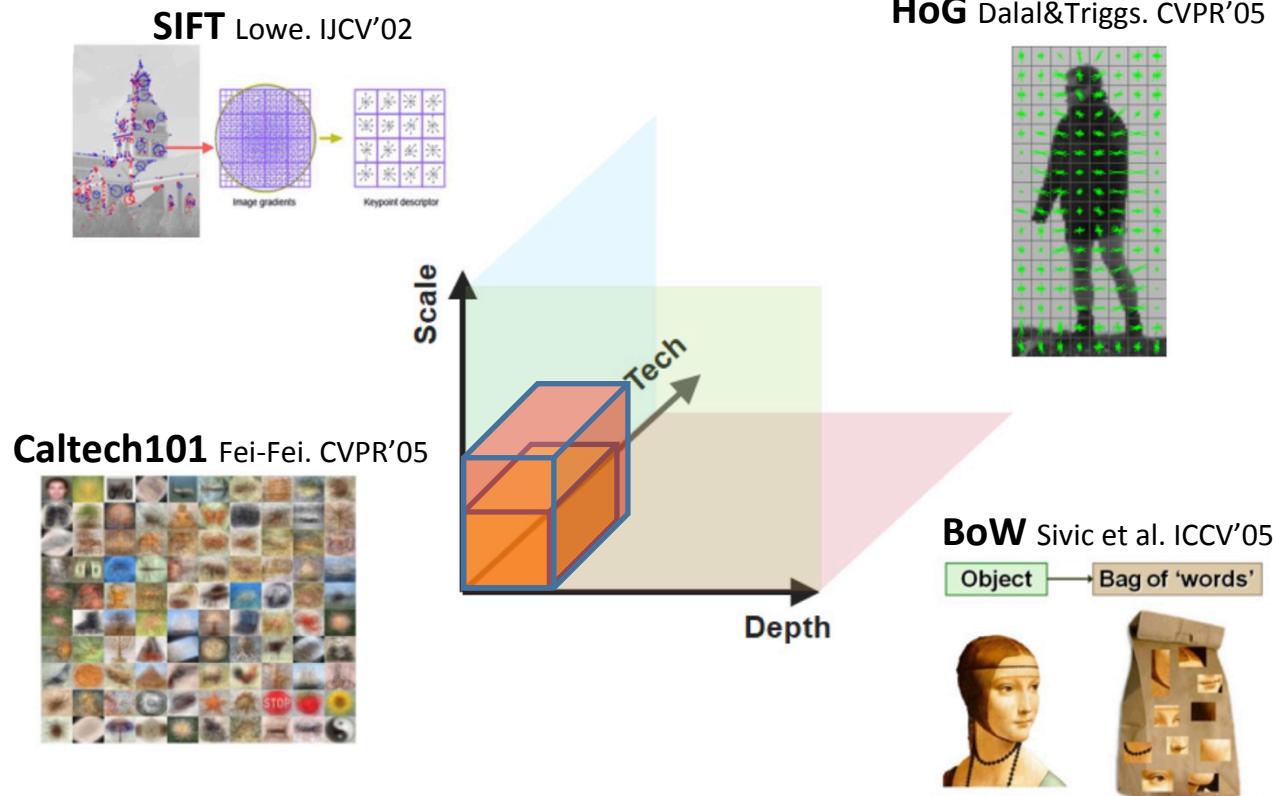
Near 2000: At the end of the early stage

Smeulders et al. TPAMI'00

Scale: Corel 10s cat. 1,000 s images
Tech: kNN, SVM, color histogram
Depth: Image Level

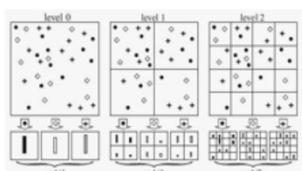


2000~2005: The Rise of Local Features

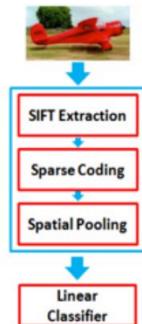


2005~2011: The Mature of Shallow Methods

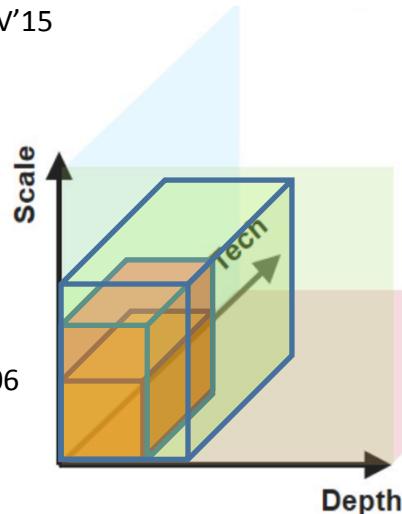
VOC Everingham et al. IJCV'15



SPM Lazebnik et al. CVPR'06



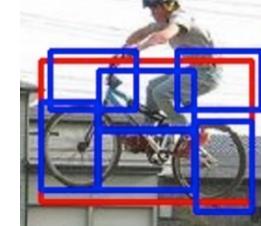
ScSP Yang et al. CVPR'09



ILSVRC Deng et al. CVPR'09



DPM Felzenszwalb et al. CVPR'08



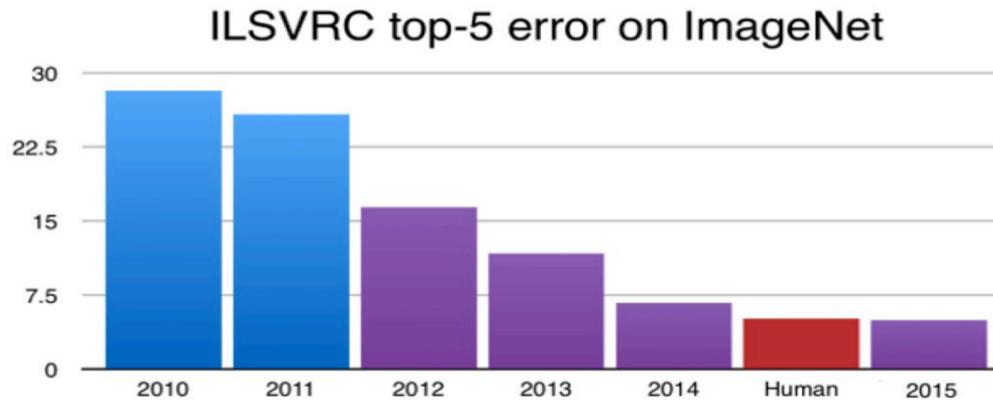
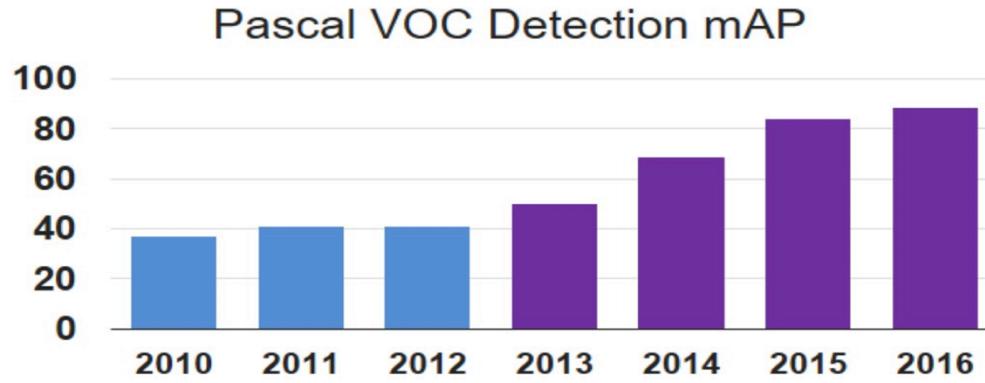
VLAD Jegou et al. CVPR'10

FisherVector Perronnin et al. ECCV'10

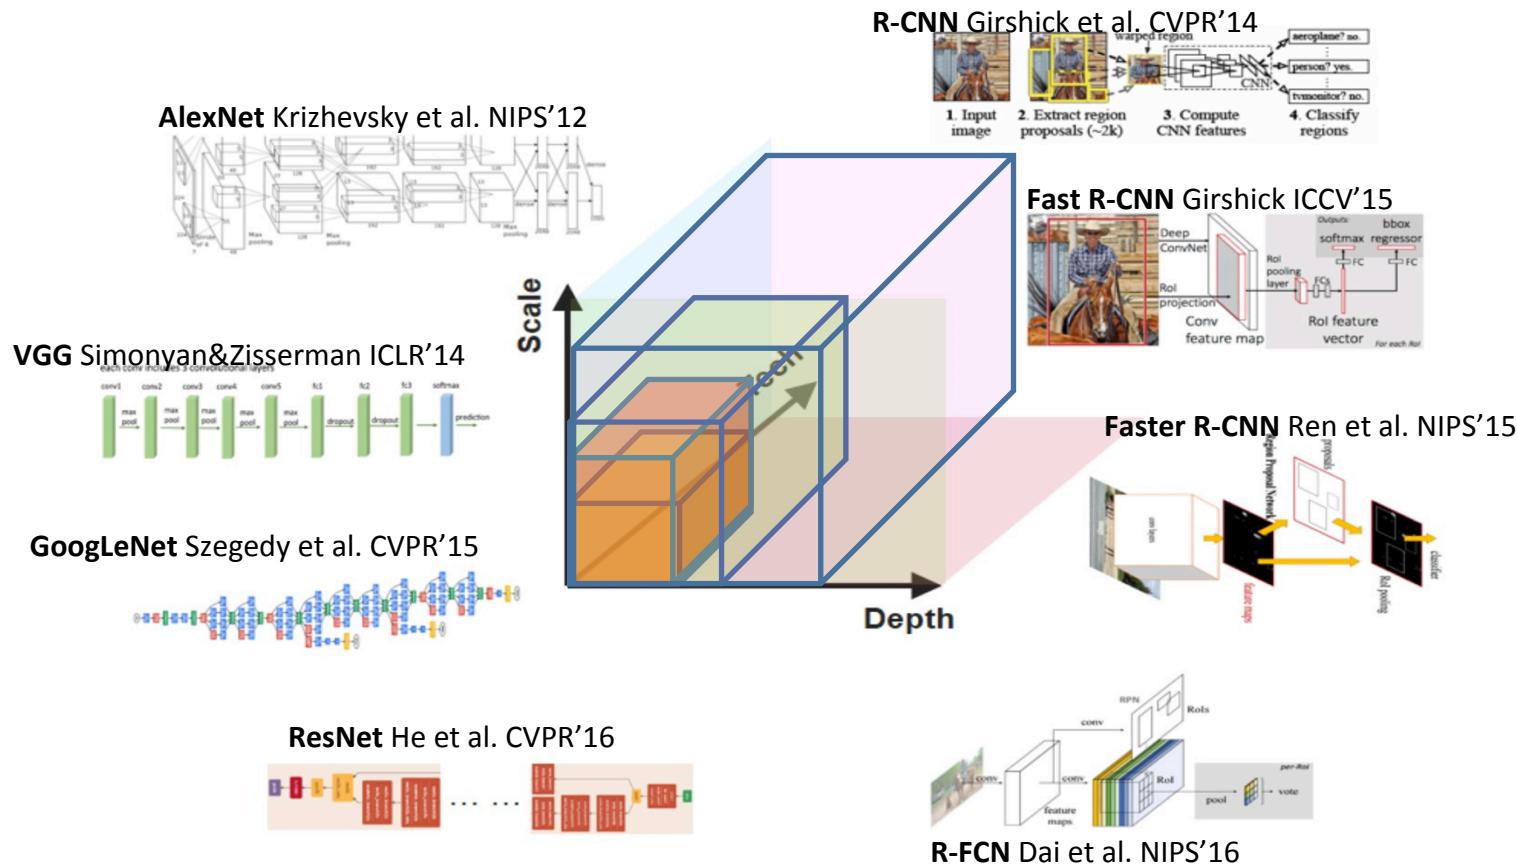


Tricks Chatfield et al. BMVC'11

2012: The Revolution



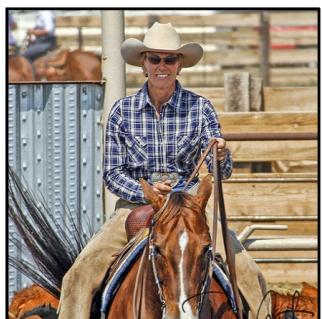
2012~Now: The DL Republic



Object Detection

--- *On the way of sharing context features*

R-CNN [Girshick et al. CVPR'14]

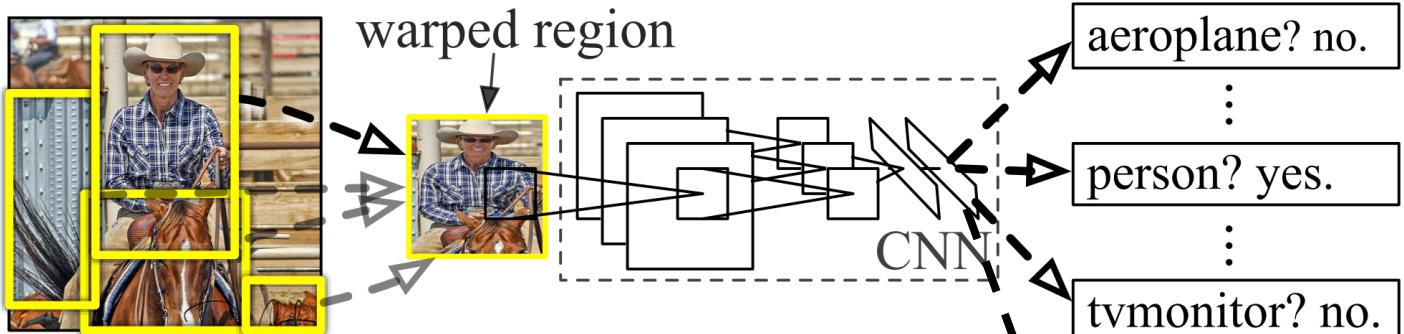


1. Input image

2. Extract region proposals (~2k)

No context shared (66.0% mAP)

Too many CNN pass (0.02 fps)



3. Compute CNN features

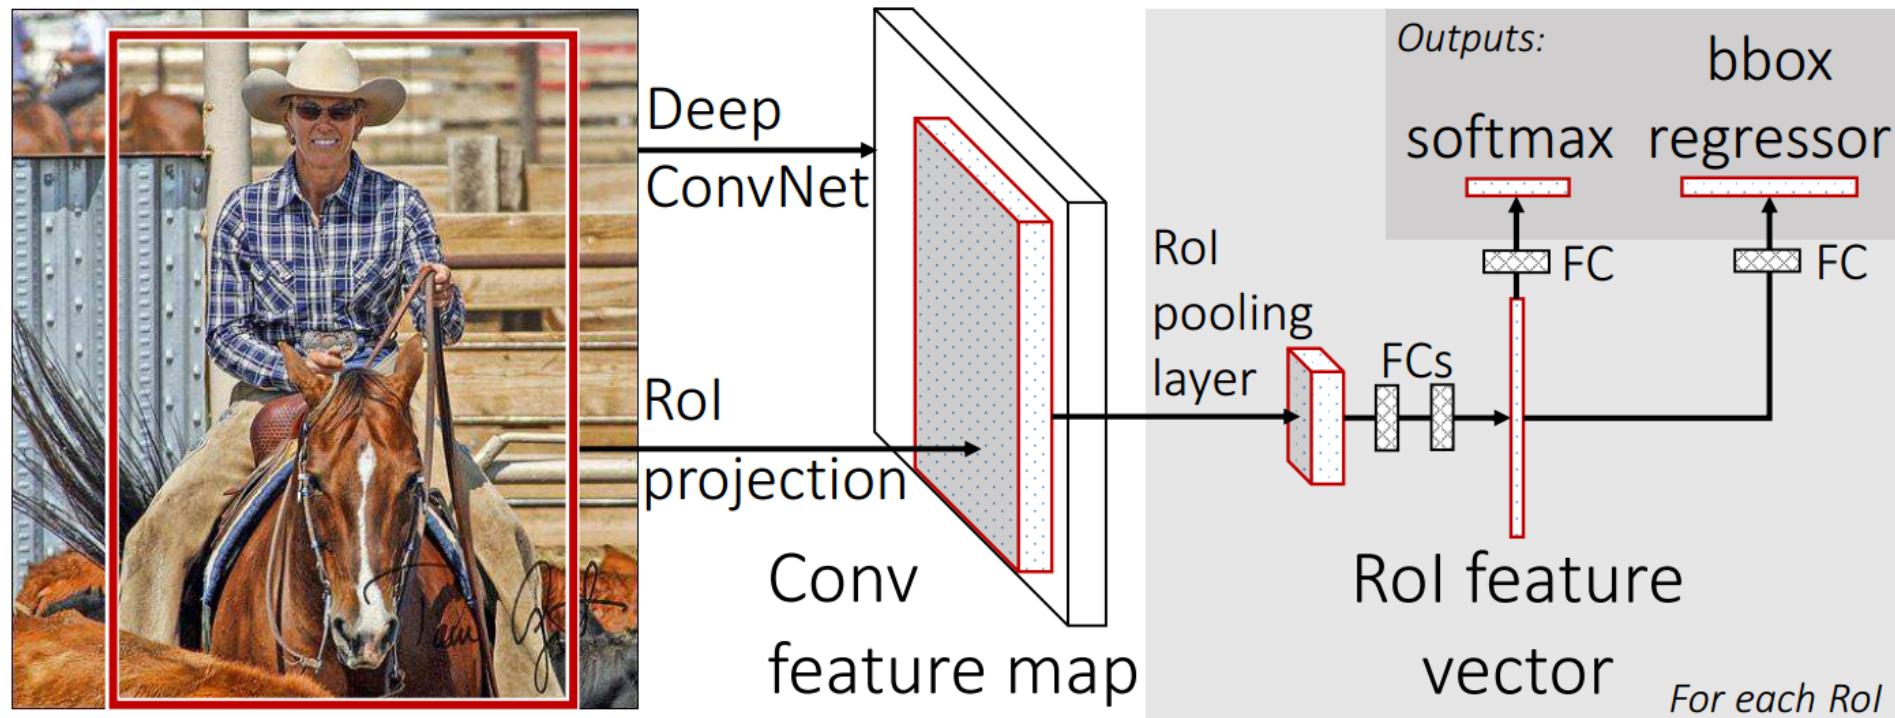
4. Classify regions

5. Bounding Box Regression

Object Detection

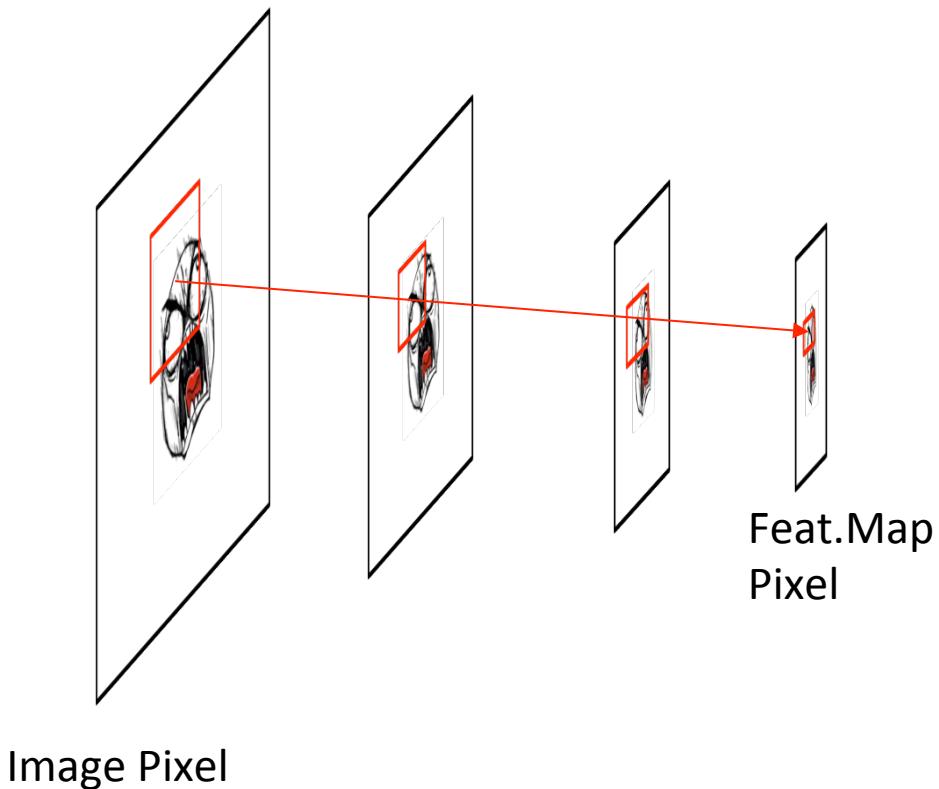
--- *On the way of sharing context features*

Fast R-CNN [Girshick et al. ICCV'15]



Object Detection

--- *On the way of sharing context features*

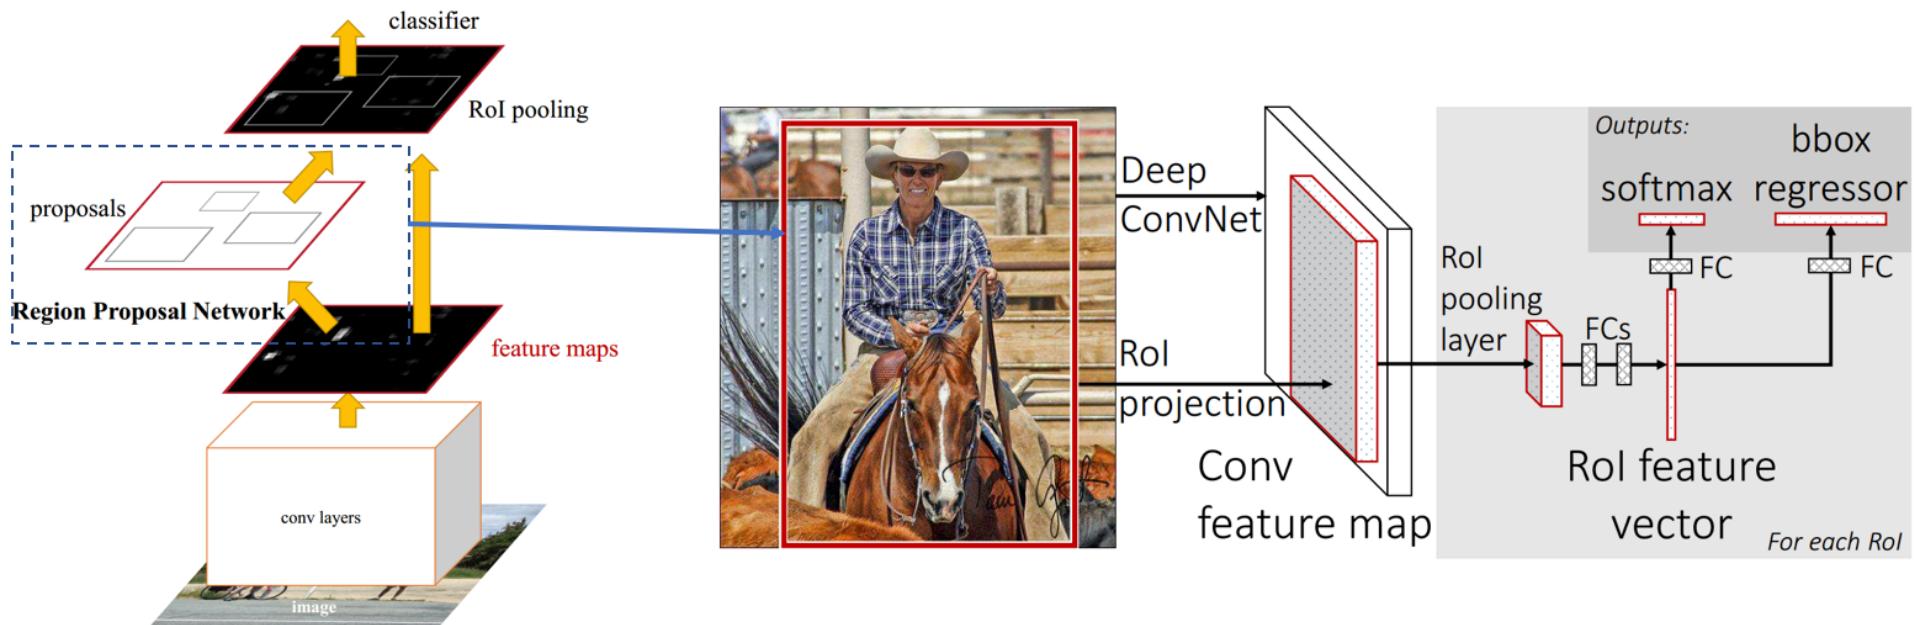


RoI pooling encapsulate context

Object Detection

--- *On the way of sharing context features*

Faster R-CNN [Ren et al. NIPS'16]

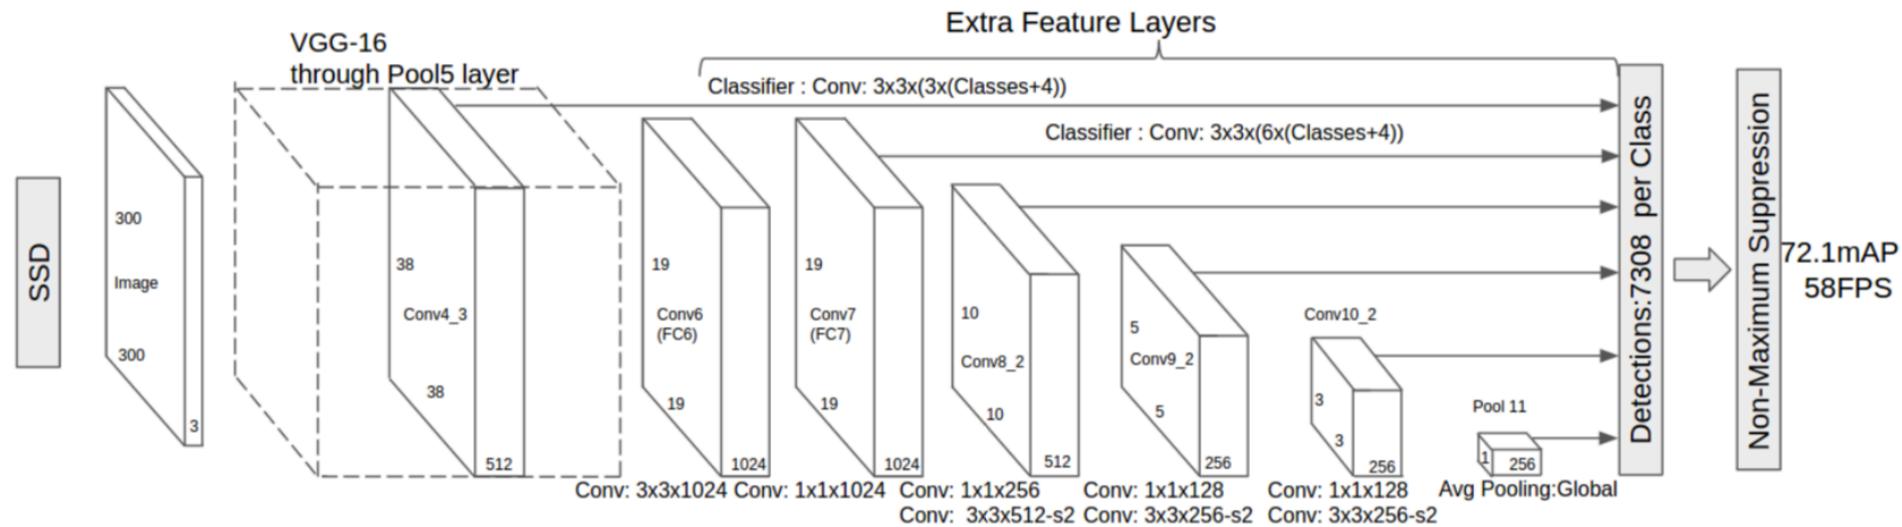


Context can help to know “where”: conv to generate proposals (73.2% mAP, 5 fps)

Object Detection

--- *On the way of sharing context features*

SSD [Liu et al. ECCV'16] (see also YOLO9K [Redmon et al CVPR'17])

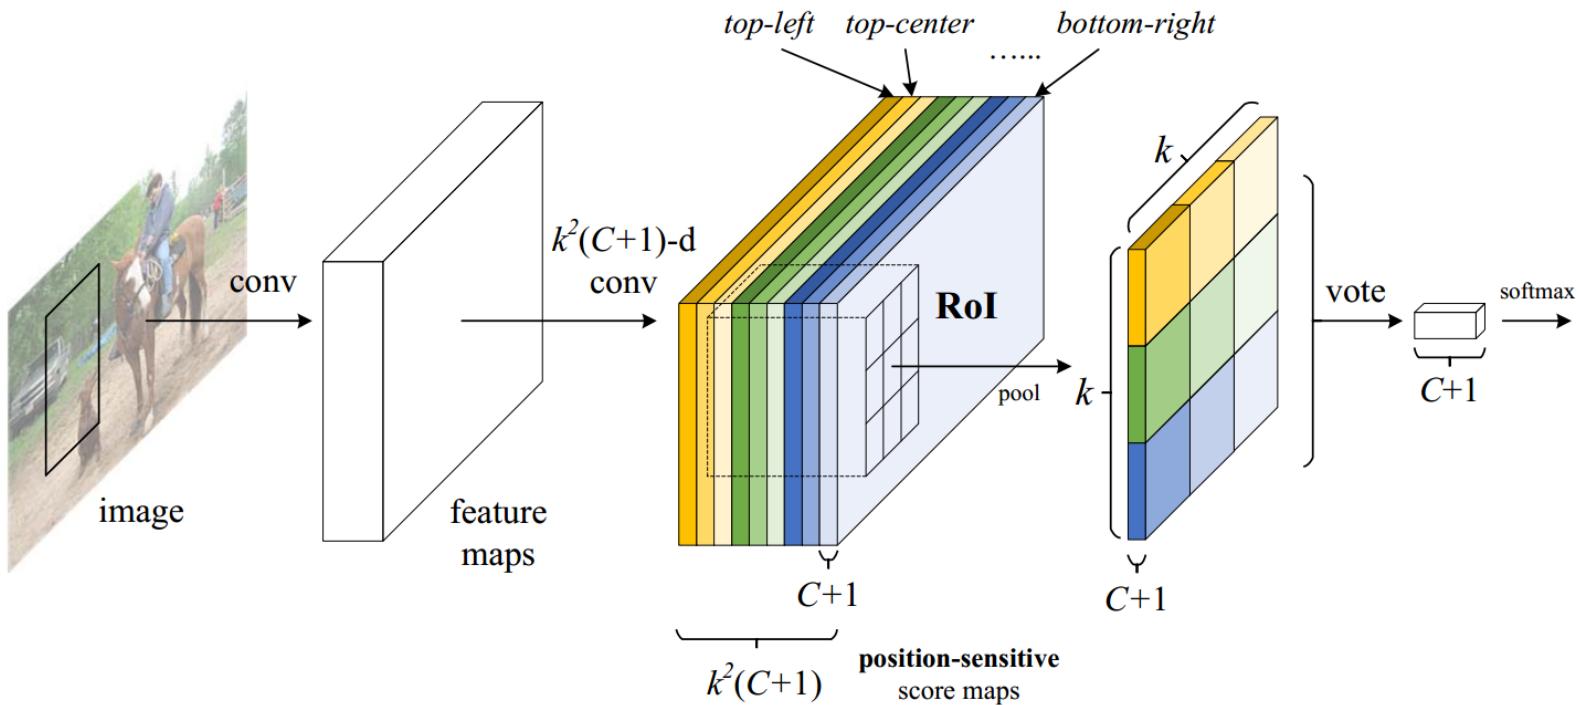


Context can help to know “where” and “what” simultaneously:
Conv at multi-layers to generate bbox and classes (75.1% mAP, 58 fps)

Object Detection

---- *On the way of sharing context features*

R-FCN[Dai et al. NIPS'16]



Context can help to know “where” and “what” **separately**:

Conv to generate bbox (RPN) and classes (RoI pooling)

80.5% mAP, 6fps

Message To-Go !

- Image feature
 - Apply CNN
 - Take the last layer output
- Object feature
 - Apply CNN-based object detector
 - Get the boxes (or masks)
 - Take the RoI feature who generates them

Video Representation

- The best hand-crafted: iDT
- The best spatial-temporal convolution: (2+1)D + Two-Stream
- The best choice: Frame Pooling



Body up?



Kissed?



Weapon or arm forward?

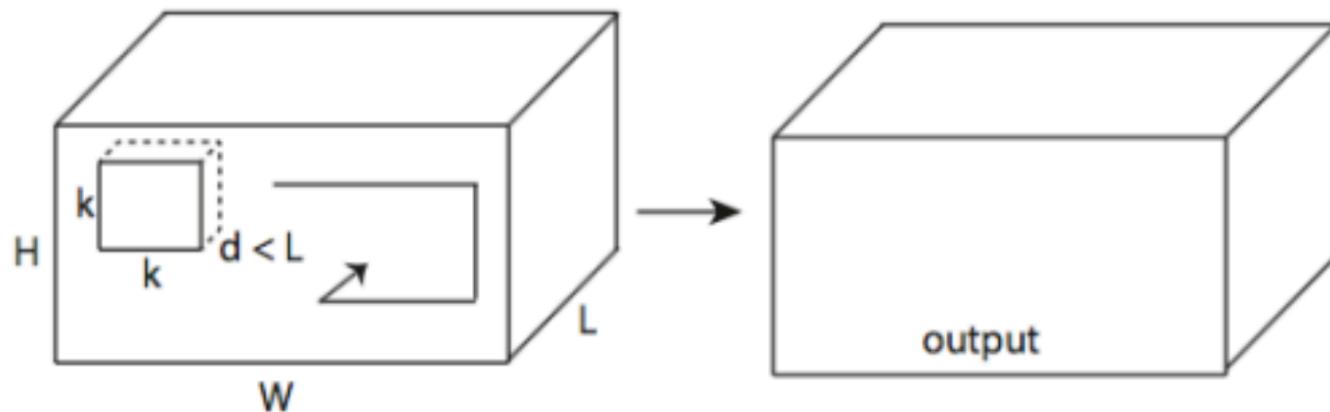
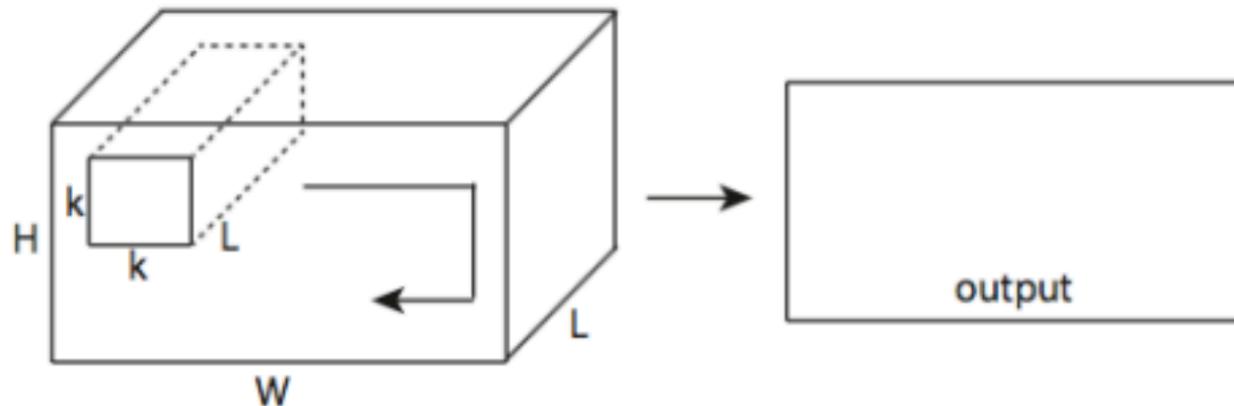
Improved Dense Trajectories (iDT)

[Wang & Schmid. ICCV'13]



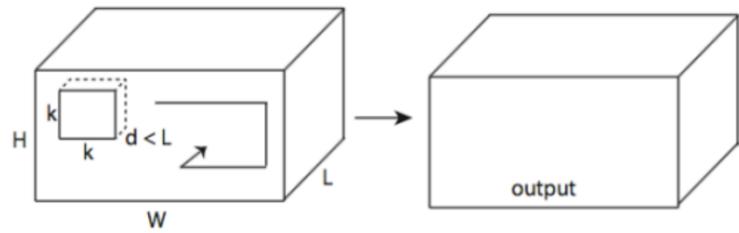
- Trajectory feature point + HOG descriptor + Fisher Vector Encoding + BoW \approx 10K dimensions
- TRECVID MED challenge 2013 and THUMOS'13 action recognition challenge winners
- Slow and high-dimensional

Spatial-Temporal Convolution

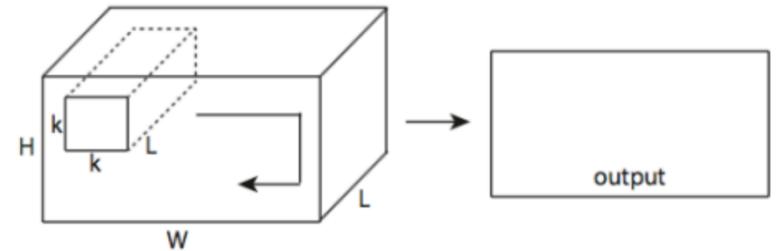


Pre-Training for Features

Sports 1M [Karpathy et al. CVPR'14]
Videos: 1.1M Class: 487



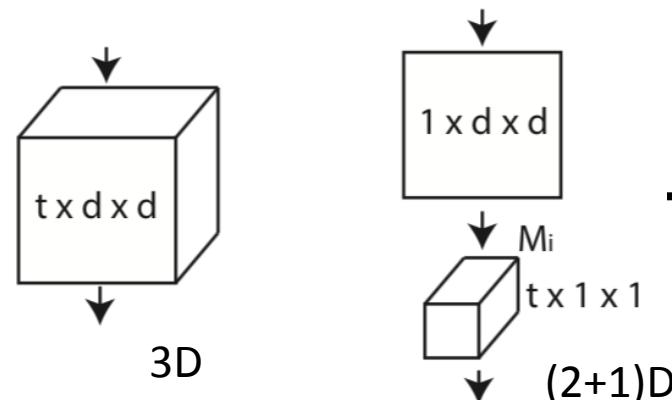
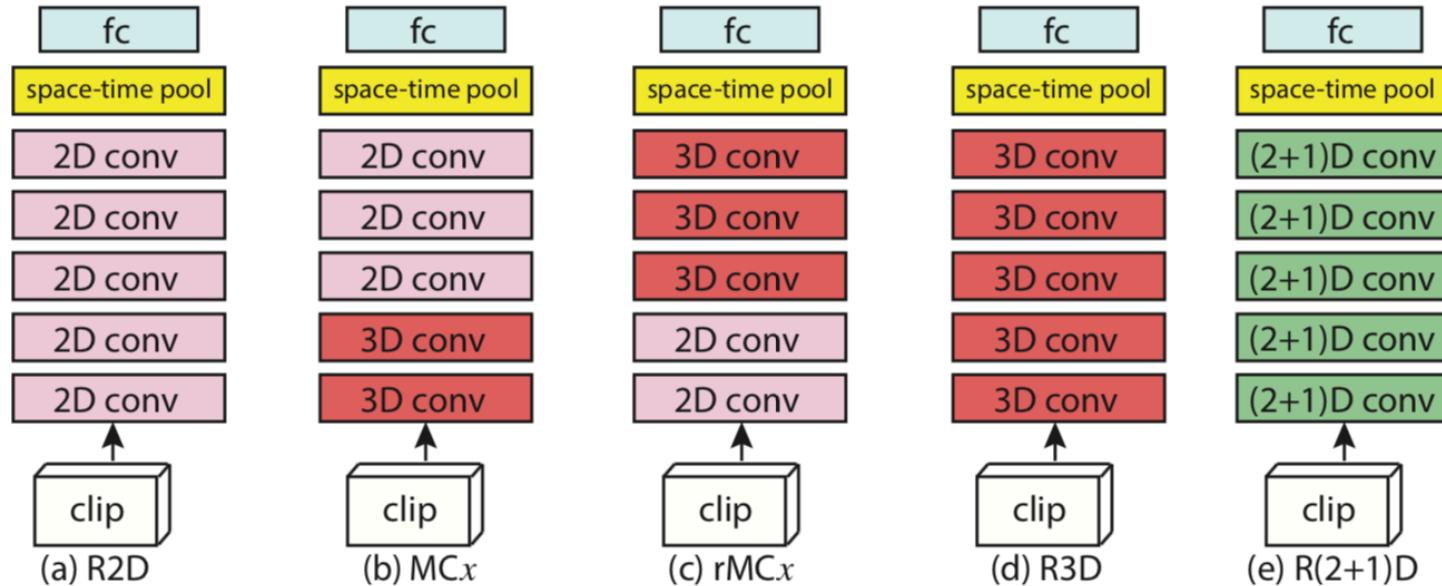
ImageNet [Russakovsky et al. IJCV'15]
Images: 1.2M Class: 1000



Bostrapping 3D from 2D [Mansimov et al. arXiv'15]
Boring video from ImageNet + Target video set

More ST Conv. Designs

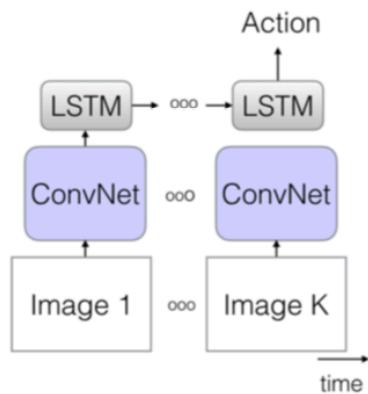
[Tran et al. CVPR'18]



More ST Conv. Designs

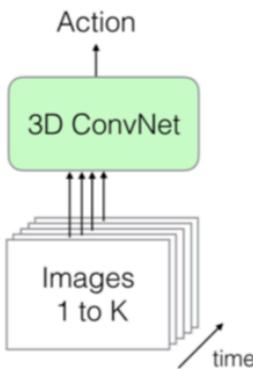
[Donahue et al. CVPR'15]
[Srivastava et al. ICML'15]

a) LSTM



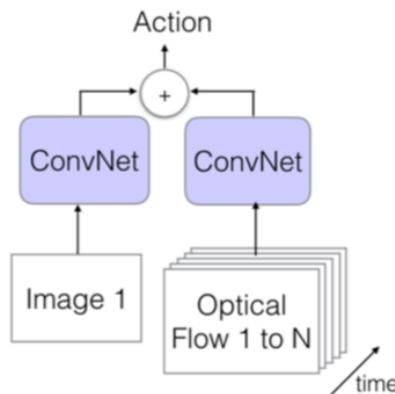
[Tran et al. CVPR'15]

b) 3D-ConvNet



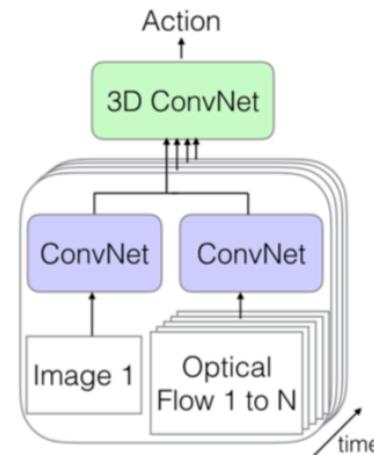
[Simonyan&Zisserman.
NIPS'15]

c) Two-Stream



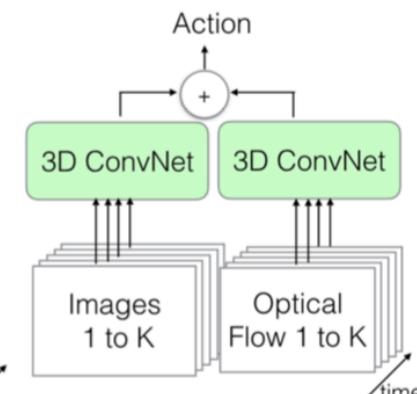
[Feichtenhofer et al. CVPR'16]

d) 3D-Fused
Two-Stream



[Carreira &Zisserman. CVPR'17]

e) Two-Stream
3D-ConvNet



Message To-Go!

Performances on Sports 1M

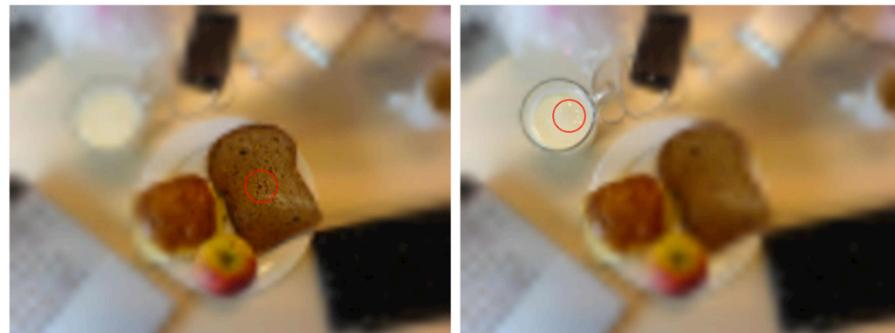
Method	C3D	2D+ Pooling	(2+1) D	Two-Stream	(2+1)D+ Two-Stream
Video@5	85.2%	90.4%	91.5%	87.4%	91.9%

- Image-level is very strong
- Video-level requires tricky pre-training
- If you are lazy, just use image-level+pooling

Visual Attention

--- dynamic visual feature

- Inspired from cognitive science, visual saliency is a special case with prior.



[Jiang et al. CVPR'16]

- Pioneer works on signal band-width savings [Mnih et al. NIPS'15]
- Great success in Machine Translation [Bahdanau et al. ICLR'15] and Image Captioning [Xu et al. ICML'15].

Dynamic Feature



Question?

Generic Formulations

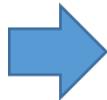
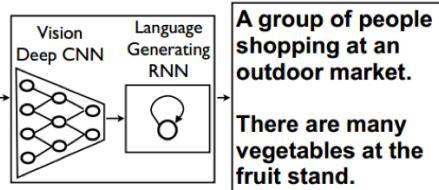
- Get a (local, partial) feature X_i
- Get a contextual information (env.) h
- Calculate an attention probability $p_i \propto F(X_i, h)$
- Calculate the new feature X
 - *Probabilistic (hard):* $X = \text{Monte Carlo}(p_i)$
 - *Deterministic (soft):* $X = \sum p_i X_i$

Notes

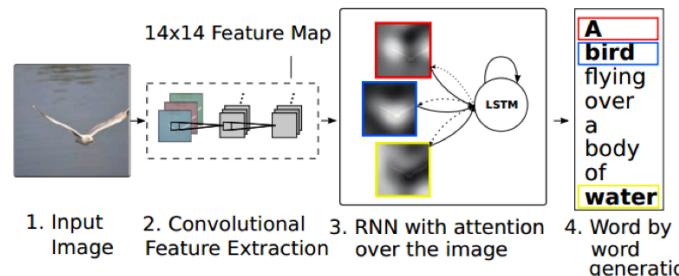
- Attention is not just “weighted sum” (soft), it is originally “discrete visual policy” (hard).
- Reinforcement learning is its natural solution.

Shallow Visual Attention

GoogleNIC (Vinyals et al. 2014)



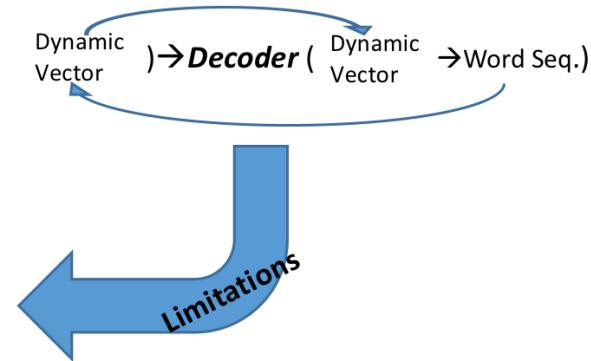
Attention (Xu et al. 2015)



Encoder (Image → CNN → Vector) → *Decoder* (Vector → Word Seq.)

Encoder (Image → CNN → Dynamic Vector) → *Decoder* (Dynamic Vector → Word Seq.)

1. Only spatial attention. What about various feature maps?
2. Only single layer. What about hierarchical attention?

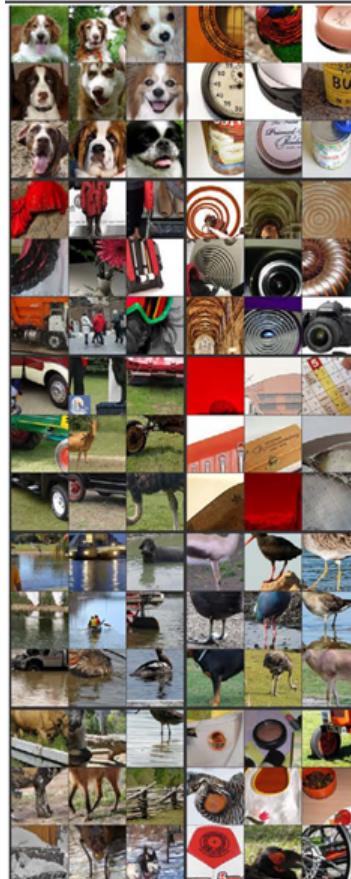


CNN Revisit

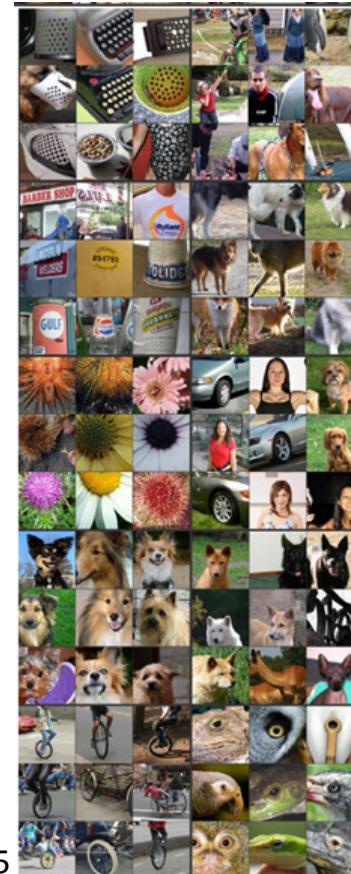
[Zeiler and Fergus. ECCV'14]



Layer3



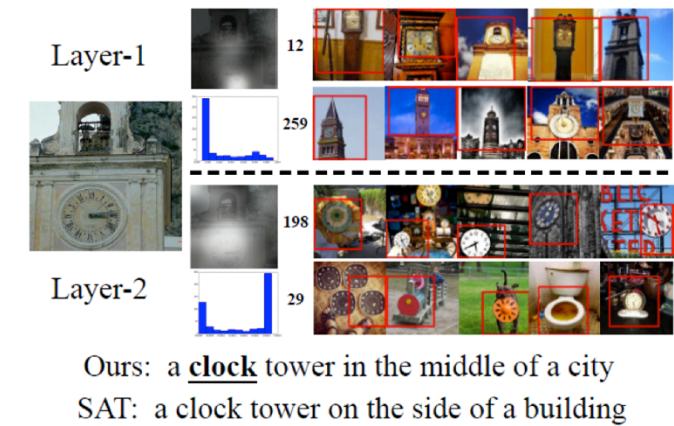
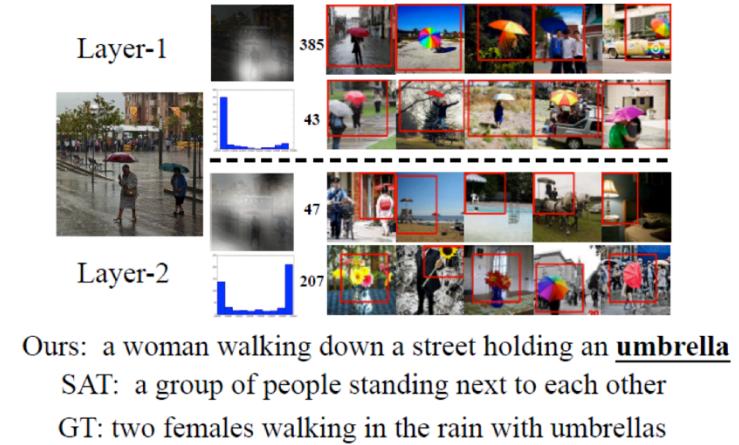
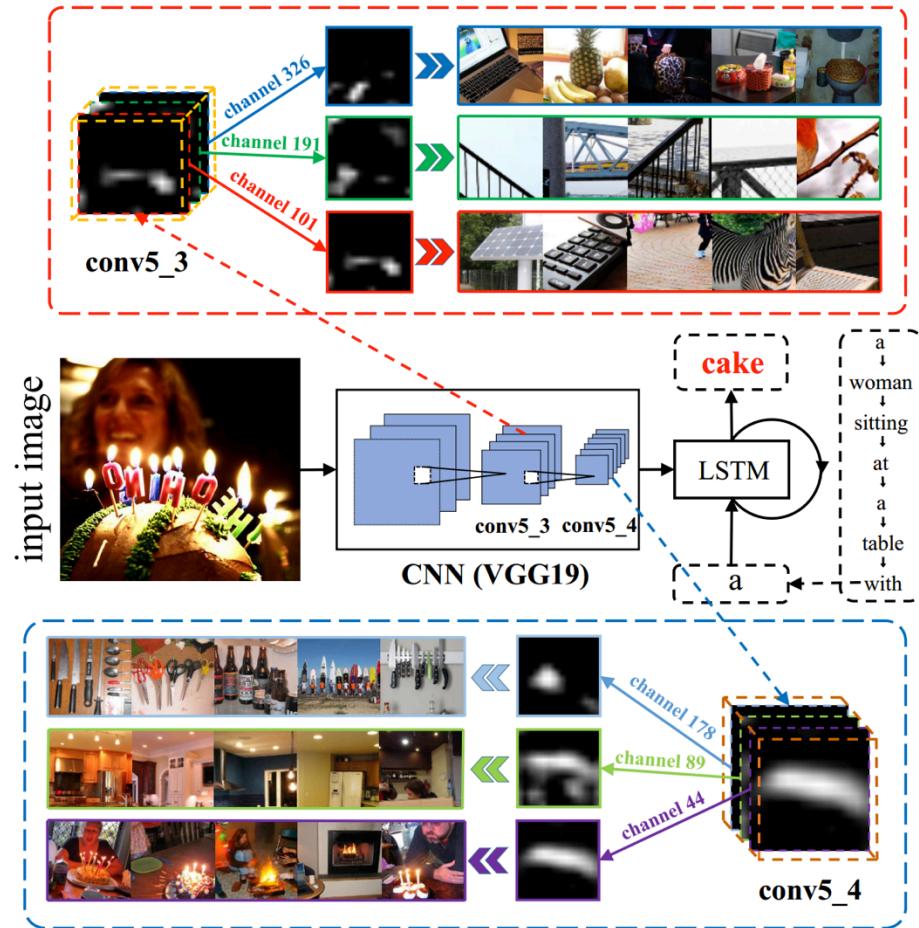
Layer4



Layer5

Spatial- and Channel-wise Attention

[Chen et al. CVPR'17]



Temporal Attention in Video

[Yao et al. ICCV'15]

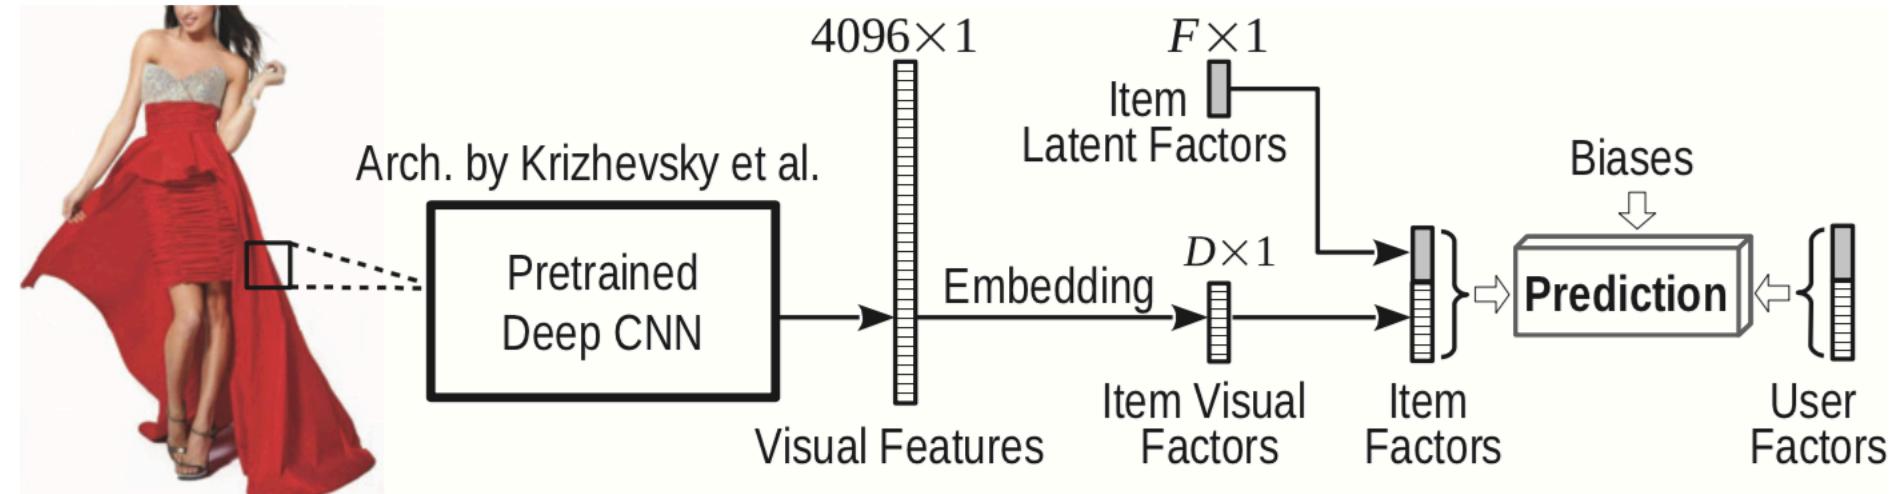


Outline

- State-of-the-art visual representations (35min)
 - Image, object feature
 - Video feature
 - Dynamic representation (visual attention)
- Deep embedding visual representation into user-item matrix (35min)
 - Weakly, semi-supervised learning

How to embed features into CF?

Visual BPR [He & McAuley. AAAI'16]



How to embed features into CF?

- Van den Oord et all. Deep content-based music recommendation.
NIPS'13
- Geng et al. One of a Kind: User Profiling by Social Curation. *MM'14*
- Geng et al. Learning Image and User Features for Recommendation in Social Networks. *ICCV'15*
- Lei et al. Comparative deep learning of hybrid representations for image recommendations. *CVPR'16*
- Chen et al. Attentive Collaborative Filtering. *SIGIR'17*
- Gao et al. A Unified Personalized Video Recommendation via Dynamic Recurrent Neural Networks. *MM'17*
- Niu et al. Neural Personalized Ranking for Image Recommendation. *WSDM'18*
- Yu et al. Aesthetic-based Clothing Recommendation. *WWW'18*

How to embed features into CF?

- Van den Oord et all. Deep content-based music recommendation.
NIPS'13
- Geng et al. One of a Kind: User Profiling by Social Curation. *MM'14*
- Geng et al. Learning Image and User Features for Recommendation in Social Networks. *ICCV'15*
- Lei et al. Comparative deep learning of hybrid representations for image recommendations. *CVPR'16*
- Chen et al. Attentive Collaborative Filtering. *SIGIR'17*
- Gao et al. A Unified Personalized Video Recommendation via Dynamic Recurrent Neural Networks. *MM'17*
- Niu et al. Neural Personalized Ranking for Image Recommendation. *WSDM'18*
- Yu et al. Aesthetic-based Clothing Recommendation. *WWW'18*

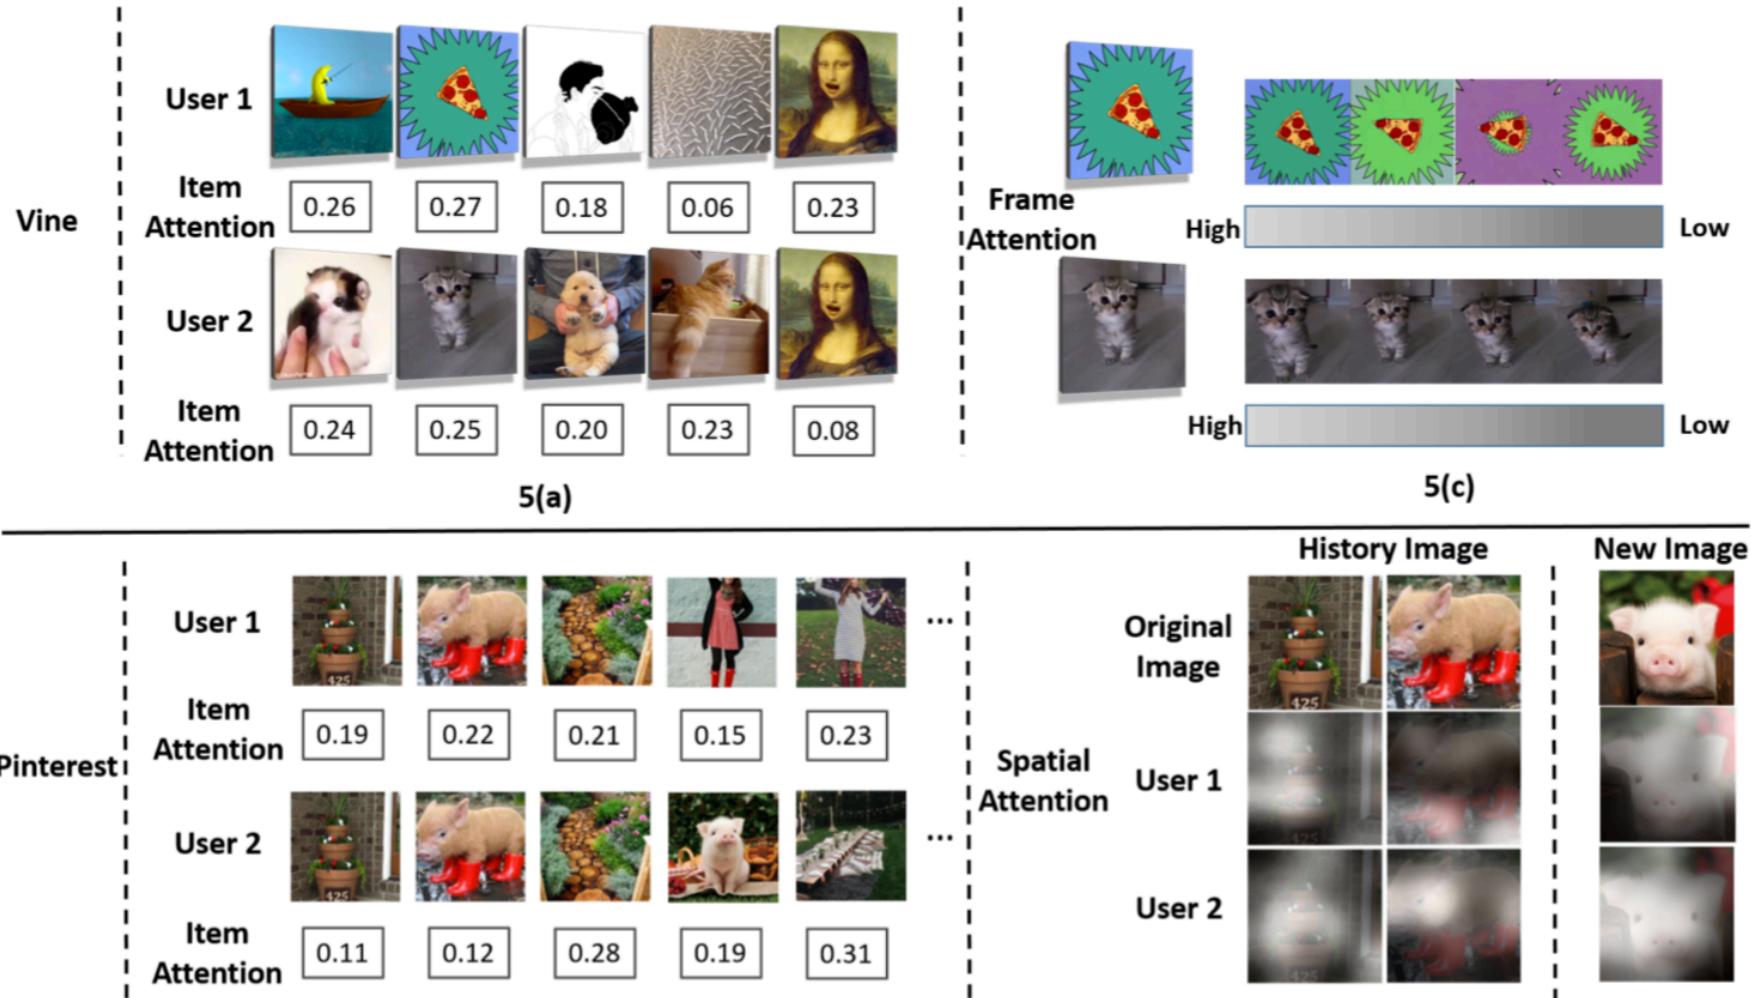
Nature: Weakly Semi-Supervised Learning

- Weak
 - Partial feedback
- Semi-
 - Implicit feedback



Towards “Weakly”: Attentive CF

[Chen et al. SIGIR’17]



ACF: Attentive Collaborative Filtering (Chen et al, SIGIR'17)

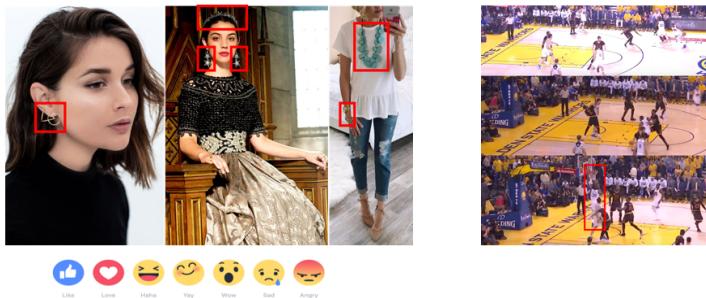
- **Input:**

user -> ID & historical interacted items.

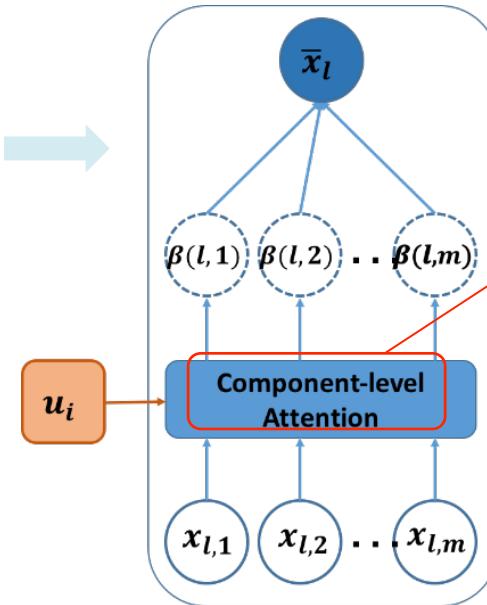
Item -> ID & visual features.

- **Item Representation:**

Component-level attention -> not all components contribute equally to an item's representation



A user's preference on different **components** of the item are not equal!



$$\bar{x}_l = \sum_{m=1}^{|\{x_{l*}\}|} \beta(i, l, m) \cdot x_{lm}$$

$$\beta(i, l, m) = \frac{\exp(b(i, l, m))}{\sum_{n=1}^{|\{x_{l*}\}|} \exp(b(i, l, n))}$$

$$b(i, l, m) = \mathbf{w}_2^T \phi(\mathbf{W}_{2u} \mathbf{u}_i + \mathbf{W}_{2x} \mathbf{x}_{lm} + \mathbf{b}_2) + c_2$$

ACF: Attentive Collaborative Filtering (Chen et al, SIGIR'17)

- **Input:**
 - user -> ID & historical interacted items.
 - item -> ID & visual features.
- **User Presentation:**
 - **Item-level attention** -> not all historical items contribute equally to a user's representation



$$\hat{R}_{ij} = \left(\mathbf{u}_i + \sum_{l \in \mathcal{R}(i)} \alpha(i, l) \mathbf{p}_l \right)^T \mathbf{v}_j$$

Attention weights learned by a neural net
↔ Attentive SVD++ model.

A user's preference on
different items of history are
not equal!

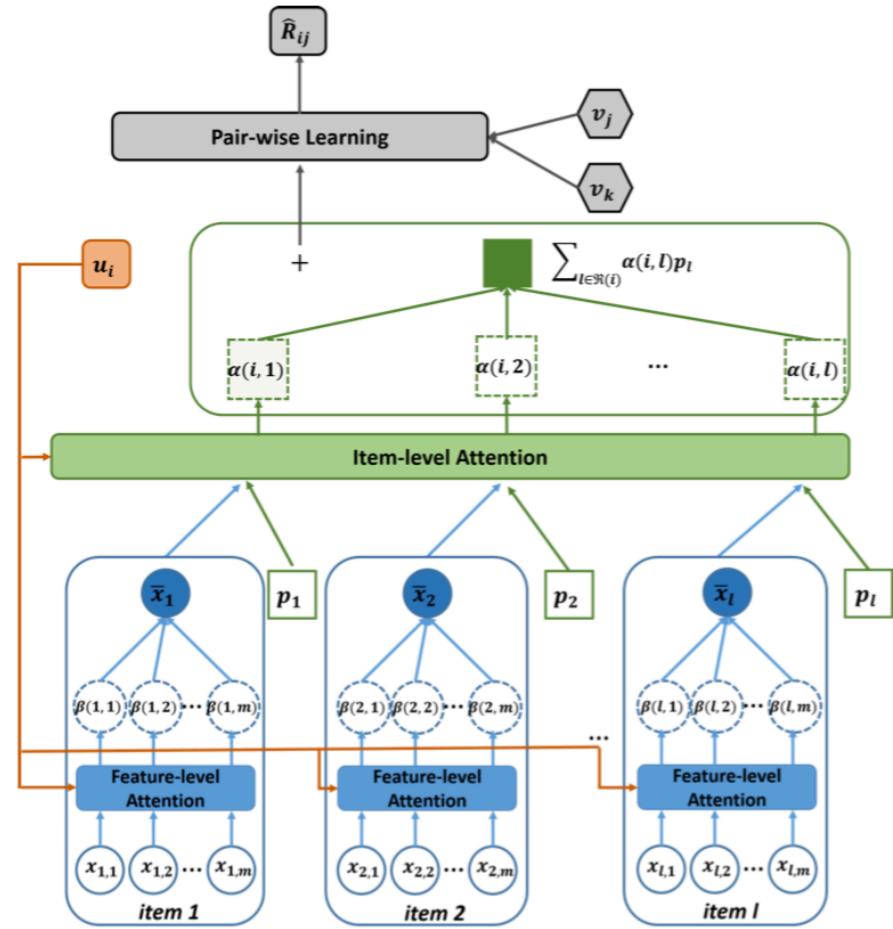
Attentive CF

$$\arg \min_{\mathbf{U}, \mathbf{V}, \mathbf{P}, \Theta} \sum_{(i, j, k) \in \mathcal{R}_B} -\ln \sigma \left\{ \left(\mathbf{u}_i + \sum_{l \in \mathcal{R}(i)} \alpha(i, l) \mathbf{p}_l \right)^T \mathbf{v}_j - \left(\mathbf{u}_i + \sum_{l \in \mathcal{R}(i)} \alpha(i, l) \mathbf{p}_l \right)^T \mathbf{v}_k \right\} + \lambda (\|\mathbf{U}\|^2 + \|\mathbf{V}\|^2 + \|\mathbf{P}\|^2),$$



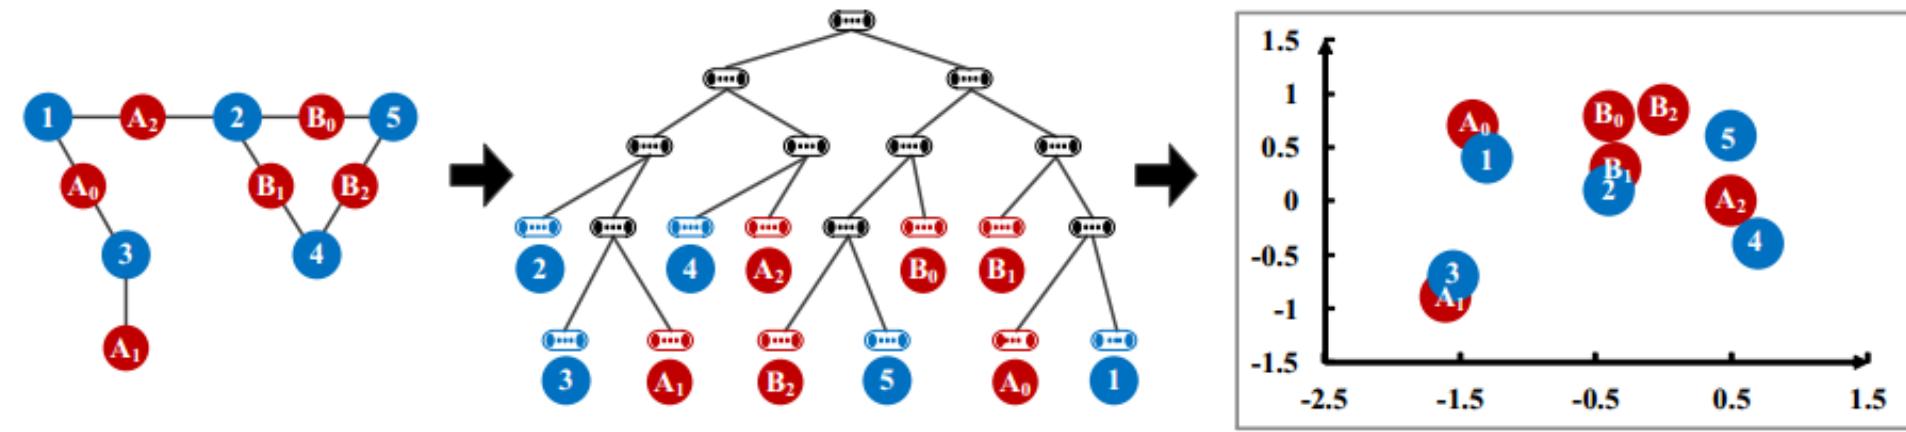
latent factor model

$$\hat{R}_{ij} = \underbrace{\mathbf{u}_i^T \mathbf{v}_j}_{\text{latent factor model}} + \underbrace{\sum_{l \in \mathcal{R}(i)} \alpha(i, l) \mathbf{p}_l^T \mathbf{v}_j}_{\text{neighborhood model}},$$



Towards “Semi-”: Deep Tree-CF

[Geng et al. ICCV’15]

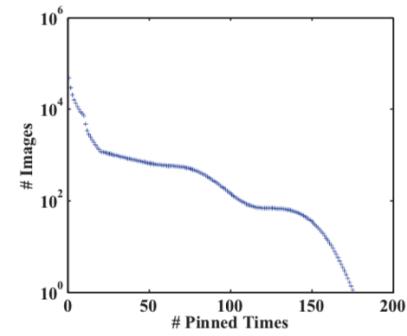
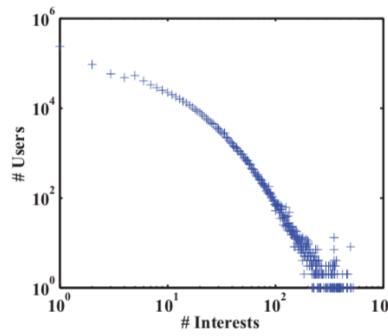


(a)

$$J = \frac{1}{2m} \sum_{ij} \left(\mathbf{G}_{ij} - \frac{d_i d_j}{2m} \right) \delta(i, j)$$

0-1 "likes" Avg "likes"

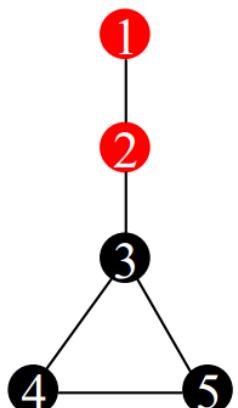
(b)



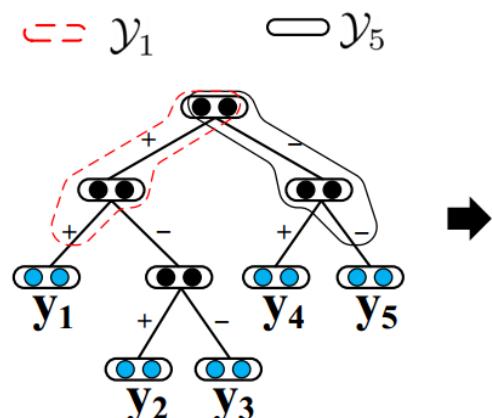
Pinterest: 100K users, 1M images

A Closer Look

[Zhang et al. TOMM'16]



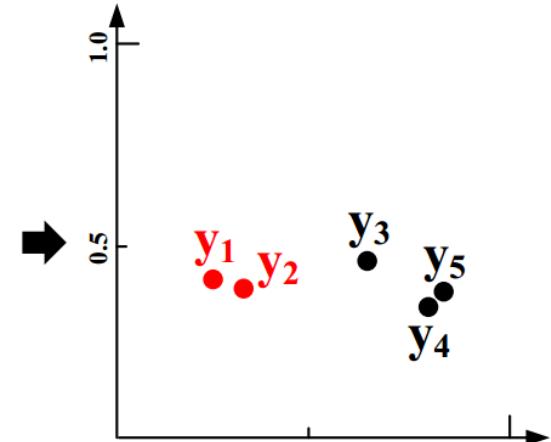
(a)



(b)

Model	Feature
\mathbf{y}_1	$\mathcal{Y}_1 \ \mathcal{Y}_2$
\mathbf{y}_2	$\mathcal{Y}_1 \ \mathcal{Y}_2 \ \mathcal{Y}_3$
\mathbf{y}_3	$\mathcal{Y}_2 \ \mathcal{Y}_3 \ \mathcal{Y}_4 \ \mathcal{Y}_5$
\mathbf{y}_4	$\mathcal{Y}_3 \ \mathcal{Y}_4 \ \mathcal{Y}_5$
\mathbf{y}_5	$\mathcal{Y}_3 \ \mathcal{Y}_4 \ \mathcal{Y}_5$

(c)

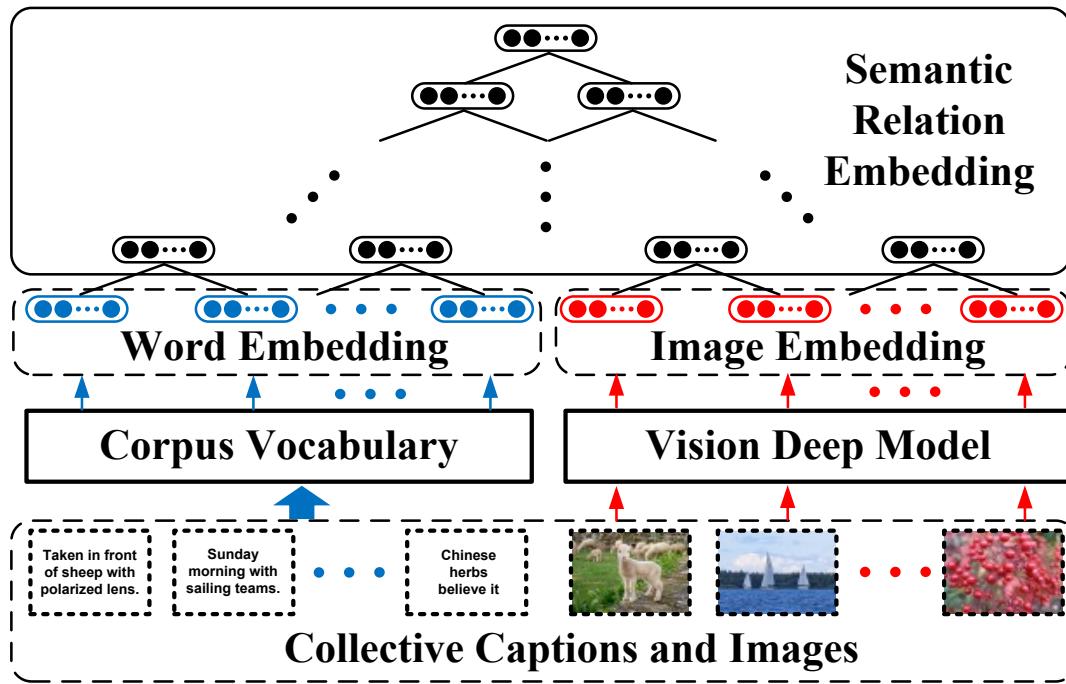


(d)

$$\min_{\mathbf{Y}} \sum_{ij} S_{ij} \log p(\mathbf{y}_i | \mathbf{y}_j)$$

Instead of $\langle \mathbf{y}_i, \mathbf{y}_j \rangle$, we use
 $\langle \mathbf{y}_i, \mathbf{Y}_j \rangle + \langle \mathbf{y}_i, \mathbf{Y}_i \rangle$

Visualization by Image and Word(User)



$$\max_{\mathbf{Y}} \sum_{ij} S_{ij} \log p(\mathbf{y}_i | \mathbf{y}_j)$$

$\mathbf{P}(\text{"dog"} | \text{dog})$



**Flickr 1M Image-Caption Pairs
8K+ Words**

Image-Word Embedding



Table 2: Performance (mAP%) of keyword-based retrieval on NUSWIDE and CCV.

Dataset/Method	S-CNN	R-Ours	HR-Ours	GoogLeNet
NUSWIDE	19.2	21.1	24.3	23.0
CCV	20.9	24.7	36.2	19.4

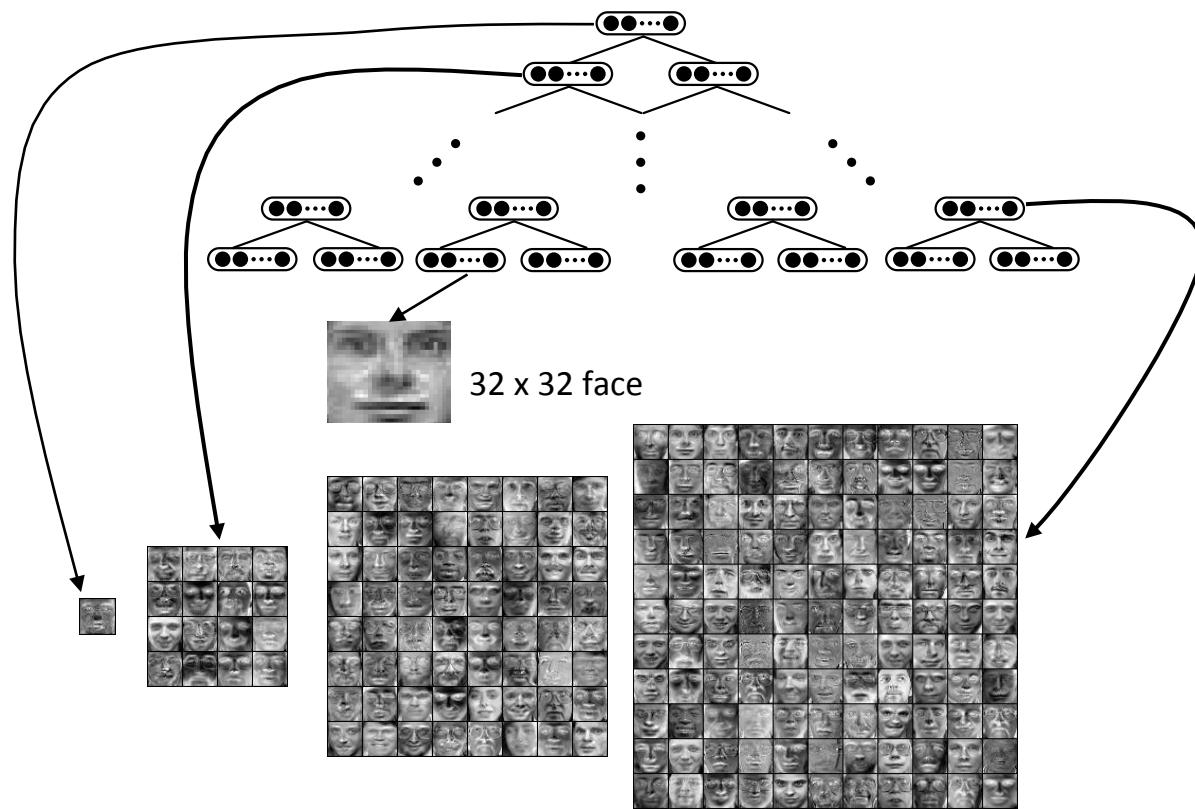
Word-Word Embedding

<i>Dog</i>	<i>foster, sleeping, service, cute, cat, puppy, stray, freshfields, spaniel, pug</i>	<i>Girl</i>	<i>dress, cute, hat, pink, boy, portrait, baby, shirt, birthday, blonde</i>
<i>Morning</i> + <i>Drink</i>	<i>drink, champagne, sky, juice, drinks, coffee, foggy, cold, soda, sun</i>	<i>Evening</i> + <i>Drink</i>	<i>evening, drink, wine, beer, bartender, alcohol, beverage, martini, liquor, bottle</i>

Table 1: Spearman correlation of 971 word-pair similarities computed by different methods and MEN human judgements. Our method learned from SBU is very close to Word2Vec learned from Google News.

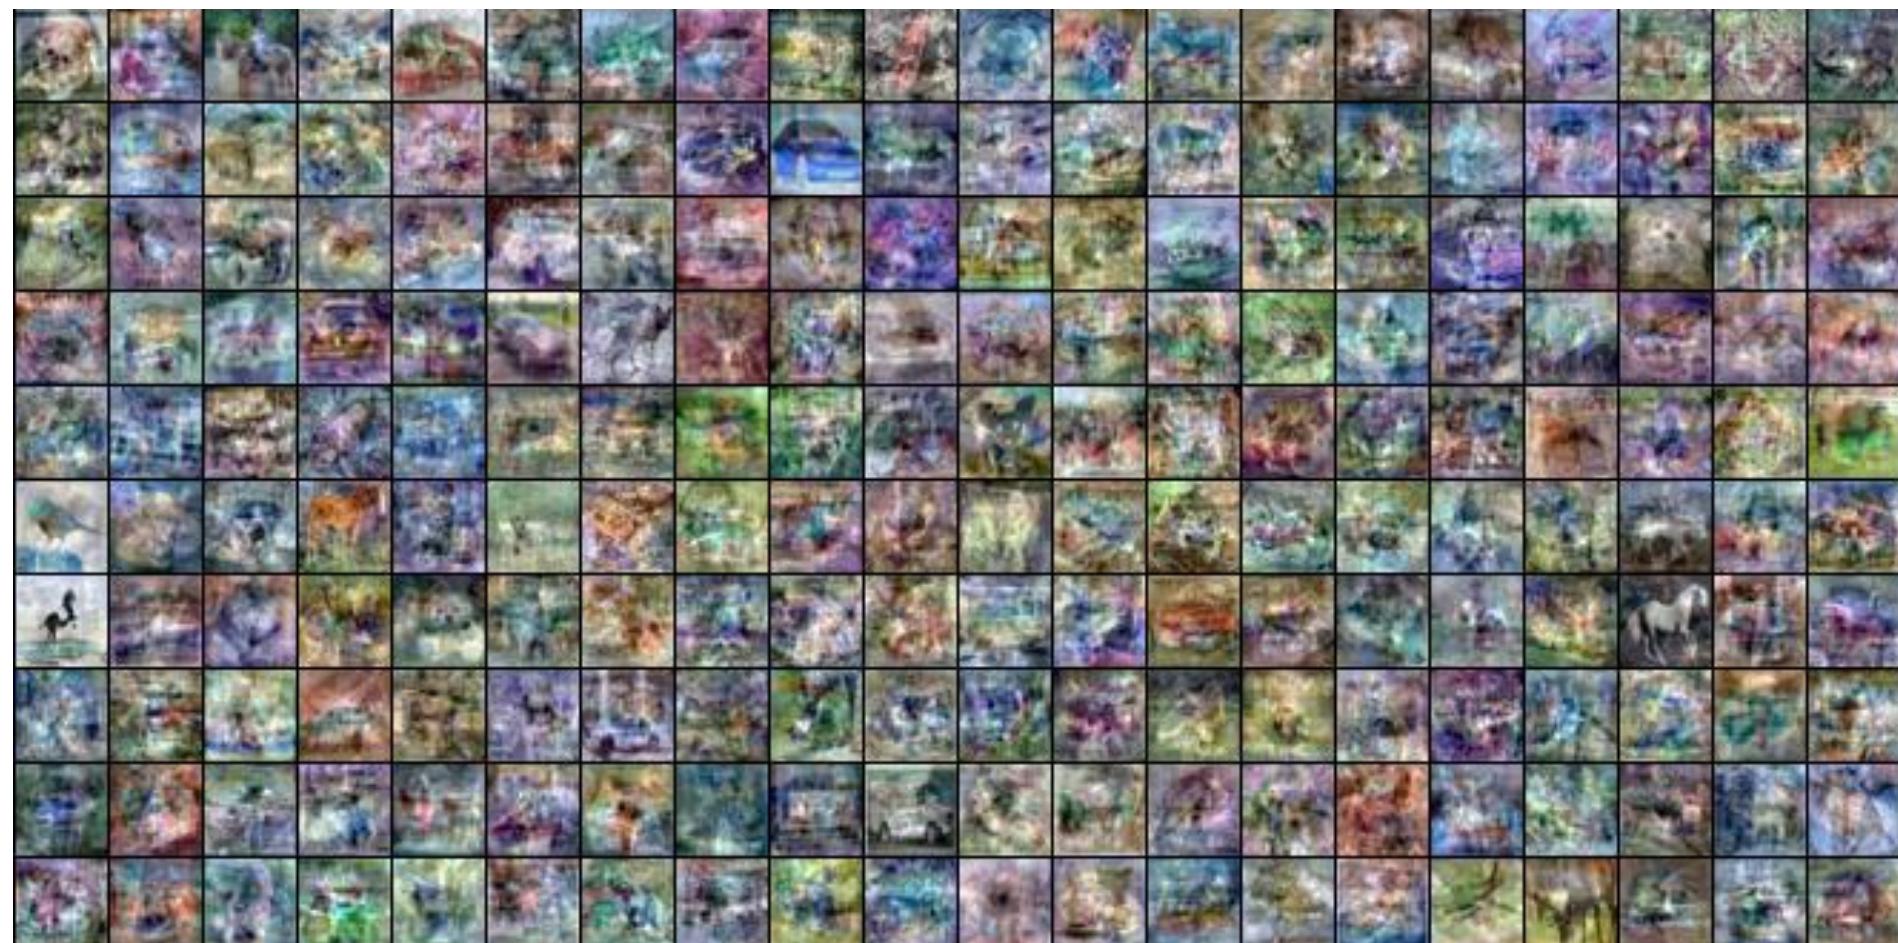
Method	Word2Vec	S-CNN	Ours
MEN	0.65	0.36	0.64

Model Visualization: Why it works?









Summary

- Many multimedia tasks benefit from GOOD visual features, so does Recommendation!
- Image: Use fine-grained and dynamic features (RoI+Visual Attention).
- Video: Better choice is to follow image, but 3d-conv is worth trying.
- Please follow ``weakly semi-supervised learning “ in ML for inspiration.

Outline of Tutorial

- Background (Xiangnan, 10 mins)
- Basics & Advances in Recommendation (Xiangnan, 50 mins)
- Visually-aware Product Recommendation (Xiangnan, 30 mins)
- Break (15 mins)
- Visual Representation (Hanwang, 45 mins)
- Image/Video Recommendation (Hanwang, 25 mins)
- Conclusion (Hanwang, 5 mins)

Conclusion

- Recommendation becomes increasingly important in our daily life
 - Age of Information overload
- Multimedia Recommendation is rich area of research.
 - ✓ Advanced recommendation technologies
 - ✓ Advanced visual representation learning
 - ✓ Representative visually-aware recommendation methods
- More research to be done in integrating state-of-the-art CV techniques into recommendation.

Challenges

- Model: data-driven + knowledge-driven
 - Most current methods are purely data-driven
 - Prior information (e.g., domain knowledge, symbolic knowledge) is helpful and should be integrated into data-driven learning in a principled way.
- Task: multiple criteria
 - Existing work have primarily focused on similarity (relevance)
 - Different scenarios may have different matching goals
 - Other criteria such as novelty, diversity, and explainability should be taken into consideration

Thanks!