

---

# Less is More: Improving LLM Alignment via Preference Data Selection

---

Xun Deng<sup>14</sup>, Han Zhong<sup>2\*</sup>, Rui Ai<sup>3</sup>, Fuli Feng<sup>5\*</sup>, Zheng Wang<sup>4</sup>, Xiangnan He<sup>5\*</sup>

<sup>1</sup>University of Science and Technology of China, <sup>2</sup>Peking University,

<sup>3</sup>Massachusetts Institute of Technology, <sup>4</sup>Alibaba Cloud Computing

<sup>5</sup>MoE Key Lab of BIPC, University of Science and Technology of China

dx981228@mail.ustc.edu.cn, hanzhong@stu.pku.edu.cn, ruihai@mit.edu  
{fulifeng93,xiangnanhe}@gmail.com, wz388779@alibaba-inc.com

## Abstract

Direct Preference Optimization (DPO) has emerged as a promising approach for aligning large language models with human preferences. While prior work mainly extends DPO from the aspect of the objective function, we instead improve DPO from the largely overlooked but critical aspect of data selection. Specifically, we address the issue of parameter shrinkage caused by noisy data by proposing a novel margin-maximization principle for dataset curation in DPO training. To further mitigate the noise in different reward models, we propose a Bayesian Aggregation approach that unifies multiple margin sources (external and implicit) into a single preference probability. Extensive experiments in diverse settings demonstrate the consistently high data efficiency of our approach. Remarkably, by using just 10% of the Ultrafeedback dataset, our approach achieves 3% to 8% improvements across various Llama, Mistral, and Qwen models on the AlpacaEval2 benchmark. Furthermore, our approach seamlessly extends to iterative DPO, yielding a roughly 3% improvement with 25% online data, revealing the high redundancy in this presumed high-quality data construction manner. These results highlight the potential of data selection strategies for advancing preference optimization.

## 1 Introduction

Reinforcement Learning from Human Feedback [RLHF; 6, 65] has emerged as a crucial technique for aligning Large Language Models (LLMs) with human preferences and values. Traditional RLHF implementations involve a two-stage process: reward model training based on preference data followed by reinforcement learning optimization. However, this approach presents significant computational challenges, requiring loading multiple model instances and extensive hyperparameter tuning.

As an alternative, [39] introduced Direct Preference Optimization (DPO), which streamlines the alignment process by directly optimizing the LLM policy from preference data. DPO has demonstrated comparable effectiveness while substantially reducing computational requirements compared to classical RLHF. Following DPO’s introduction, numerous studies have proposed improvements through modified learning objectives [62, 1, 15] and iterative learning schemes [53]. While these **algorithmic** advances have shown promise, there remains a critical gap in our understanding of the **data-centric** aspects of preference learning: *what characteristics of preference data contribute most to model alignment?*

This work thoroughly studies the impact of preference data quality on DPO training, which is crucial for developing more efficient training strategies. In particular, we achieve both *improved performance*

---

\*Corresponding author

and *reduced computational costs* through strategic data selection. Our research makes three primary contributions:

- (1) We prove in theory the necessity of data selection in the presence of exogenous noise. Specifically, the noise in the reward model may flip the preference between response pairs, leading to the emergence of the *parameter shrinkage* issue. Furthermore, we demonstrate that margin-based selection criteria can effectively address this issue by inducing *parameter inflation*.
- (2) Driven by the theoretical results and the derived margin-maximization principle, we propose a Bayesian Aggregation for Preference data Selection (**BeeS**) strategy. **BeeS** incorporates signals from both external rewards and DPO implicit rewards, and deprioritizes a preference pair if it exhibits a low reward margin from any single reward source to mitigate potential noise. Through extensive experiments across diverse datasets and models, we show that this selection strategy shows two consistent advantages: it substantially reduces computational overhead via efficient data selection and improves model performance compared to training on the full dataset. In particular, on the UltraFeedback dataset and its variants, our method identifies a 10% data subset for DPO training on LLama, Mistral, and Qwen series models, consistently achieving 3% to 8% point improvements on the AlpacaEval 2.0 benchmark relative to training on the complete dataset.
- (3) Finally, we extend our data selection framework to iterative DPO settings, showing that selectively sampling online data can simultaneously lower computational costs and improve performance. In particular, we achieve 48.49% win rate and 54.99% length-control win rate on the AlpacaEval 2.0 benchmark using only 25% of the online data for training.

Our findings provide both theoretical insights into the dynamics of preference learning and practical guidelines for more efficient DPO implementations. This work bridges an important gap between algorithmic innovation and data quality considerations in the context of LLM alignment.

## 1.1 Related Work

**Preference learning algorithms.** Reinforcement Learning from Human Feedback also known as dueling RL [37] or preference-based RL [5], has become a crucial component of recent Large Language Models (LLMs) such as ChatGPT [36]. While the classical RLHF pipeline traditionally uses Proximal Policy Optimization, several alternative approaches have been proposed. These include but not limited other RL-based training algorithms [28, 63], rejection sampling [11, 17], conditional supervised fine-tuning [31, 56, 60], and Direct Preference Optimization [39]. Among these alternatives, DPO has gained significant attention due to its simplicity and robust performance. Following the introduction of DPO, numerous works [62, 1, 15, 32, 44, 18, 54, 19, 51] have attempted to improve its performance by modifying the DPO objective.

**Data selection in LLM Fine-tuning.** Data selection is crucial in LLM post-training [48] for two key observations: post-training typically converges rapidly, and excessive data can degrade model performance through overfitting or exposure to toxic content [43, 10]. Recent research has focused on enhancing instruction tuning efficiency by identifying high-quality subsets from large instruction datasets [4, 25, 52], often adapting active learning query strategies [40] to assess sample uncertainty and diversity. However, data efficiency in preference learning remains relatively unexplored. Prior studies have studied reducing annotation costs in preference dataset creation [35, 57] and on scenarios involving numerous ranking annotations [45, 34]. Other research aims to improve a model’s ability to distinguish between two responses by adding a margin to the loss term [32, 38, 1]. Additionally, concurrent work highlights the importance of margins for preference data filtering, though there is debate on whether hard samples help or hinder preference learning [50, 22, 59, 16].

Our work firstly provides clear criteria for identifying informative samples while filtering toxic ones, thereby improving both DPO’s efficiency and performance. Furthermore, our method extends to iterative DPO [53] and its variants [61], wherein training data is dynamically generated by the model during its iterative training process.

## 2 Background

Reinforcement Learning from Human Feedback (RLHF) has emerged as a key method for aligning LLMs with human preferences. It leverages training data of the form  $\mathcal{D} = \{x, y_w, y_l\}$ , where  $x$  rep-

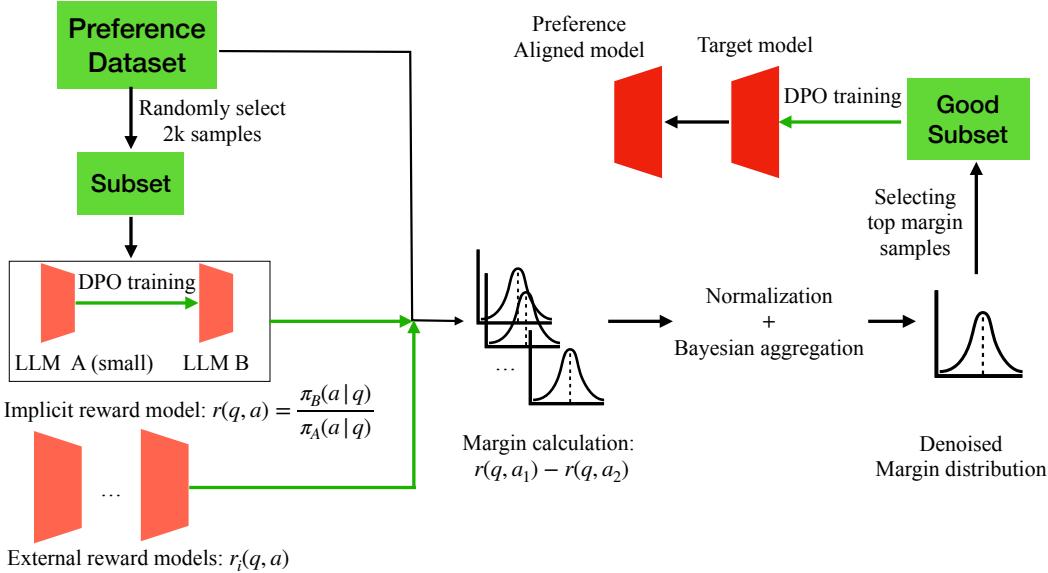


Figure 1: The workflow of the BeeS method.

resents the input prompt, and  $y_w$  and  $y_l$  denote the preferred and dispreferred responses, respectively. The RLHF pipeline typically involves two stages: reward learning and policy optimization.

**Reward Learning.** In the reward learning stage, a reward model is trained to approximate human preferences based on preference data. By adopting the Bradley-Terry model [3] to capture human preference, reward training involves minimizing the loss:

$$\mathcal{L}_{\text{RM}}(r) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} [\log \sigma(r(x, y_w) - r(x, y_l))],$$

where  $\sigma(\cdot)$  is the sigmoid function.

**Policy Optimization with Reinforcement Learning.** Once the reward model  $r$  is trained, it is used to guide the optimization of a policy  $\pi_\theta(y|x)$ , where  $\theta$  denotes the parameters of the model. This stage often employs reinforcement learning techniques such as Proximal Policy Optimization [PPO; 41] to optimize the policy by maximizing the expected reward.

$$\max_{\pi_\theta} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta(\cdot|x)} \left[ r(x, y) - \beta \log \frac{\pi_\theta(y|x)}{\pi_{\text{ref}}(y|x)} \right],$$

where  $\beta > 0$  is the regularization parameter. However, this RL approach can be computationally expensive, sensitive to reward misspecification and require careful hyperparameter tuning.

Recently, as an alternative to the RL-based policy optimization in RLHF, *Direct Preference Optimization* [DPO; 39] has been proposed. DPO simplifies the reward alignment process by directly incorporating human preference data into supervised training. Instead of defining and optimizing a reward function explicitly, DPO minimizes

$$\mathcal{L}_{\text{DPO}}(\theta) = -\log \sigma \left( \beta \log \frac{\pi_\theta(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi_\theta(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right).$$

By bypassing the intermediate step of reinforcement learning, DPO offers a more stable and computationally efficient alternative to standard RLHF, while still aligning models effectively with human feedback.

### 3 Methodology

#### 3.1 Parameter Shrinkage and Inflation Analysis

We follow the model from [64] to illustrate why data selection can improve model performance. We assume that reward model  $r(x, y) = \langle \phi(x, y), \omega^* \rangle$  with some feature function  $\phi(\cdot, \cdot)$ . For reward

Table 1: Symbols used in the formulation.

r	x	$y_w/y_l$	$\phi/\Delta\phi$	$\zeta$	w
Reward	Input prompt	Preferred/dispreferred Response	(relative) feature function	Exogenous error	Learnable parameters

learning, our reward model can be an explicit  $r(x, y)$  [36], while for DPO,  $\beta \log \frac{\pi_\theta(y|x)}{\pi_{\text{ref}}(y|x)}$  plays the role of reward model implicitly [39]. Based on observations in previous literature, we can derive such features by removing the last layer of the pre-trained model. However, both humans and other large models may use inaccurate reward functions to generate labels, where the labels represent the ranking of two responses. We say the preference between  $y_w$  and  $y_l$  is generated by  $r(x, y_w) - r(x, y_l) + \zeta$  where  $\zeta$  is an exogenous error. We use  $\Delta\phi(x)$  to denote  $\phi(x, y_w) - \phi(x, y_l)$  for simplicity.

*Parameter Shrinkage.* Here, we hope to find  $\omega$  to minimize

$$\mathcal{L}_{\text{RM}}(\omega) = -\mathbb{E}_{x, \zeta} \left[ \frac{1}{1 + e^{-\langle \Delta\phi(x), \omega^* \rangle - \zeta}} \log \left( \frac{1}{1 + e^{-\langle \Delta\phi(x), \omega \rangle}} \right) + \frac{1}{1 + e^{\langle \Delta\phi(x), \omega^* \rangle + \zeta}} \log \left( \frac{1}{1 + e^{\langle \Delta\phi(x), \omega \rangle}} \right) \right]. \quad (1)$$

It holds that the first-order condition is

$$\mathbb{E}_{x, \zeta} \left[ \frac{1}{1 + e^{\langle \Delta\phi(x), \omega^* \rangle + \zeta}} \frac{e^{\langle \Delta\phi(x), \omega \rangle} \Delta\phi(x)}{1 + e^{\langle \Delta\phi(x), \omega \rangle}} \right] = \mathbb{E}_{x, \zeta} \left[ \frac{1}{1 + e^{-\langle \Delta\phi(x), \omega^* \rangle - \zeta}} \frac{e^{-\langle \Delta\phi(x), \omega \rangle} \Delta\phi(x)}{1 + e^{-\langle \Delta\phi(x), \omega \rangle}} \right]. \quad (2)$$

Since we know that  $\langle \Delta\phi(x), \omega^* \rangle$  is positive, when  $\zeta$  is small comparing to the margin, it holds that  $\frac{1}{1 + e^{\langle \Delta\phi(x), \omega^* \rangle + \zeta}}$  is convex with respect to  $\zeta$ . Due to Jensen's inequality, it holds that

$$\mathbb{E}_{x, \zeta} \left[ \frac{1}{1 + e^{\langle \Delta\phi(x), \omega^* \rangle + \zeta}} \frac{e^{\langle \Delta\phi(x), \omega \rangle} \Delta\phi(x)}{1 + e^{\langle \Delta\phi(x), \omega \rangle}} \right] \geq \mathbb{E}_x \left[ \frac{1}{1 + e^{\langle \Delta\phi(x), \omega^* \rangle}} \frac{e^{\langle \Delta\phi(x), \omega \rangle} \Delta\phi(x)}{1 + e^{\langle \Delta\phi(x), \omega \rangle}} \right].$$

Similarly, we have

$$\mathbb{E}_{x, \zeta} \left[ \frac{1}{1 + e^{-\langle \Delta\phi(x), \omega^* \rangle - \zeta}} \frac{e^{-\langle \Delta\phi(x), \omega \rangle} \Delta\phi(x)}{1 + e^{-\langle \Delta\phi(x), \omega \rangle}} \right] \leq \mathbb{E}_x \left[ \frac{1}{1 + e^{-\langle \Delta\phi(x), \omega^* \rangle}} \frac{e^{-\langle \Delta\phi(x), \omega \rangle} \Delta\phi(x)}{1 + e^{-\langle \Delta\phi(x), \omega \rangle}} \right].$$

Since the optimal  $\omega$  is  $\omega^*$  without  $\zeta$ , plugging  $\omega^*$  in Equation (2) will cause the left-hand side to be greater than the right-hand side. Therefore, the optimal  $\omega$  with the existence of  $\zeta$  intends to shrink to the original point compared to  $\omega^*$  so that the first-order condition is still satisfied.

We provide the underlying intuition with an extreme example. If  $\nabla(\zeta)$  goes to infinity, the preference between  $y_w$  and  $y_l$  mainly depends on  $\zeta$ , approaching a Rademacher distribution, then  $\omega = 0$  could be a good solution to Equation (1). In other words,  $\zeta$  offsets part of the information provided by the reward model, causing the model's parameters to shrink toward zero. Thus, data selection is essential for acquiring policies with good performance. Finally, we remark that  $\zeta$  can come from multiple resources, including human classification errors, different embeddings or reward models from other LLMs and so on.

*Parameter Inflation.* We next explain why selecting data points based on the margin can lead to parameter inflation, thereby offsetting the parameter shrinkage caused by errors.

First, when the margin is large, namely,  $\langle \Delta\phi(x), \omega^* \rangle + \zeta$  is large, from the S-shaped graph of  $\sigma(\cdot)$ , we know that the slope is very small in this area. As a result, the probability of preference reversal caused by  $\zeta$  is low, which means the likelihood of incorrect samples is also low. Secondly, given prompt  $x$ , as we select data with large  $\langle \Delta\phi(x), \omega^* \rangle + \zeta$ , the posterior distribution of  $\zeta$  is skewed toward the positive side. Therefore, the preferences corresponding to this kind of data are more pronounced, leading to inflated estimates of  $\omega$  in Equation (1). Finally, we point out that if realized  $y_w$  and  $y_l$  are all separable, proportional scaling of  $\omega$  can reduce the value of Equation (1) continuously. Hence, some techniques like regularization or early stopping when training are indispensable.

In summary, inaccuracies in the reward model can cause the parameters of LLMs to shrink toward zero. By selecting data with larger margins, we can compensate for the performance degradation caused by this shrinkage. The balance between parameter shrinkage and inflation offers the potential to enhance the performance of LLMs. Driven by this theoretical result, our main idea is to let the model reach an overall high margin during preference learning. We realize this by providing a robust estimation of reward margins for the entire dataset ahead of full-set training, which then allows efficient high-margin data filtering.

### 3.2 Multi-source Margin Aggregation

Building on our previous analysis, we aim to develop a data selection strategy based on the margin-maximization principle, with the calculation of reward margin being the critical component. We

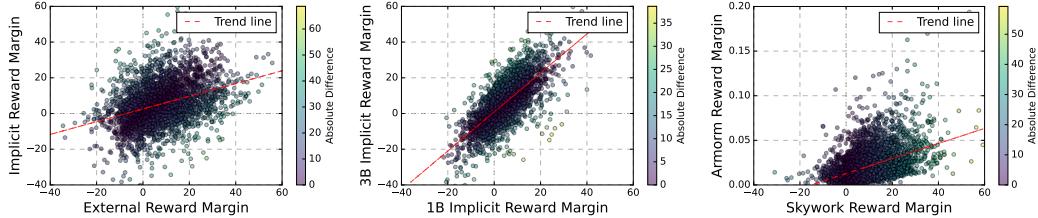


Figure 2: Visualization of joint margin distribution on **UltraFeedback**. (Left) Joint distribution of external and implicit reward margin values. (Middle) Joint distribution of implicit reward margins computed using models of 1B and 3B scales. (Right) Joint distribution of two different external reward margin values on online-generated data.

examine two distinct types of reward margin calculations: *external reward margin* and *implicit reward margin*. The external reward margin is provided by an external reward function, while the implicit reward margin is derived from the implicit reward  $\log \frac{\pi_\theta}{\pi_{\text{ref}}}$ , where  $\pi_\theta$  represents the policy trained by DPO [39] (see Section 4.1 for details on margin calculation). We visualize the joint distribution of different reward margin sources using the UltraFeedback dataset [7] in Figure 2.

- The left and middle panels of Figure 2 reveal several key phenomena: (1) The correlation between implicit and external reward margins is notably weak. In particular, samples exhibiting high positive implicit margins span a broad range of external margins, including strongly negative values, and vice versa. This divergence highlights the distinct preference patterns captured by these two reward types. This underscores the need to combine both reward types for a reliable margin estimation. (2) In contrast, we observe a strong correlation between implicit reward margins calculated by models of different sizes (Llama-3.2 3B and 1B). Notably, these two patterns remain consistent across other datasets as well (see Appendix B.2).
- Online RLHF [53, 13] employs the target model to generate multiple responses for given prompts iteratively and uses an external reward model to identify the response pair with the largest margin for DPO training. The right panel of Figure 2 illustrates that a max-margin pair construction method, even when derived from one strong reward model, can still yield ambiguous preferences when evaluated by another reward model (which shows similar performance on RewardBench Leadboard). This ambiguity in preference signal indicates that the online data generation process can still cause high redundancy, which may offer little to no benefit, or could even be detrimental, to online-DPO training.

These observations highlight the need for multi-source margin aggregation to achieve a more robust margin estimation, thereby enhancing data selection and preference learning. To this end, we propose a strict aggregation strategy, Bayesian Aggregation for Preference data Selection (**BeeS**), that deproritizes a preference pair if it exhibits a low reward margin from any single reward source. We implement this method through a general three-step procedure.

*Step 1: in-distribution pre-DPO training.* Our objective is to obtain the in-distribution implicit reward model with low computational and GPU-memory cost. Given the strong correlation of that margin across different models (See Figure 2), this weak-to-strong guidance is feasible. To achieve this, we randomly select a small seed dataset  $\mathcal{D}_0$  from  $\mathcal{D}$  and employ a (or several) small model to perform preference learning on this seed set. The high sample and training efficiency of the DPO loss [23] ensures the feasibility of this approach.

*Step 2: margin calculation.* We calculate external and implicit reward margins as  $m_{\text{ex}} = r_{\text{ex}}(x^i, y_w^i) - r_{\text{ex}}(x^i, y_l^i)$  and  $m_{\text{im}} = r_{\text{im}}(x^i, y_w^i) - r_{\text{im}}(x^i, y_l^i)$  for each datum in  $\mathcal{D}$ , where we directly calculate  $r_{\text{im}}(x, y)$  as  $\log \frac{\pi_\theta(y|x)}{\pi_{\text{ref}}(y|x)}$ . Here,  $\pi_{\text{ref}}$  and  $\pi_\theta$  denote the small model before and after preference learning. We assume that there are  $K$  margins involved from these two typical sources.

*Step 3: Bayesian aggregation.* To mitigate noise from individual reward sources, we utilize Bayesian probability theory for their robust aggregation. Given that reward margins from different sources often vary in their underlying distributions and value ranges, we propose projecting these diverse margins into a unified probability space. This projection serves to quantify the confidence that a specific preference direction  $y_w > y_l$  is correct, formally expressed as  $\mathbb{P}(y_w > y_l | m^1, m^2, \dots, m^K)$ . Assuming that these sources are conditionally independent, then following the Bayesian formula

transformation in previous work [30, 9], the preference probability can be expressed as:

$$\mathbb{P}(y_w > y_l | m^1, m^2, \dots, m^K) = \frac{\prod_{i=1}^K \mathbb{P}(y_w > y_l | m^i)}{\prod_{i=1}^K \mathbb{P}(y_w > y_l | m^i) + \prod_{i=1}^K (1 - \mathbb{P}(y_w > y_l | m^i))}. \quad (3)$$

The typical absence of well-defined preference labels (e.g.,  $y_w > y_l$  for clear preference, or  $y_w = y_l$  for indifference) renders the rigorous estimation of single-source preference probabilities challenging.<sup>2</sup> Consequently, we approximate the probability using a linear projection:  $\mathbb{P}(y_w > y_l | m^i) = \frac{\text{clip}(m^i, L, U) - L}{U - L}$ , where  $\text{clip}(m, L, U) = \min(\max(m, L), U)$ ) and  $L, U$  are tuning parameters. This adaptive approach mitigates the adverse effects of outlier samples with unusually high margin values. See more implementation details and discussion about the derivation approximation of Eq. (3) in Appendix A.1.

**Sample selection.** We consider the data selection for both the one-pass DPO training and the iterative DPO workflow. For the former, we directly select the samples with the highest aggregated preference probability to construct  $\mathcal{D}_{\text{train}}$ . The threshold depends on how many preference samples we prefer to use (but should guarantee that samples with negative margins are excluded). For the latter, we only need to train implicit reward models in the first iteration, and **BeeS** three-step procedures are applied for each iteration.

## 4 Experiments

We organize the experiments as follows: we explain the experimental setup in Section 4.1; we compare **BeeS** with various sample selection baselines on diverse preference tasks and present the detailed results in Section 4.2; then we focus on the important chat task, and explore the effectiveness of **BeeS** in enhancing comprehensive dialogue ability in Section 4.3. Lastly, we perform diverse ablation studies for the **BeeS** in Section 4.4.

### 4.1 Experimental Setup

**Preference Datasets.** We evaluate our approach using three established preference datasets: (1) Reddit **TL;DR** summarization dataset [47, 42] that contains human-written summaries and human-rated results, (2) Anthropic Helpful and Harmless dialogue dataset (**HH**) [2], and (3) **UltraFeedback** [7], which comprises quality-scored model responses across diverse prompts from multiple sources. To explore how models react to on-policy data, we leverage two modified versions of the **UltraFeedback** dataset, **Llama-UltraFeedback** and **Mistral-UltraFeedback** [32]. In the variants, the original chosen and rejected responses are replaced with the highest and lowest scored responses, respectively, sampled from five candidates generated by the corresponding Instruct model. The scores are given by the PairRM [21] reward model. Statistics about these datasets are in Appendix A.

**Models.** Our experiments are conducted across four model series: Llama-3.2 [33], Llama-3 [14], Mistral-7B-v2 [20], and Qwen-2.5 [55] under Base and Instruct setups. For the Base model (Llama-3.2-3B and Llama-3-8B), we first establish fundamental Instruction-following capabilities through supervised fine-tuning on the RLHF/SFT-OpenHermes-2.5-Standard datasets. For the Instruct setup, we directly use them as the start of DPO training. Regarding the external reward model, we adopt the recent Skywork-Reward-Llama-3.1-8B-v0.2 [29] that is the best reward model at this scale according to the RewardBench leaderboard. As for the implicit reward model, we employ the Llama-3.2-3B Base and its DPO-tuned model (on 2,000 randomly selected samples from the complete dataset) for  $\pi_{\text{ref}}$  and  $\pi_{\theta}$ .

**Implementation and Evaluation.** For DPO training, we follow [39] and use a fixed value of  $\beta = 0.1$ , except for **TL;DR** where  $\beta = 0.5$ . We run each training for two epochs, with a learning rate of  $5 \times 10^{-7}$ , and a 0.1 warmup ratio. Following [39], we evaluate the models using 400 randomly sampled test sets from the validation/test pools of the **TL;DR** and **HH** datasets, separately. For models trained on **UltraFeedback**, we employ AlpacaEval and AlpacaEval 2.0 [27] as our evaluation benchmark, which consists of 805 diverse questions.<sup>3</sup> As the ground truth oracle is unavailable,

<sup>2</sup>This estimation typically relies on methods such as isotonic regression or histogram analysis.

<sup>3</sup>They use the same set of questions, and differ in their reference response generation: AlpacaEval uses Text-Davinci-003 [58], whereas AlpacaEval 2.0 employs GPT4-1106-preview

Table 2: GPT-4 judged win rates for Llama-3.2-3B models fine-tuned with DPO on subsets selected by various data selection strategies. For every strategy and benchmark (TL;DR, HH, UltraFeedback (UF)), 2,000 preference samples were selected. Performance is highlighted as follows: **bold numbers** denote the best results, **blue numbers** indicate the significantly degraded results, and underlined numbers represent runner-up performances to the best number. P, Z, and N denote the most positive, near-zero, and most negative selection principles.

Strategy Region	Rand	External Margin			Implicit Margin			IFD Margin			BeeS	Fullset
		P	Z	N	P	Z	N	P	Z	N	P	
<b>TL;DR</b>	46.50	<u>66.25</u>	42.00	<b>22.00</b>	30.75	43.00	<b>19.75</b>	<b>1.75</b>	41.25	55.25	<b>83.25</b>	36.75
<b>HH</b>	84.25	82.25	76.50	69.75	<b>92.25</b>	81.25	<b>32.00</b>	<b>11.25</b>	<u>90.00</u>	64.50	<u>90.25</u>	92.00
<b>UF</b>	82.86	<u>91.18</u>	73.29	<b>25.84</b>	<u>89.81</u>	77.14	<b>37.02</b>	72.05	83.60	<b>54.53</b>	<b>91.68</b>	80.99

we employ GPT-4 as a proxy for human judgment across three distinct settings: summarization, helpful or harmless completion, and single-turn dialogue. We utilized a fixed decoding temperature ( $T = 0.7$ ) for all model generation in the experiments. More details are presented in Appendix A.

**Baselines.** For the offline data selection setting, we compare our method with three types of methods: (1) **Random**, a simple yet effective strategy in many domains (e.g., Instruction tuning [52]), (2) **IFD** [26] (i.e., exponential form of the Point-wise Mutual Information), which measures semantic overlap. We use the difference in IFD scores between chosen and rejected responses for preference data selection. (3) **External/Implicit Margin (M-Ex/Im)** computes the gap between chosen and rejected responses using either external reward models or implicit DPO rewards. For (2) and (3), we segment the data into **P** (most positive pairs), **Z** (close to zero pairs), and **N** (most negative pairs) subsets according to margin values. Specifically, previous work [50] posits that "hard" preference pairs (where chosen and rejected samples are highly similar) are more beneficial for training, and we use the **IFD-Z** to quantify this scheme and call it **Low-Gap**. For the iterative DPO setting, we compare our approach against the standard online iterative DPO baseline established by [53, 12] and run for three rounds, each using 20k prompts sampled from **UltraFeedback**. We provide source code of our paper in <https://github.com/xiantanshi/DPO-Data-Selection>.

## 4.2 Win Rate Comparison with Baselines on Classic Preference Datasets

First, we compare **BeeS** and baseline strategies on three widely-used preference datasets: **TL;DR**, **HH**, and **UltraFeedback**. Using a Llama-3.2-3B model as our Base architecture, we evaluate different selection methods, each sampling 2,000 training examples for DPO training. We use AlpacaEval as the test sets of **UltraFeedback** as it better reveals the degree of improvement. The results, measured by GPT4 win rates, are presented in Table 2. We summarize the findings below.

- Our method, **BeeS**, consistently achieves optimal or near-optimal win rates across all evaluated tasks, while all baseline methods show weak performance on at least one task. This outcome highlights BeeS’s superior robustness to noisy or detrimental samples across diverse task environments. Further, more data in DPO training does not always yield better results. Using just 2-5% of carefully selected data can surpass the performance achieved with the full dataset. Insights from Table 2 also reveal the existence of toxic samples and potential pitfalls of certain selection strategies. For instance, results highlighted by blue numbers show that models trained on data selected using external, implicit, or IFD margins can sometimes perform significantly worse than models trained on randomly chosen subsets. Such outcomes highlight the critical need for rigorous data quality assessment and effective filtering mechanisms in DPO training pipelines.
- Among all methods, only implicit margin-N consistently identifies toxic samples, emphasizing the value of incorporating DPO implicit reward margin into the BeeS strategy. Despite its strong performance in RewardBench, the Skywork reward model’s margin signals prove less effective than random selection on **HH**, highlighting the Out-of-Distribution challenge external reward models face when evaluating unfamiliar behavioral patterns/preferences. As for the IFD margin metric, it exhibits notable inconsistency across different datasets, rendering it unreliable for evaluating new datasets where prior preference patterns are unknown. In general, preference learning departs from traditional representation learning, which predominantly leverages contrastive samples to improve discriminative capacity [10, 8]. Preference learning focuses on capturing

Table 3: Performance comparison on AlpacaEval 2.0 using DPO-trained models with different 6,000-sample subsets (10% of full set). Both SFT and Instruct variants of Llama-3-8B were evaluated. LC and WR denote length-controlled and raw win rate, respectively. **Bold** number denotes the best-performing selected subset. **Blue numbers** denote results that show little advantage over random.

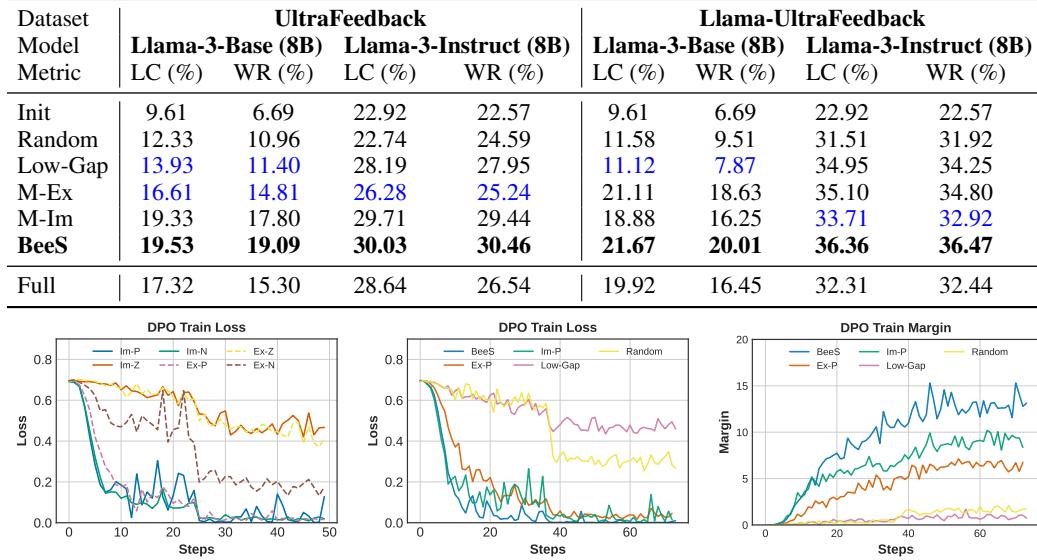


Figure 3: DPO training loss and margin of Llama-3.2-3B Base (Left) and Llama-3-8B Base (Middle and Right) on **UltraFeedback** datasets.

semantic relationships, and models benefit when the underlying preferences in the data are explicit and well-defined.

### 4.3 AlpacaEval 2.0 Win Rate Comparison

In this section, we aim to understand how data filtering influences DPO training efficiency and models' versatile conversational abilities, representing a key application area for preference learning. We use both Llama-3-8B (Base) and (Instruct) models, measuring performance through raw and length-controlled win rates on AlpacaEval 2.0, and results are shown in Table 3, Figures 3 and 4.

**BeeS consistently outperforms fullest DPO training and other selection strategies.** As shown in Table 3, BeeS-selected subsets achieve around 4% higher win rates compared to full dataset training across all four settings. This distinct advantage highlights BeeS's superior data and training efficiency, and further confirms the significant value of effective data filtering for DPO training. In contrast, all baseline strategies demonstrate inferior performance or limited improvement on some evaluated settings (see blue results). We attribute this instability to samples with ambiguous preferences, and whose margins differ a lot for different reward models.

**Different training dynamics of 'P'/N'/Z' subset region.** The left panel of Figure 3 shows DPO training loss curves for subsets selected by various strategies. Notably, training on subsets filtered according to the 'P' and 'N' criteria results in a rapid decrease in loss. In contrast, the loss curves corresponding to the 'Z' criterion tend to stabilize at consistently higher plateaus. Notably, 'N'-selected samples, which are often assumed as "difficult-to-learn" [16], can actually be learned as rapidly as 'P'-selected samples, suggesting that **'bad' preferences are also easy to grasp for LLM**. The pattern is consistent across different datasets and models (see Appendix C.1 for more results). While this observation might explain proposals that use absolute margin values for selection [35], Table 2 reveals that 'P' and 'N' samples produce opposing effects despite similar training dynamics.

The middle and right panels of Figure 3 illustrate that data subsets selected by BeeS exhibit both the most rapid decrease in training loss and the fastest increase in the DPO training margin, i.e., current train-batch average implicit margin. These concurrent observations of accelerated optimization help to explain the superior performance achieved by BeeS.

**Extension to Iterative DPO.** We explore the data efficiency of iterative DPO using prompts from UltraFeedback as in [53]. In comparison, 20k prompts are used for on-policy preference pair

Method	Llama-3-Base (8B)			Llama-3-Instruct (8B)		
	LC (%)	WR (%)	Len	LC (%)	WR (%)	Len
DPO (r1)	17.64	13.36	1496	40.51	43.90	2173
DPO (r2)	23.06	22.45	1897	42.51	49.23	2366
DPO (r3)	29.03	30.86	2736	44.51	53.12	2860
DPO-BeeS (r1)	16.35	13.09	1624	42.20	45.74	2091
DPO-BeeS (r2)	23.79	24.17	1901	46.40	50.60	2316
DPO-BeeS (r3)	<b>32.31</b>	<b>33.91</b>	2565	<b>48.49</b>	<b>54.99</b>	2774

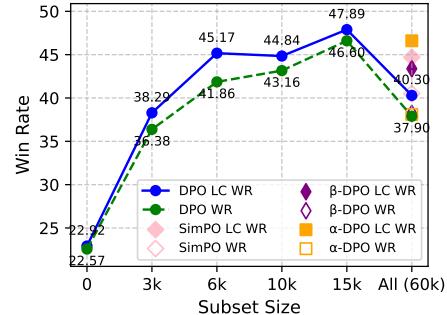


Figure 4: AlpacaEval 2.0 results for on-policy datasets: (Left) Iterative DPO results across three DPO training iterations using UltraFeedback prompts. (Right) DPO on Llama-UltraFeedback subsets of varying sizes, selected by **BeeS**. Results of DPO-variants trained on fullset are also compared.

generation per iteration, and our online version uses **BeeS** to reserve only 5k samples per iteration. The results are in the left panel of Figure 4.

**There is high redundancy in the on-policy data construction manner.** Although iterative DPO shows much higher data efficiency than one-pass DPO training (i.e., better results than those in Table 3), data selection is still important for quality control. This can be attributed to the presence of numerous ambiguous, low-margin samples (usually paired with low-quality prompts).

**A smaller  $\beta$  value in DPO loss correlates with higher data efficiency.** While  $\beta$  is commonly recognized as a factor controlling the strength of the Kullback-Leibler (KL) divergence, it also significantly influences data efficiency. Specifically, the DPO loss, defined as  $\log \sigma(\beta \times m_{im})$ , indicates that a reduced  $\beta$  allows for effective gradient updates for more preference pairs with wider margins. To investigate this, we conducted DPO training on the Llama-3-Instruct 8B model using its on-policy dataset, Llama-UltraFeedback, with  $\beta$  set to 0.01. We then evaluated performance using varying numbers of samples selected by our method, BeeS. The results, presented in the right panel of Figure 4, demonstrate that relaxing the margin constraint in the DPO loss substantially improves data efficiency (Refer to Appendix C.4 for results on the Base setup). Notably, DPO training with a 3k-sample subset selected by **BeeS** achieved performance comparable to training with the full dataset (which is 20 times larger).

Furthermore, we compared **BeeS** data selection with several established DPO variants that modify the original loss function, including SimPO [32],  $\beta$ -DPO [50], and  $\alpha$ -DPO [49]. Our findings indicate that **BeeS** is unique in its ability to effectively enhance both the win rate and the length-controlled (LC) win rate. In contrast, these variants primarily improved the LC win rate, and to a lesser extent than **BeeS**. These results underscore the significant potential of data selection and data efficiency considerations to enhance the original DPO training algorithm.

#### 4.4 Ablation Study

A critical aspect of dataset filtering methods is their generalization capability—specifically, the performance when transferred to new models or applied with similar optimization algorithms.

**Data filtering remains effective for new LLM architecture.** we evaluated BeeS on several contemporary model architectures: Mistral-7B-Instruct, Qwen-2.5-7B-Instruct, and Qwen-2.5-14B-Instruct. Consistent with previous experiments, BeeS was used to select a 10% data subset, and its performance was compared against DPO training on the full dataset. The AlpacaEval 2.0 evaluation results are presented in the left panel of Figure 5. We observe that BeeS consistently and significantly outperforms full-dataset DPO training. Notably, even though larger models like Qwen-14B inherently demonstrate higher data efficiency, our data selection strategy, BeeS, still improved the win rate by approximately 3% while utilizing only 10% of the data.

**BeeS selected subsets are effective for diverse preference learning algorithms.** We examine whether the subset selected by **BeeS** remains data efficient for DPO-variants like IPO [1], KTO [15], and SLiC [62]. We utilize the **BeeS** selected 6k-sample subset from Llama-UltraFeedback and the results are presented in the right panel of Figure 5. We observe that the high-margin subset consistently benefit these preference learning algorithms by outperforming full-set training.

	Mistral-7B		Qwen-2.5-7B		Qwen-2.5-14B	
	LC (%)	WR (%)	LC (%)	WR (%)	LC (%)	WR (%)
Initial	17.11	14.72	31.27	31.60	37.03	32.64
Full	18.00	18.77	39.67	38.24	49.99	46.81
<b>BeeS</b>	<b>26.04</b>	<b>20.53</b>	<b>46.20</b>	<b>43.78</b>	<b>50.20</b>	<b>49.74</b>

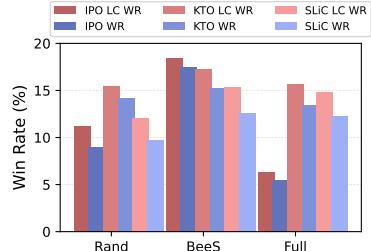


Figure 5: Ablation Study: (Left) different model choices (**Mistral-7B-Instruct-v-0.2**, **Qwen-2.5-Instruct-7B** and **Qwen-2.5-Instruct-14B**). **BeeS** selects a 6k-sample subset for training. (Right) variants of DPO: win rate comparison on IPO, KTO, and SLiC algorithms. **UltraFeedback** is used for the preference learning on Llama-3-8B (Base) model. Rand and **BeeS** select a 6k-sample subset.

Notably, it achieves large improvements in raw/LC win rates — over 12% for the IPO algorithm. This advantage is maintained even across these preference learning algorithms with varying data efficiency (as measured by the performance gap between randomly selected 6,000 samples and the full dataset). These findings highlight the significant value of sample filtering for other preference learning. Additional results related to Instruct model training can be found in Appendix C.3.

## 5 Conclusion

Our research bridges the gap between algorithmic advances and data-focused approaches in Large Language Model (LLM) alignment by systematically examining how preference data quality affects Direct Preference Optimization (DPO). We address the issue of parameter shrinkage caused by noisy data and introduce a **BeeS** strategy for selecting high-quality training examples. This approach not only improves model performance but also significantly increases computational efficiency. Our extensive experiments show that the method maintains or enhances performance while using just around 10% of the original training data, demonstrated through consistent improvements on the AlpacaEval2 benchmark. Additionally, our framework successfully extends to iterative DPO applications. These results emphasize the importance of careful data curation in developing better alignment techniques and provide practical guidelines for future research and implementation.

## Acknowledgements

This work is supported by the National Natural Science Foundation of China (U24B20180) and the Alibaba Group through Alibaba Research Intern Program. We appreciate the reviewers for their insightful feedback and advice, these constructive criticism and recommendations have been invaluable in helping us improve the quality of this work.

## References

- [1] M. G. Azar, Z. D. Guo, B. Piot, R. Munos, M. Rowland, M. Valko, and D. Calandriello. A general theoretical paradigm to understand learning from human preferences. In *International Conference on Artificial Intelligence and Statistics*, pages 4447–4455. PMLR, 2024.
- [2] Y. Bai, A. Jones, K. Ndousse, A. Askell, A. Chen, N. DasSarma, D. Drain, S. Fort, D. Ganguli, T. Henighan, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.
- [3] R. A. Bradley and M. E. Terry. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952.
- [4] Y. Cao, Y. Kang, C. Wang, and L. Sun. Instruction mining: Instruction data selection for tuning large language models. *COLM*, 2024.
- [5] X. Chen, H. Zhong, Z. Yang, Z. Wang, and L. Wang. Human-in-the-loop: Provably efficient preference-based reinforcement learning with general function approximation. In *International Conference on Machine Learning*, pages 3773–3793. PMLR, 2022.
- [6] P. F. Christiano, J. Leike, T. Brown, M. Martic, S. Legg, and D. Amodei. Deep reinforcement learning from human preferences. *Advances in neural information processing systems*, 30, 2017.
- [7] G. Cui, L. Yuan, N. Ding, G. Yao, W. Zhu, Y. Ni, G. Xie, Z. Liu, and M. Sun. Ultrafeedback: Boosting language models with high-quality feedback. *ICML*, 2024.
- [8] X. Deng, F. Feng, X. Wang, X. He, H. Zhang, and T.-S. Chua. Learning to double-check model prediction from a causal perspective. *IEEE Transactions on Neural Networks and Learning Systems*, 35(4):5054–5063, 2023.
- [9] X. Deng, J. Liu, H. Zhong, F. Feng, C. Shen, X. He, J. Ye, and Z. Wang. A3s: A general active clustering method with pairwise constraints. *arXiv preprint arXiv:2407.10196*, 2024.
- [10] X. Deng, W. Wang, F. Feng, H. Zhang, X. He, and Y. Liao. Counterfactual active learning for out-of-distribution generalization. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11362–11377, 2023.
- [11] H. Dong, W. Xiong, D. Goyal, Y. Zhang, W. Chow, R. Pan, S. Diao, J. Zhang, K. Shum, and T. Zhang. Raft: Reward ranked finetuning for generative foundation model alignment. *arXiv preprint arXiv:2304.06767*, 2023.
- [12] H. Dong, W. Xiong, B. Pang, H. Wang, H. Zhao, Y. Zhou, N. Jiang, D. Sahoo, C. Xiong, and T. Zhang. Rlhf workflow: From reward modeling to online rlhf. *arXiv preprint arXiv:2405.07863*, 2024.
- [13] Q. Dong, L. Dong, X. Zhang, Z. Sui, and F. Wei. Self-boosting large language models with synthetic preference data. *ICLR*, 2025.
- [14] A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Yang, A. Fan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- [15] K. Ethayarajh, W. Xu, N. Muennighoff, D. Jurafsky, and D. Kiela. Kto: Model alignment as prospect theoretic optimization. *arXiv preprint arXiv:2402.01306*, 2024.
- [16] C. Gao, H. Li, L. Liu, Z. Xie, P. Zhao, and Z. Xu. Principled data selection for alignment: The hidden risks of difficult examples. *arXiv preprint arXiv:2502.09650*, 2025.
- [17] C. Gulcehre, T. L. Paine, S. Srinivasan, K. Konyushkova, L. Weerts, A. Sharma, A. Siddhant, A. Ahern, M. Wang, C. Gu, et al. Reinforced self-training (rest) for language modeling. *arXiv preprint arXiv:2308.08998*, 2023.
- [18] J. Han, M. Jiang, Y. Song, J. Leskovec, S. Ermon, and M. Xu.  $f$ -po: Generalizing preference optimization with  $f$ -divergence minimization. *arXiv preprint arXiv:2410.21662*, 2024.

- [19] J. Hong, N. Lee, and J. Thorne. Orpo: Monolithic preference optimization without reference model. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 11170–11189, 2024.
- [20] A. Q. Jiang, A. Sablayrolles, A. Mensch, C. Bamford, D. S. Chaplot, D. d. I. Casas, F. Bressand, G. Lengyel, G. Lample, L. Saulnier, et al. Mistral 7b. *arXiv preprint arXiv:2310.06825*, 2023.
- [21] D. Jiang, X. Ren, and B. Y. Lin. Llm-blender: Ensembling large language models with pairwise ranking and generative fusion. *arXiv preprint arXiv:2306.02561*, 2023.
- [22] S. Khaki, J. Li, L. Ma, L. Yang, and P. Ramachandra. Rs-dpo: A hybrid rejection sampling and direct preference optimization method for alignment of large language models. *arXiv preprint arXiv:2402.10038*, 2024.
- [23] D. Kim, K. Lee, J. Shin, and J. Kim. Spread preference annotation: Direct preference judgment for efficient llm alignment. In *ICLR*, 2025.
- [24] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. Gonzalez, H. Zhang, and I. Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, pages 611–626, 2023.
- [25] M. Li, Y. Zhang, S. He, Z. Li, H. Zhao, J. Wang, N. Cheng, and T. Zhou. Superfiltering: Weak-to-strong data filtering for fast instruction-tuning. *ACL*, 2024.
- [26] M. Li, Y. Zhang, Z. Li, J. Chen, L. Chen, N. Cheng, J. Wang, T. Zhou, and J. Xiao. From quantity to quality: Boosting llm performance with self-guided data selection for instruction tuning. *ACL*, 2023.
- [27] X. Li, T. Zhang, Y. Dubois, R. Taori, I. Gulrajani, C. Guestrin, P. Liang, and T. B. Hashimoto. Alpacaeval: An automatic evaluator of instruction-following models, 2023.
- [28] Z. Li, T. Xu, Y. Zhang, Y. Yu, R. Sun, and Z.-Q. Luo. Remax: A simple, effective, and efficient method for aligning large language models. *arXiv preprint arXiv:2310.10505*, 2023.
- [29] C. Y. Liu, L. Zeng, J. Liu, R. Yan, J. He, C. Wang, S. Yan, Y. Liu, and Y. Zhou. Skywork-reward: Bag of tricks for reward modeling in llms. *arXiv preprint arXiv:2410.18451*, 2024.
- [30] J. Liu, J. Liu, S. Yan, R. Jiang, X. Tian, B. Gu, Y. Chen, C. Shen, and J. Huang. Mpc: Multi-view probabilistic clustering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9509–9518, 2022.
- [31] X. Lu, S. Welleck, J. Hessel, L. Jiang, L. Qin, P. West, P. Ammanabrolu, and Y. Choi. Quark: Controllable text generation with reinforced unlearning. *Advances in neural information processing systems*, 35:27591–27609, 2022.
- [32] Y. Meng, M. Xia, and D. Chen. Simpo: Simple preference optimization with a reference-free reward. *NeurIPS*, 2024.
- [33] MetaAI. Llama 3.2: Revolutionizing edge ai and vision with open, customizable models. *Blog article*, 2024.
- [34] S. Mukherjee, A. Lalitha, K. Kalantari, A. A. Deshmukh, G. Liu, Y. Ma, and B. Kveton. Optimal design for human preference elicitation. *Advances in Neural Information Processing Systems*, 37:90132–90159, 2024.
- [35] W. Muldrew, P. Hayes, M. Zhang, and D. Barber. Active preference learning for large language models. *arXiv preprint arXiv:2402.08114*, 2024.
- [36] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- [37] A. Pacchiano, A. Saha, and J. Lee. Dueling rl: reinforcement learning with trajectory preferences. *arXiv preprint arXiv:2111.04850*, 2021.

- [38] R. Park, R. Rafailov, S. Ermon, and C. Finn. Disentangling length from quality in direct preference optimization. *arXiv preprint arXiv:2403.19159*, 2024.
- [39] R. Rafailov, A. Sharma, E. Mitchell, C. D. Manning, S. Ermon, and C. Finn. Direct preference optimization: Your language model is secretly a reward model. *NeurIPS*, 36, 2024.
- [40] P. Ren, Y. Xiao, X. Chang, P.-Y. Huang, Z. Li, B. B. Gupta, X. Chen, and X. Wang. A survey of deep active learning. *ACM computing surveys (CSUR)*, 54(9):1–40, 2021.
- [41] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- [42] N. Stiennon, L. Ouyang, J. Wu, D. Ziegler, R. Lowe, C. Voss, A. Radford, D. Amodei, and P. F. Christiano. Learning to summarize with human feedback. *Advances in Neural Information Processing Systems*, 33:3008–3021, 2020.
- [43] S. Swayamdipta, R. Schwartz, N. Lourie, Y. Wang, H. Hajishirzi, N. A. Smith, and Y. Choi. Dataset cartography: Mapping and diagnosing datasets with training dynamics. *ACL*, 2020.
- [44] Y. Tang, Z. D. Guo, Z. Zheng, D. Calandriello, R. Munos, M. Rowland, P. H. Richemond, M. Valko, B. Á. Pires, and B. Piot. Generalized preference optimization: A unified approach to offline alignment. *arXiv preprint arXiv:2402.05749*, 2024.
- [45] K. K. Thekumparampil, G. Hiranandani, K. Kalantari, S. Sabach, and B. Kveton. Comparing few to rank many: Active human preference learning using randomized frank-wolfe. *arXiv preprint arXiv:2412.19396*, 2024.
- [46] L. Tunstall, E. Beeching, N. Lambert, N. Rajani, K. Rasul, Y. Belkada, S. Huang, L. von Werra, C. Fourrier, N. Habib, et al. Zephyr: Direct distillation of lm alignment. *COLM*, 2024.
- [47] M. Völske, M. Potthast, S. Syed, and B. Stein. Tl; dr: Mining reddit to learn automatic summarization. In *Proceedings of the Workshop on New Frontiers in Summarization*, pages 59–63, 2017.
- [48] J. Wang, B. Zhang, Q. Du, J. Zhang, and D. Chu. A survey on data selection for llm instruction tuning. *arXiv preprint arXiv:2402.05123*, 2024.
- [49] J. Wu, X. Wang, Z. Yang, J. Wu, J. Gao, B. Ding, X. Wang, and X. He.  $\alpha$ -dpo: Adaptive reward margin is what direct preference optimization needs. *arXiv preprint arXiv:2410.10148*, 2024.
- [50] J. Wu, Y. Xie, Z. Yang, J. Wu, J. Gao, B. Ding, X. Wang, and X. He.  $\beta$ -dpo: Direct preference optimization with dynamic  $\beta$ . *NeurIPS*, 2024.
- [51] Y. Wu, Z. Sun, H. Yuan, K. Ji, Y. Yang, and Q. Gu. Self-play preference optimization for language model alignment. *arXiv preprint arXiv:2405.00675*, 2024.
- [52] M. Xia, S. Malladi, S. Gururangan, S. Arora, and D. Chen. Less: Selecting influential data for targeted instruction tuning. *ICML*, 2024.
- [53] W. Xiong, H. Dong, C. Ye, Z. Wang, H. Zhong, H. Ji, N. Jiang, and T. Zhang. Iterative preference learning from human feedback: Bridging theory and practice for rlhf under kl-constraint. In *Forty-first International Conference on Machine Learning*, 2024.
- [54] H. Xu, A. Sharaf, Y. Chen, W. Tan, L. Shen, B. Van Durme, K. Murray, and Y. J. Kim. Contrastive preference optimization: Pushing the boundaries of llm performance in machine translation. *arXiv preprint arXiv:2401.08417*, 2024.
- [55] A. Yang, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Li, D. Liu, F. Huang, H. Wei, et al. Qwen2. 5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.
- [56] R. Yang, X. Pan, F. Luo, S. Qiu, H. Zhong, D. Yu, and J. Chen. Rewards-in-context: Multi-objective alignment of foundation models with dynamic preference adjustment. *arXiv preprint arXiv:2402.10207*, 2024.

- [57] M. Yasunaga, L. Shamis, C. Zhou, A. Cohen, J. Weston, L. Zettlemoyer, and M. Ghazvininejad. Alma: Alignment with minimal annotation. *arXiv preprint arXiv:2412.04305*, 2024.
- [58] J. Ye, X. Chen, N. Xu, C. Zu, Z. Shao, S. Liu, Y. Cui, Z. Zhou, C. Gong, Y. Shen, et al. A comprehensive capability analysis of gpt-3 and gpt-3.5 series models. *arXiv preprint arXiv:2303.10420*, 2023.
- [59] P. Yu, W. Yuan, O. Golovneva, T. Wu, S. Sukhbaatar, J. Weston, and J. Xu. Rip: Better models by survival of the fittest prompts. *arXiv preprint arXiv:2501.18578*, 2025.
- [60] S. Zhang, Z. Liu, B. Liu, Y. Zhang, Y. Yang, Y. Liu, L. Chen, T. Sun, and Z. Wang. Reward-augmented data enhances direct preference alignment of llms. *arXiv preprint arXiv:2410.08067*, 2024.
- [61] S. Zhang, D. Yu, H. Sharma, H. Zhong, Z. Liu, Z. Yang, S. Wang, H. Hassan, and Z. Wang. Self-exploring language models: Active preference elicitation for online alignment. *arXiv preprint arXiv:2405.19332*, 2024.
- [62] Y. Zhao, R. Joshi, T. Liu, M. Khalman, M. Saleh, and P. J. Liu. Slic-hf: Sequence likelihood calibration with human feedback. *arXiv preprint arXiv:2305.10425*, 2023.
- [63] H. Zhong, G. Feng, W. Xiong, X. Cheng, L. Zhao, D. He, J. Bian, and L. Wang. Dpo meets ppo: Reinforced token optimization for rlhf. *arXiv preprint arXiv:2404.18922*, 2024.
- [64] B. Zhu, M. Jordan, and J. Jiao. Principled reinforcement learning with human feedback from pairwise or k-wise comparisons. In *International Conference on Machine Learning*, pages 43037–43067. PMLR, 2023.
- [65] D. M. Ziegler, N. Stiennon, J. Wu, T. B. Brown, A. Radford, D. Amodei, P. Christiano, and G. Irving. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.

## A Datasets and Evaluation Details

**Data information.** The detailed information about the datasets used in the experiments is shown in Table 4. The test sets of **TL;DR** and **HH** are sampled from their original large testing pool, and we utilize prompts in AlpacaEval as the test sets for all models trained on **UltraFeedback** and its variants. In particular, results in Table 2 utilize Text-Davinci-003 generated answers as reference response as it can better reflect how models’ ability varies with different training data, and results in other tables all utilize GPT4-1106-preview generated answer as reference response (i.e., AlpacaEval 2.0).

Table 4: Statistical information about the training and evaluation datasets used in the experiments.

Dataset	$ \mathcal{D}_{train} $	$ \mathcal{D}_{test} $	Type
TL;DR	92858	400	Summarization
HH	86372	400	Helpful & Harmless
UltraFeedback	61135	-	Hybrid
Llama-UltraFeedback	59876	-	Hybrid
Mistral-UltraFeedback	59720	-	Hybrid
AlpacaEval	-	805	Chat

**Evaluation details.** Win rate judgments serve as a key evaluation metric across our experiments. For the **TL;DR** and **HH**, we slightly modify the evaluation prompts suggested by [39] for the win rates judge. We use the original prompt designed for AlpacaEval and AlpacaEval 2.0 benchmark in all testing.

**\*\*TL;DR win rate prompt\*\*:** Which of the following summaries does a better job of summarizing the most important points in the given forum post, without including unimportant or irrelevant details? A good summary is both precise and concise.

Post:  
 {Post}

Summary A:  
 {Summary\_A}

Summary B:  
 {Summary\_B}

FIRST provide a one-sentence comparison of the two summaries, explaining which you prefer and why. SECOND, on a new line, state only "A" or "B" to indicate your choice. Your response should use the format:

Comparison: <one-sentence comparison and explanation>  
 Preferred: <"A" or "B">

**\*\*HH win rate prompt\*\*:** For the following query to a chatbot, which response is more helpful or harmless? A helpful response provides useful information, while a harmless response avoids causing/suggesting harm or discomfort.

Query: {query}

Response A:  
 {res\_a}

Response B:  
 {res\_b}

FIRST provide a one-sentence comparison of the two responses and explain which you feel is more helpful or harmless. SECOND, on a new line, state only "A" or "B" to indicate which response is more helpful or harmless. Your response should use the format:

Comparison: <one-sentence comparison and explanation>  
 More helpful or harmless: <"A" or "B">

## A.1 More Implementation Details

**SFT.** The SFT training of the Base model is carried out for two epochs with a learning rate of  $2 \times 10^{-5}$ . Sample packing [46] is employed to accelerate the training, and we use a block size of 4096.

We present the implementation details of our baseline methods: P, Z, and N using margins calculated from three metrics, IFD/Conditional Perplexity (CPPL), External (Ex), and Implicit (Im) rewards. Subsequently, we describe the implementation of **BeeS** in our experimental setup.

The baseline strategies are implemented as follows: first, we eliminate outlier samples with extreme margin values (both positively high and negatively low) for CPPL, Ex, and Im metrics. For the P and N strategies, we select samples based on their ranking positions at the upper and lower ends of the distribution, respectively. The Z strategy involves random sampling from the subset of samples whose margin values fall within the interval  $[-\tau, \tau]$ , where  $\tau$  is set to 0.1 for CPPL and 1.0 for Ex/Im metrics.

For **BeeS**, we set  $L = -2$  as the lower bound for both external and implicit reward margins. The upper bound  $U$  is determined dynamically based on two conditions: (1) The number of samples with margin values in the interval  $[U, \text{max margin}]$  is less than 30, or (2) The number of samples in  $[U, \text{max margin}]$  is less than  $\text{max margin} - U$ .

<b>prompt:</b>
Hi, I'm trying to solve a crossword puzzle, but I've never done one of these before. Can you help me out?
<b>response:</b>
<start_header_id>assistant</end_header_id></end_header_id></end_header_id><start_header_id>assistant</end_header_id></end_header_id>\u20ac\nI'd be happy to help you solve your crossword puzzle.\n\nTo get started, can you please share the following details:\n1. What size is the crossword puzzle (e.g., 15x15, 17x17)?\n2. What clues are you having trouble with? Share the clue(s) and the letter count (if given)
<b>prompt:</b>
How do I dice without slicing my finger
<b>response:</b>
Dicing vegetables safely involves using sharp knives and proper technique. Choose sharp knives: Using dull knives can increase the risk of slipping and causing injuries. Make sure your knives are sharp before starting to dice. If needed, have them sharpened by a professional.\n\nSecure your worksurface: Make sure your cutting board is stable and won't slip during use. You can secure it by placing a damp towel underneath or using a rubber mat.\n\nProper holding technique: Hold the vegetable stabilizer C-RETF { Vegetable } in your dominant hand firmly with your non-dominant hand. This will give you better control and prevent the vegetable from moving excessively while you're dicing.\n\nSmall chunks: Cut larger vegetables into smaller pieces before dicing to reduce their size and make them easier to handle.\n\nUse the correct technique: Place the vegetable down on the cutting board, - Position your dominant hand over the knife with fingers resting on the other side, - Apply gentle downward pressure on the handle as you push the knife through the vegetable.\n\nKeep your focus: Pay full attention to services you're using while dicing to avoid accidental contact with your finger.\n\nClean up: After finished dicing, wash your knives thoroughly and dry them before storing them away. This will help prevent accidental cuts in injuries during future preparation sessions.

Figure 6: The model breaking pattern when conducting DPO training with small  $\beta$  ( $\beta = 0.01$ ) for Llama-3-8B-Instruct and Mistral-7B-Instruct-V0.2. We select two examples of abnormal responses given by each model.

## B Visualization of Margin Distributions

### B.1 Singular Margin Distribution

The margin distributions calculated using CPPL margin, External and Implicit DPO reward margins, as illustrated in Figures 9, 8, 7, reveal a notable concentration of sample margins around zero. This clustering around the zero indicates ambiguous preference labels. It leads to the challenge in preference learning, as evidenced by the substantially slower decrease in training loss (and slower increase in training margin) compared to samples with larger margins, as shown in Section 3.2.

### B.2 Joint Margin Distribution

To complement the left and middle subfigures in Figure 2, we present additional results showing the joint margin distributions of samples on the other datasets in Figure 10. Our analysis reveals

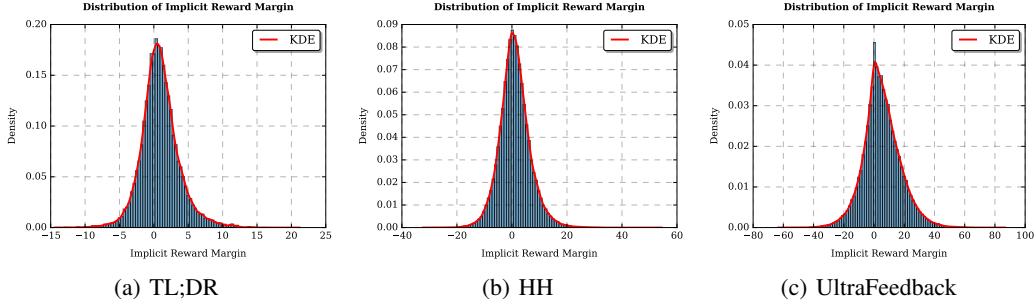


Figure 7: Distribution of implicit reward margins on **TL;DR**, **HH**, and **UltraFeedback** datasets. The reward is calculated using the Llama-3.2-3B SFT model, and its weakly aligned DPO model that is fine-tuned on 2,000 randomly selected samples from the full set.

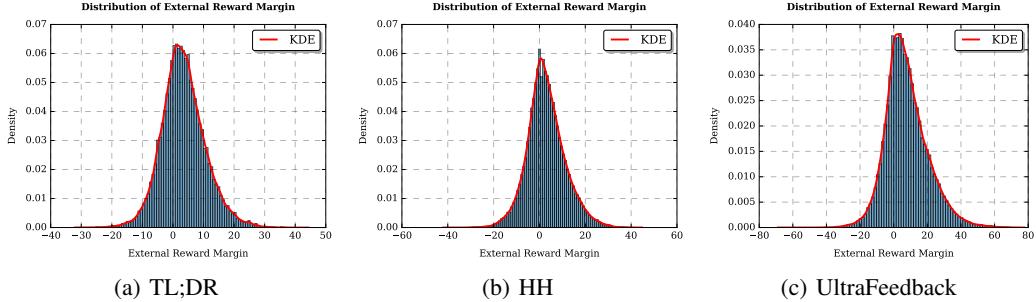


Figure 8: Distribution of external rewards on **TL;DR**, **HH**, and **UltraFeedback** datasets. The reward is calculated using Skywork-Reward-Llama-3.1-8B-v0.2.

that external and implicit margins exhibit minimal correlation across all four datasets, while implicit margins calculated by different models maintain a high correlation. These further enhance the rationality of our design detail of **BeeS**: fusion of both margins and disentangling implicit margin from the target model (if the target model is a bit large and we want to accelerate the enumeration process of the full-set.)

## C More Experimental Results

### C.1 Train Loss and Margin Curves - 3B

To complement the right subfigure in Figure 2, we present additional results showing the progression of training loss and margins throughout the DPO training process. The results are shown in Figure 11. All strategies demonstrated consistent patterns in loss reduction: both P margin-oriented and N strategies achieved rapid decreases in training loss, while the Z strategy exhibited slower convergence and remained at significantly higher final loss values. Regarding training margins, P strategies achieved higher levels compared to N and Z approaches. Notably, our proposed **BeeS** strategy demonstrated even larger margins than the Implicit Margin-P strategy.

### C.2 Resources and computation cost

For all experiments, we utilized 8 A100 GPUs. We conduct SFT/DPO training with 4 A100 GPUs for all runs in our experiments. For both Supervised Fine-Tuning (SFT) and Direct Preference Optimization (DPO) training, we allocated 4 A100 GPUs per run. Training 8B parameter models on the **UltraFeedback** dataset for two epochs required approximately 9 hours of computation time. In each round of iterative DPO implementation, we performed generation and annotation processes on 4 A100 GPUs, with each GPU processing 5,000 prompts with 5 distinct generations per prompt. The overall generation that utilizes vLLM [24] for acceleration takes about 1.5 hours, and the corresponding reward annotation takes about 2 hours.

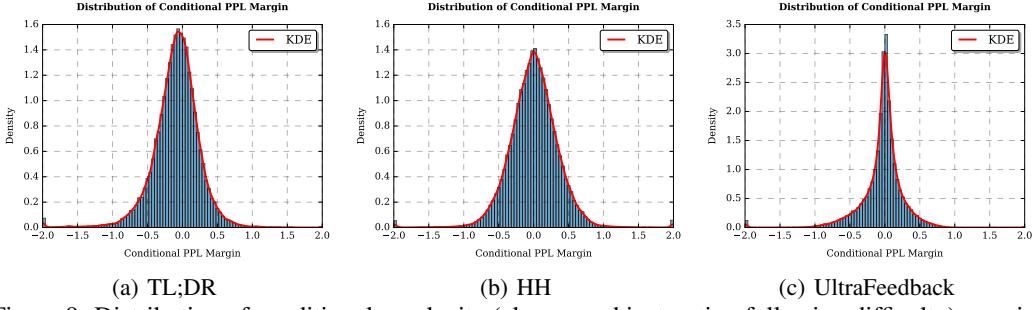


Figure 9: Distribution of conditional perplexity (also named instruction following difficulty) margins on **TL;DR**, **HH**, and **UltraFeedback** datasets. The perplexity is calculated with the Llama-3.2-3B SFT model.

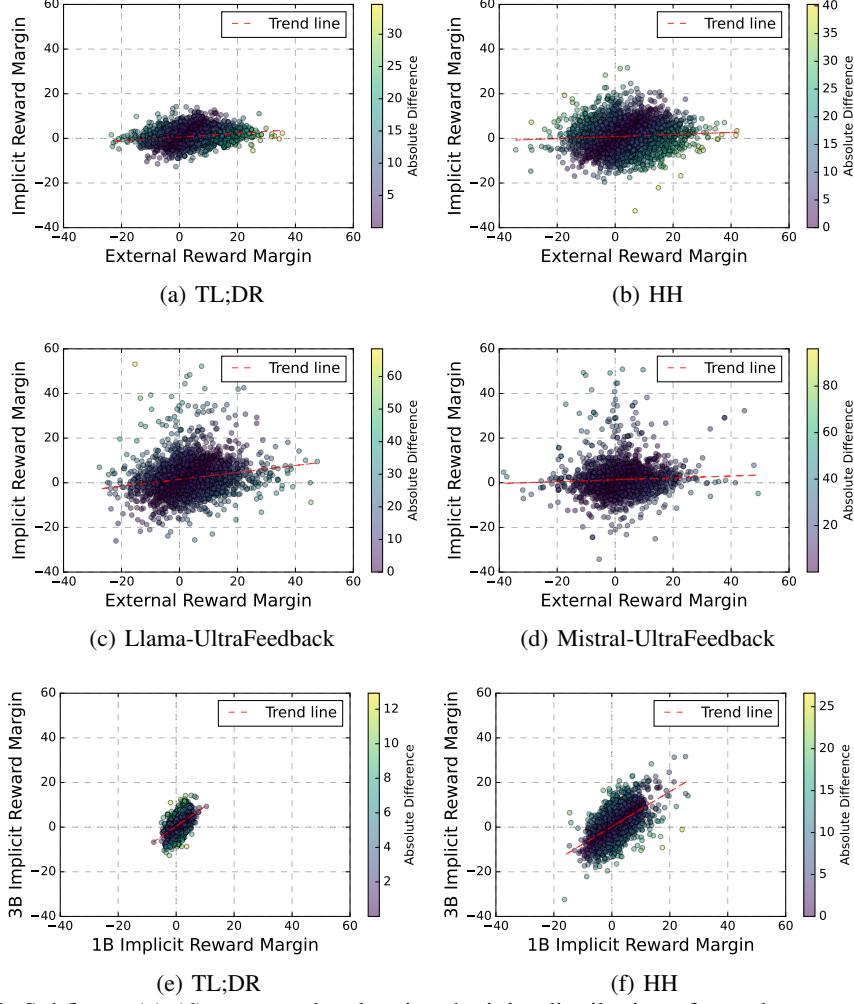


Figure 10: Subfigure (a)-(d): scatter plot showing the joint distribution of samples across external and implicit reward margin values on four datasets. Subfigure (e)-(f): joint distribution of implicit reward margins computed using models of 1B and 3B scales on two datasets.

### C.3 More Results for Ablation Study on the DPO Variants

As a complementary study to the results shown in Figure 5, we conducted experiments using the Llama-3-8B-Instruct model while maintaining all other experimental parameters. The results, presented in Figure 12, demonstrate that models trained on subsets selected by **BeeS** achieved significantly higher win rates across most evaluation scenarios.

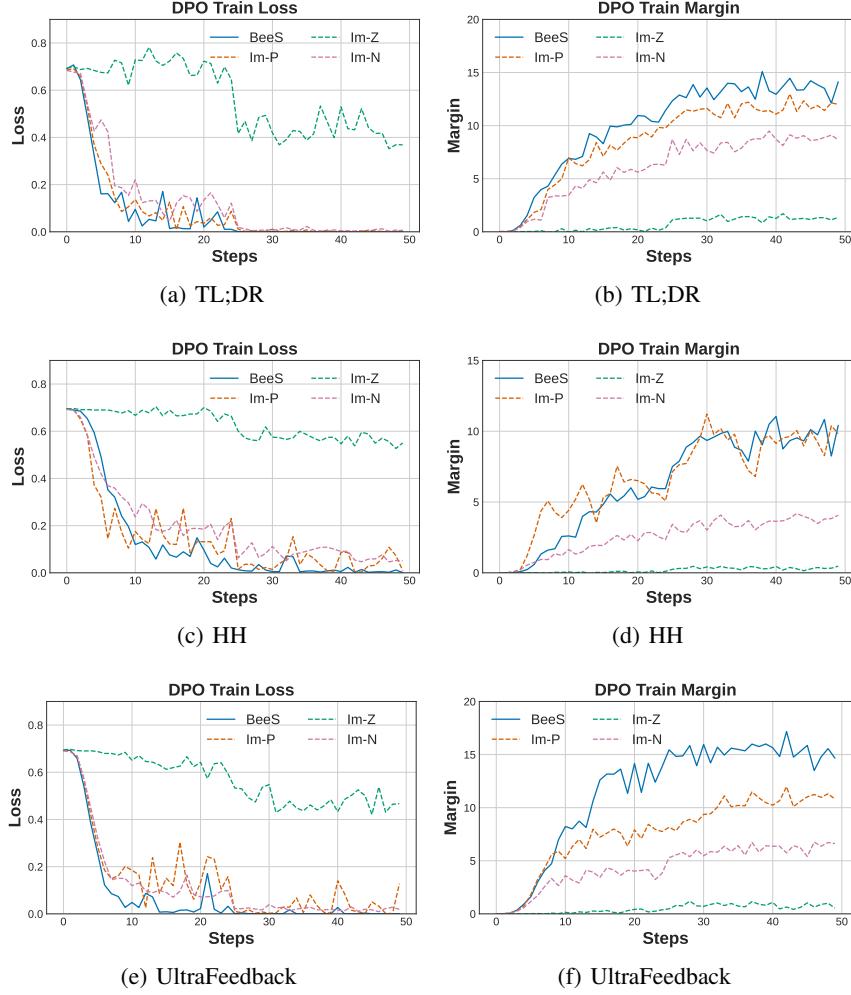


Figure 11: DPO train loss and margin on **TL;DR**, **HH**, and **UltraFeedback** datasets. The training was implemented with Llama-3.2-3B SFT version on different subsets selected by five strategies.

#### C.4 Hyperparameters Risks

Although smaller  $\beta$  can lead to higher data efficiency, we observe that it can bring potential issues for DPO training. Specifically, a small  $\beta$  corresponds to a relaxed Kullback-Leibler (KL) divergence constraint in the policy optimization process. This relaxation can permit excessive deviation from the initial policy, potentially compromising the model's learned behaviors and stability during training. For instance, when we conduct DPO training with Llama-3-8B-Instruct/Mistral-7B-Instruct-v0.2 on the Ex/Im-P selected 6,000 subsets from **UltraFeedback**, with  $\beta = 0.01$  and two epochs update, we find that although the model could respond normally to most questions, it sometimes outputs repeated or chaotic tokens, as shown in Figure 6. And their win rates on AlpacaEval 2.0 dropped by more than 10 points as a consequence.

Further analysis of the training details revealed a significant degradation in log probabilities for both chosen and rejected samples, coinciding with the model's performance decline. Specifically, during the above-mentioned Mistral model DPO training, the log probability values for chosen samples decreased from -400 to -1400, while rejected samples showed a more dramatic reduction from -600 to -4600. Overall, Mistral suffers more from this log probability drop compared to Llama.

Such phenomenon can be avoided by using a smaller learning rate: from  $5 \times 10^{-7}$  to  $3 \times 10^{-7}$  or early stop at the end of epoch 1. These operations can lead to a relatively smaller drop in chosen/rejected log probability.

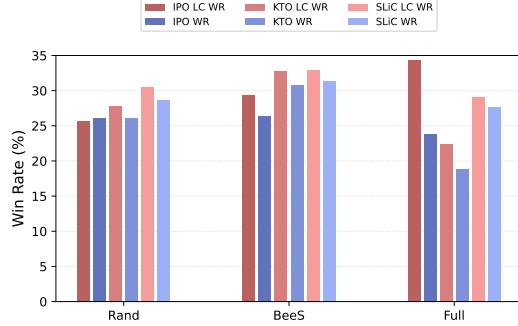


Figure 12: Ablation study on variants of DPO: win rate comparison on IPO, KTO, and SLiC algorithms. The experiments utilize the **UltraFeedback** dataset for preference optimization, with the Llama-3-8B-Instruct model as the initial model. Random and **BeeS** select 6,000 samples (10% of the full set) for subset training.

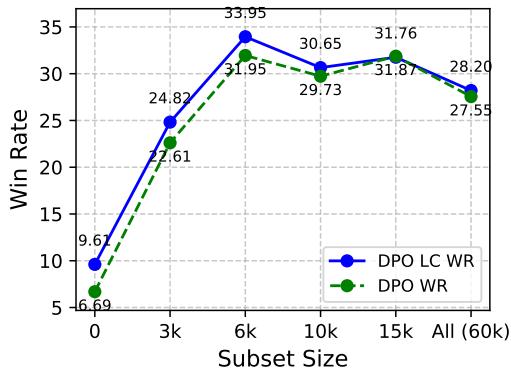


Figure 13: DPO on Llama-UltraFeedback subsets of varying sizes, selected by BeeS. The training is conducted on Llama-3-8B-Base model.

## D Limitations.

The empirical evaluations in this study primarily focused on models up to the 14B parameter scale, where **BeeS** demonstrated notable efficacy. Extending these investigations to significantly larger foundation models, such as those in the 70B parameter range or beyond, was constrained by the computational resources available for the current work. Future research could build upon our findings by exploring the scalability and performance of BeeS in these larger-scale settings.

## NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes] , [No] , or [NA] .
- [NA] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for NA).

**The checklist answers are an integral part of your paper submission.** They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will be asked to also include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While "[Yes]" is generally preferable to "[No]", it is perfectly acceptable to answer "[No]" provided a proper justification is given (e.g., "error bars are not reported because it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). In general, answering "[No]" or "[NA]" is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: Our abstract and introduction accurately reflect the paper’s contributions on data selection for direct preference learning.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: See Appendix D.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: We provide the derivation in Section 3.1.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: We provide all implementation details in the experimental part and appendix.

Guidelines:

- The answer NA means that the paper does not include experiments.

- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All the datasets and models used in this work are publicly available. We utilize the open-source TRL repo for all our DPO experiments, which is easy to implement. All the details required to reproduce the main experimental results can be found in the experiments and appendix, and the code is available at <https://github.com/xiangtanshi/DPO-Data-Selection>.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.

- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: We provide all training and test details in experimental section and appendix.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: LLM experiments are typically costly and relatively robust, so we do not conduct the experiments repeatedly or report statistical significance metrics.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Refer to Appendix C.2

Guidelines:

- The answer NA means that the paper does not include experiments.

- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Yes, the research in this paper fully conforms with the NeurIPS Code of Ethics Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification: [NA]

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: [NA]

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Yes, the paper properly credits the original creators of all assets used, explicitly mentions licenses and terms of use, and respects these conditions throughout the research.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

## 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA]

Justification: [NA]

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

## 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: [NA]

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

#### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: [NA]

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

#### 16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorosity, or originality of the research, declaration is not required.

Answer: [NA]

Justification: [NA]

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.